

# On-device Artificial Intelligence Neural Style Transfer

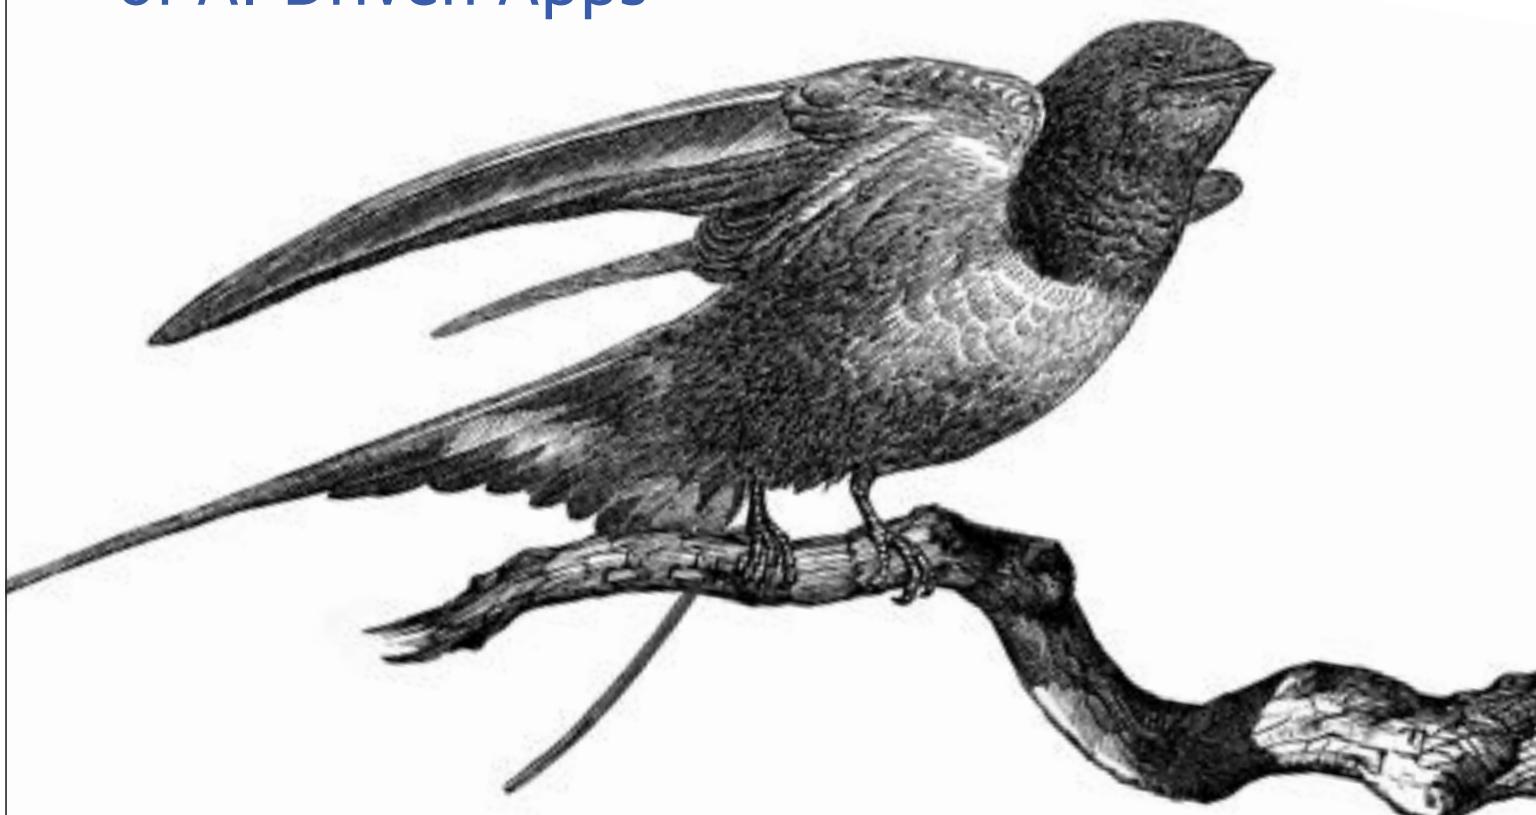
with Mars Geldard &  
Dr Paris Buttfield-Addison



O'REILLY®

# Practical Artificial Intelligence with Swift

From Fundamental Theory to Development  
of AI-Driven Apps



Mars Geldard, Jonathon Manning,  
Paris Buttfield-Addison & Tim Nugent



@parisba

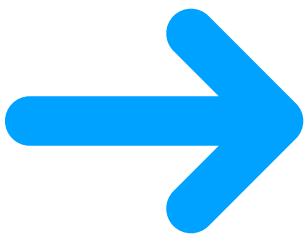
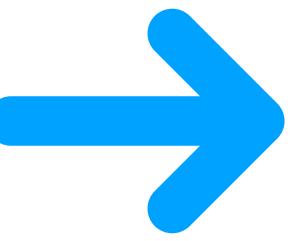


@themartianlife

# Neural Style Transfer

# Neural Style Transfer

(Image A Content + Image B Style = Image C)



# **NST is just our case study**

This is really a talk about on-device ML

# This isn't really an Apple talk

They're just the ~~biggest~~ most interesting player... for now







# Goals of this talk

- Discuss how useful, effective, and sensible the idea of on-device artificial intelligence is
- Have a brief diversion to discuss what Style Transfer is, and gloss over how it works
- Demonstrate that on-device AI works, by implementing Style Transfer for it on iOS
- Show off how useful task-focused approaches to AI can be, and how useful Turi Create is

*Part 1*

# On-device Machine Learning

*Part 2*

# Neural Style Transfer

*Part 3*

# Demonstration

*Part 4*

# The Future

*Part 1*

# On-device Machine Learning

**This is the big takeaway of our talk...**

# On-device AI

# On-device artificial intelligence...

... usually means on-device Machine Learning

# ML refresher

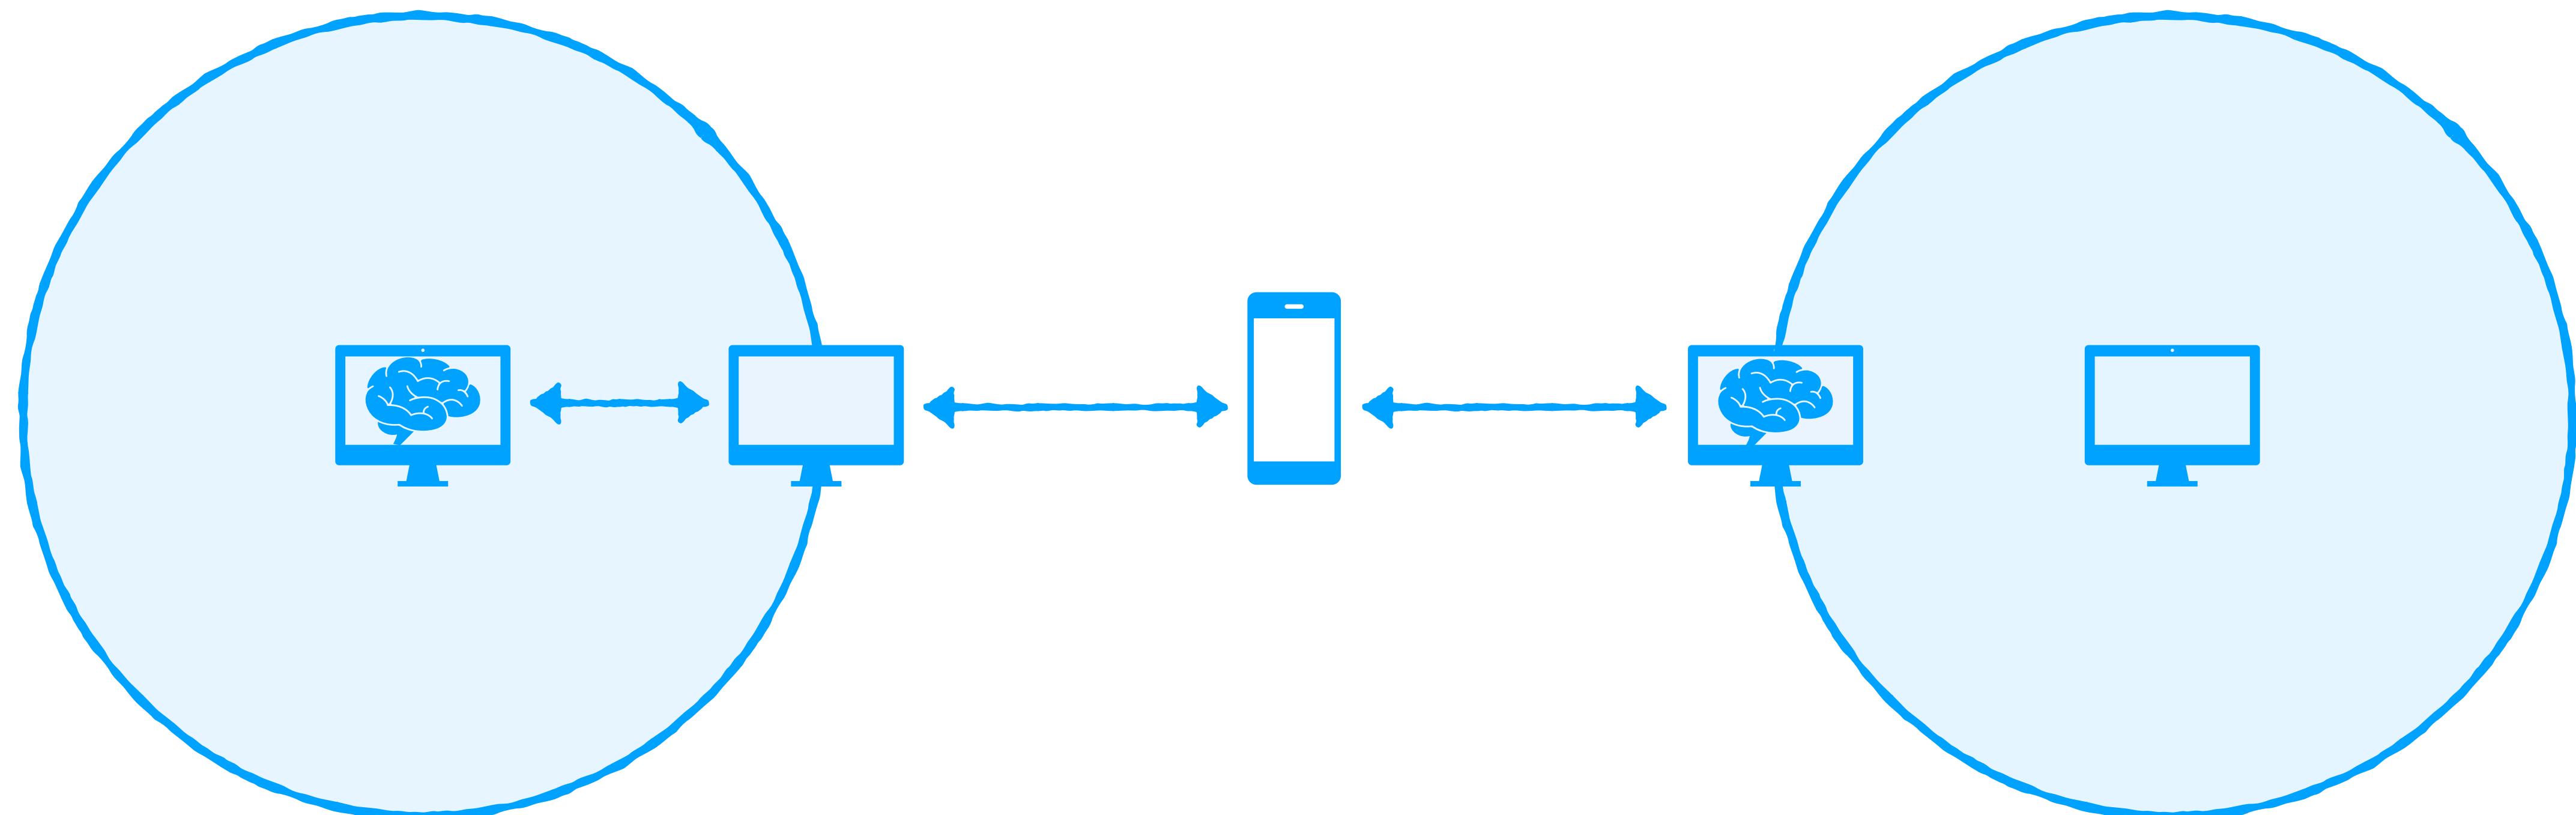
# **Training and Inference**

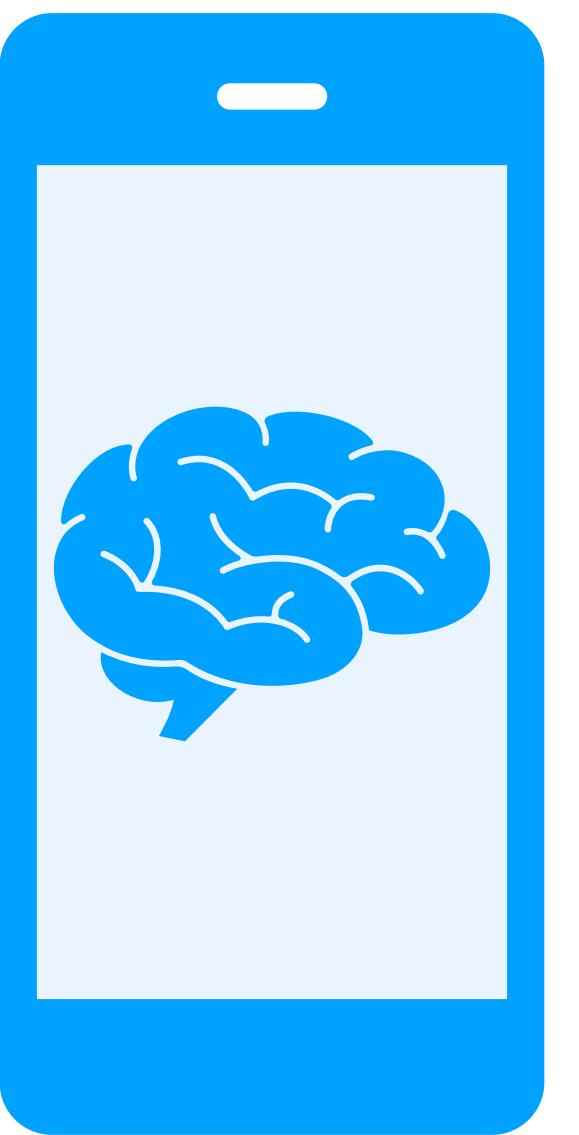
# On-device machine learning...

... usually means on-device consultation of a pre-trained model (inference)

# On-device inference

# This is just edge computing





# **“Inference on the edge”**

... is another phrase for “on-device ML”

# What can you do on device?

# What can you do on device?

- Image recognition
- Object localisation
- Speech/character recognition/translation
- Gesture recognition
- Text prediction/classification

# Benefits?

- Less (laggy) network requests
- High-speed lag-free user experiences
- Reliability (works offline)
- Cheap(er)
- Privacy



# Negatives

- Bigger apps
- Battery use, maybe
- Creation of initial training data
- Model training/retraining
- App updates need to ship the model (or download the model)



# Training AND Inference?

- Much harder to do with current hardware, but still possible
- iOS's Photos app
- FaceID + TouchID
- Set of conditions that must be met for the model to attempt training with new data

*Part 1*

# On-device Machine Learning

*Part 2*

# Neural Style Transfer

*Part 3*

# Demonstration

*Part 4*

# The Future

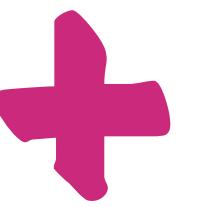
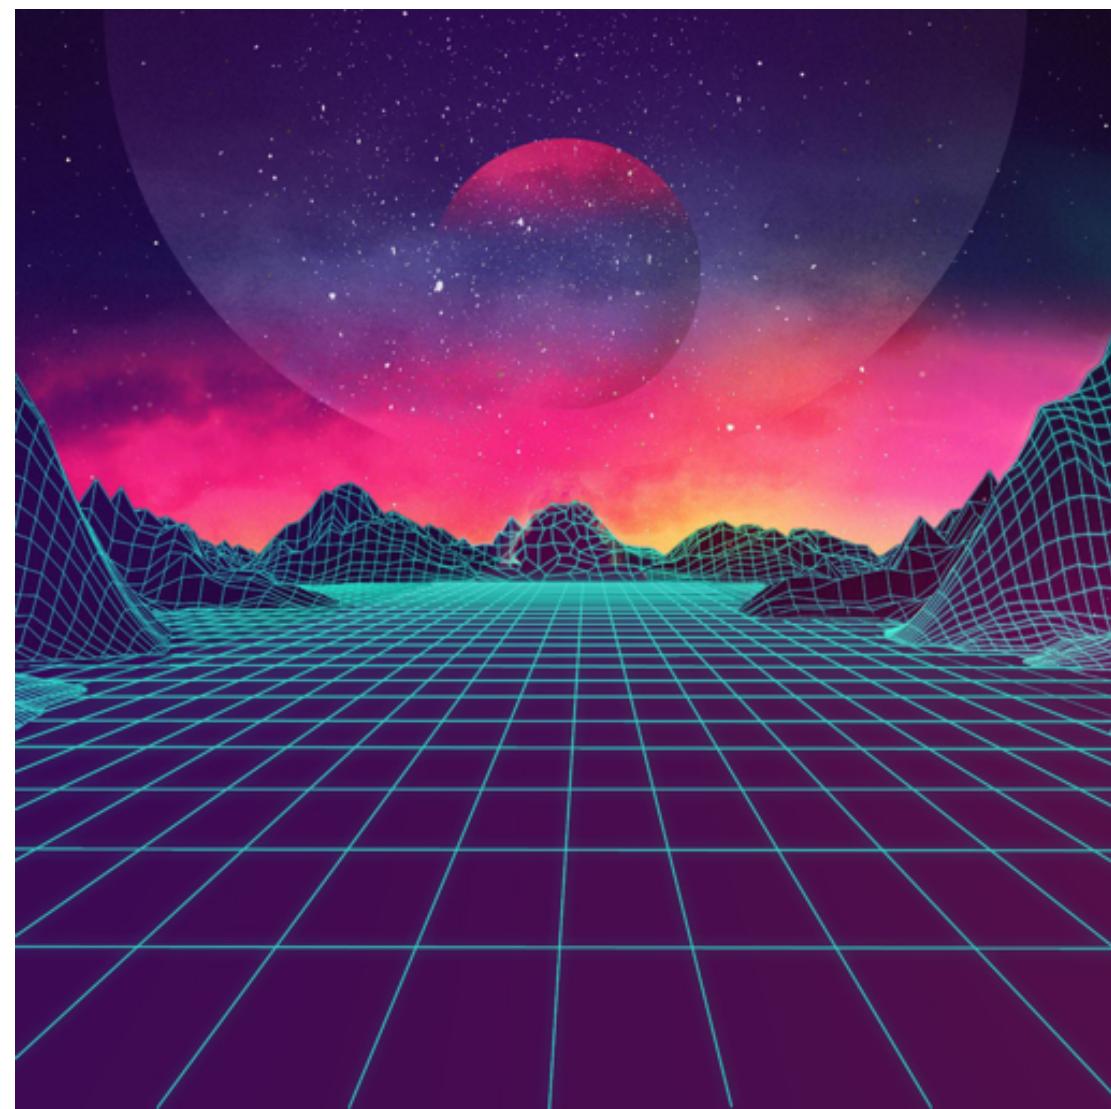
*Part 2*

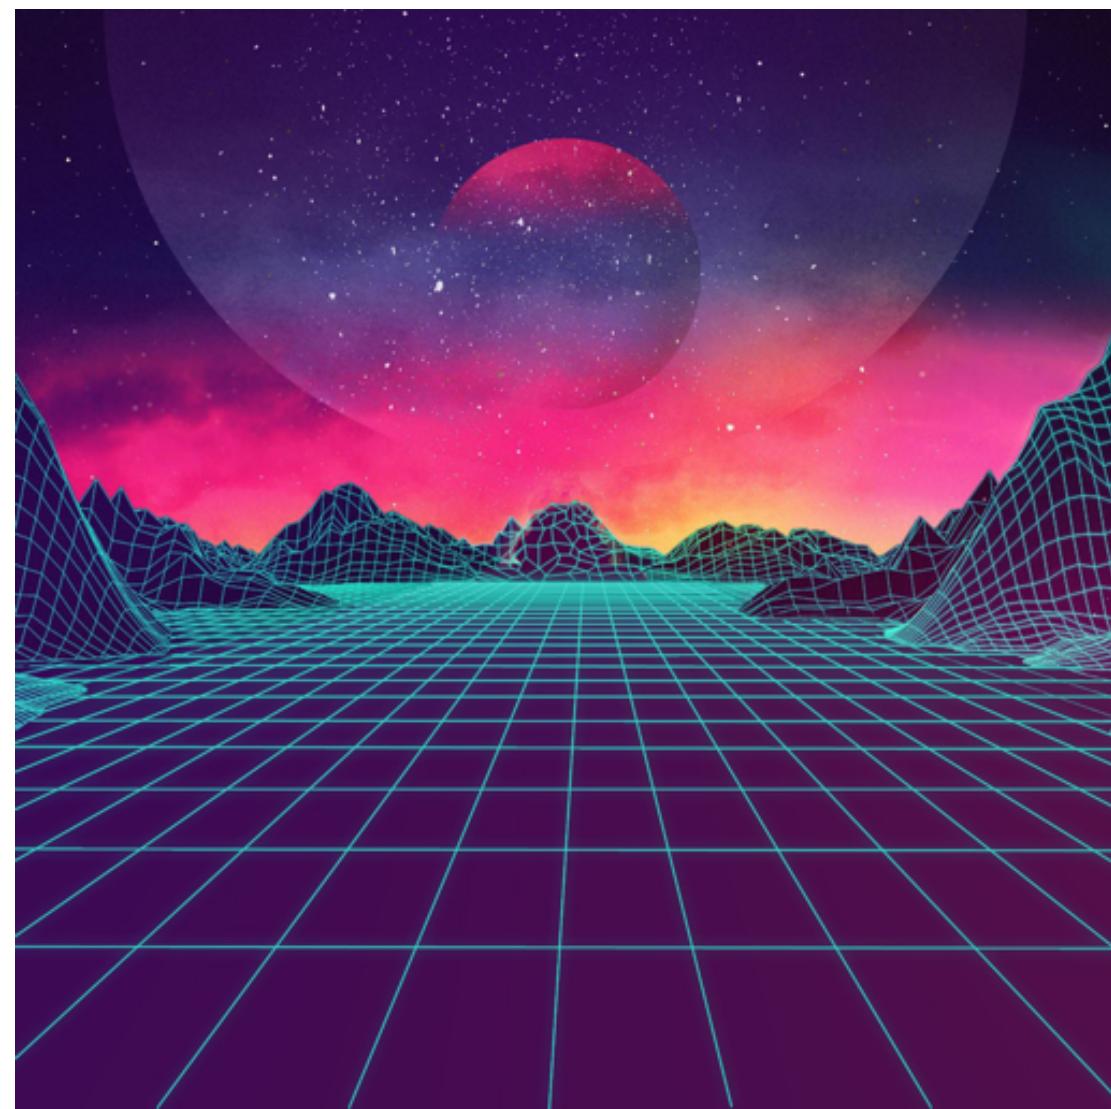
# Neural Style Transfer

**“Style Transfer is a task wherein the stylistic elements of a style image are imitated onto a new image while preserving the content of the new image.”**

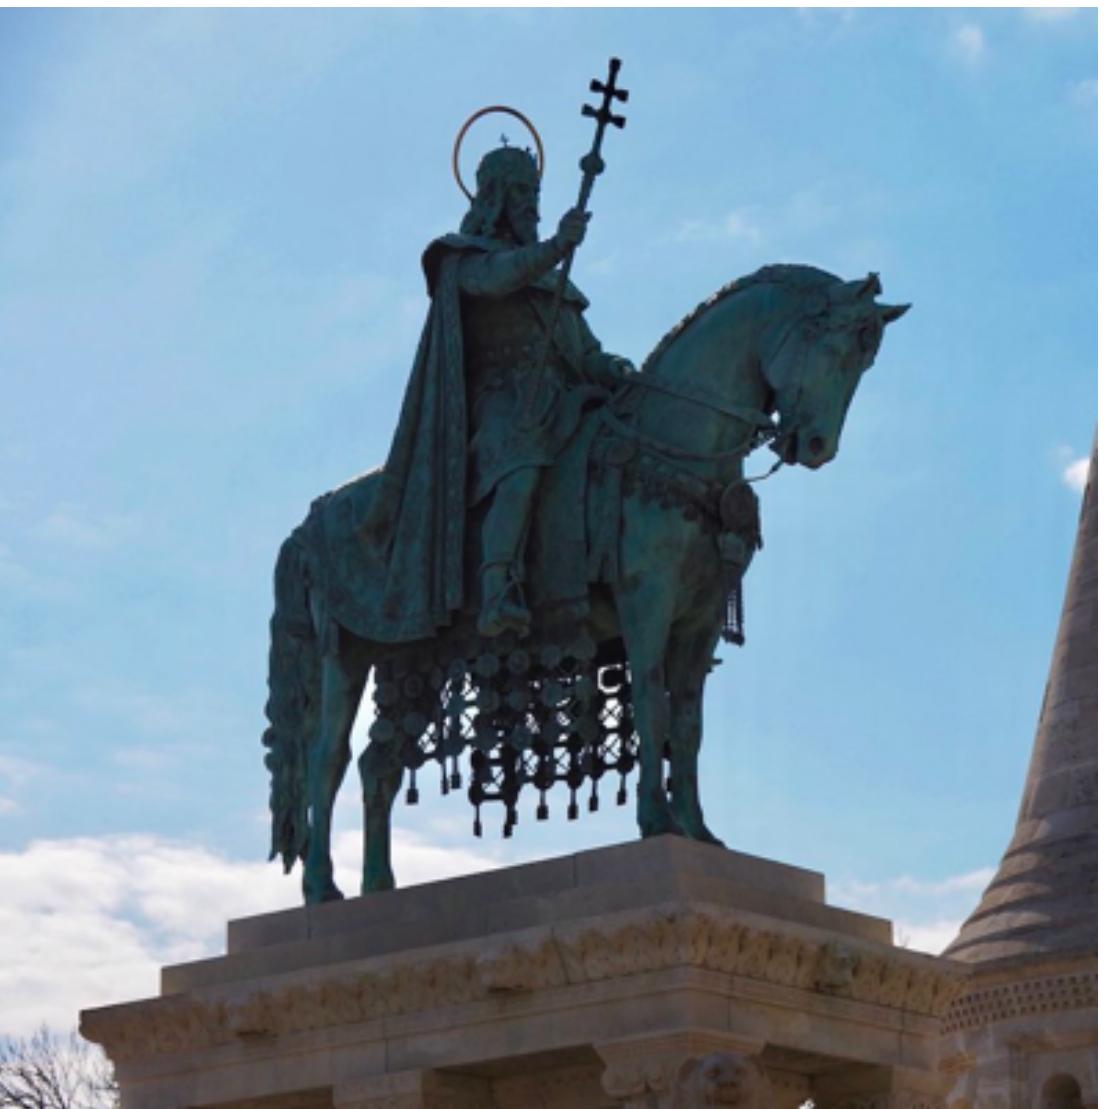
*–TuriCreate documentation*



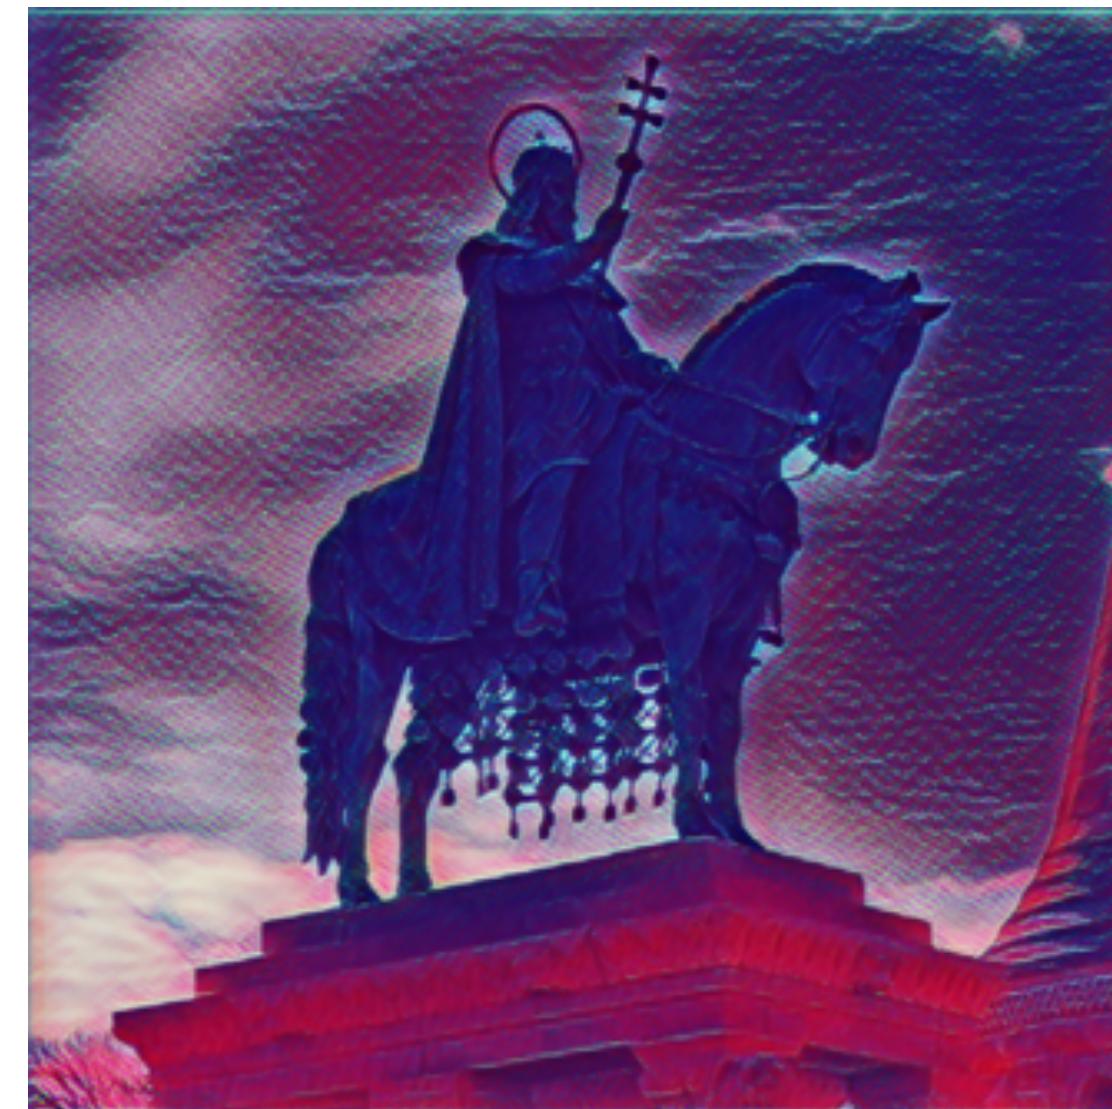




+



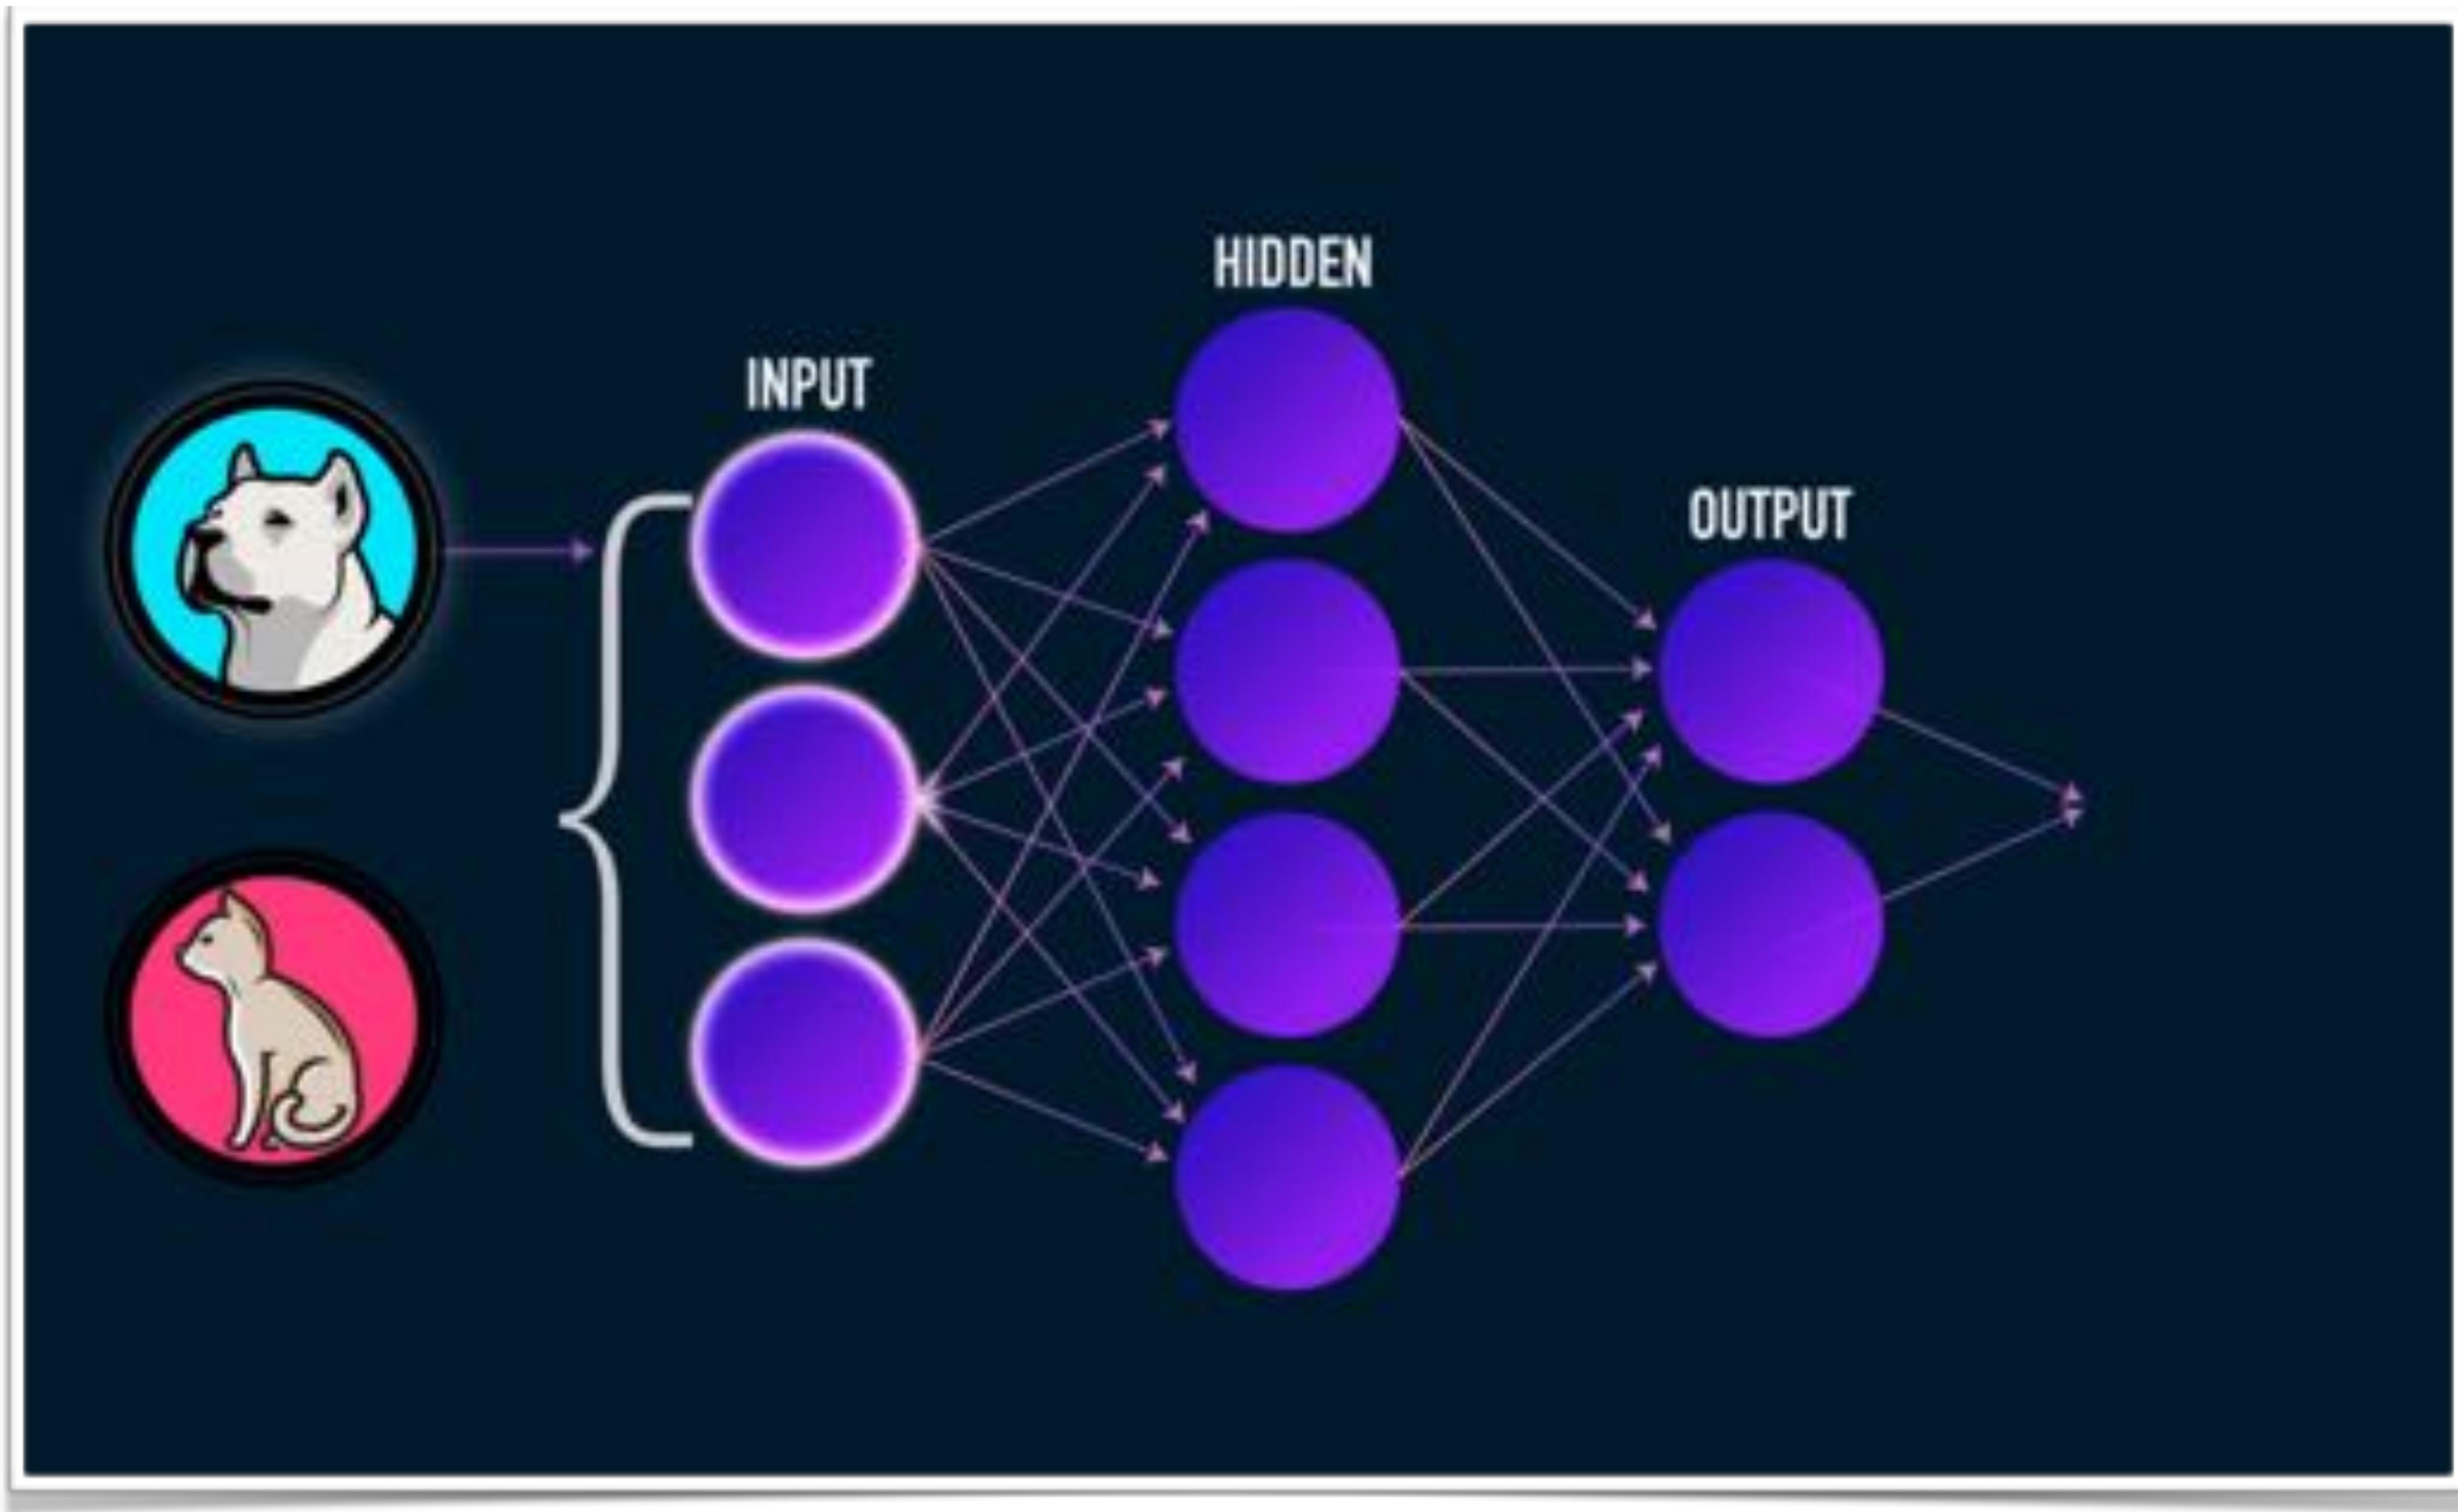
=



# OK, great. How's it work?

Technically, what's going on?

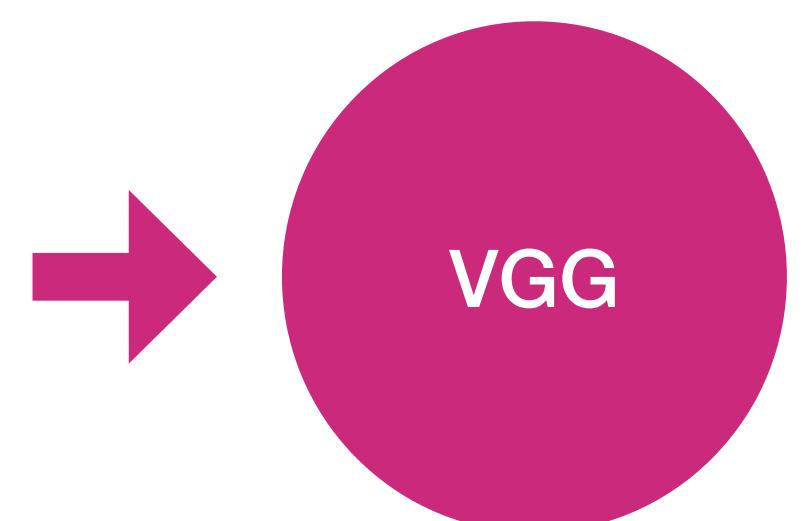
# Neural Network refresher

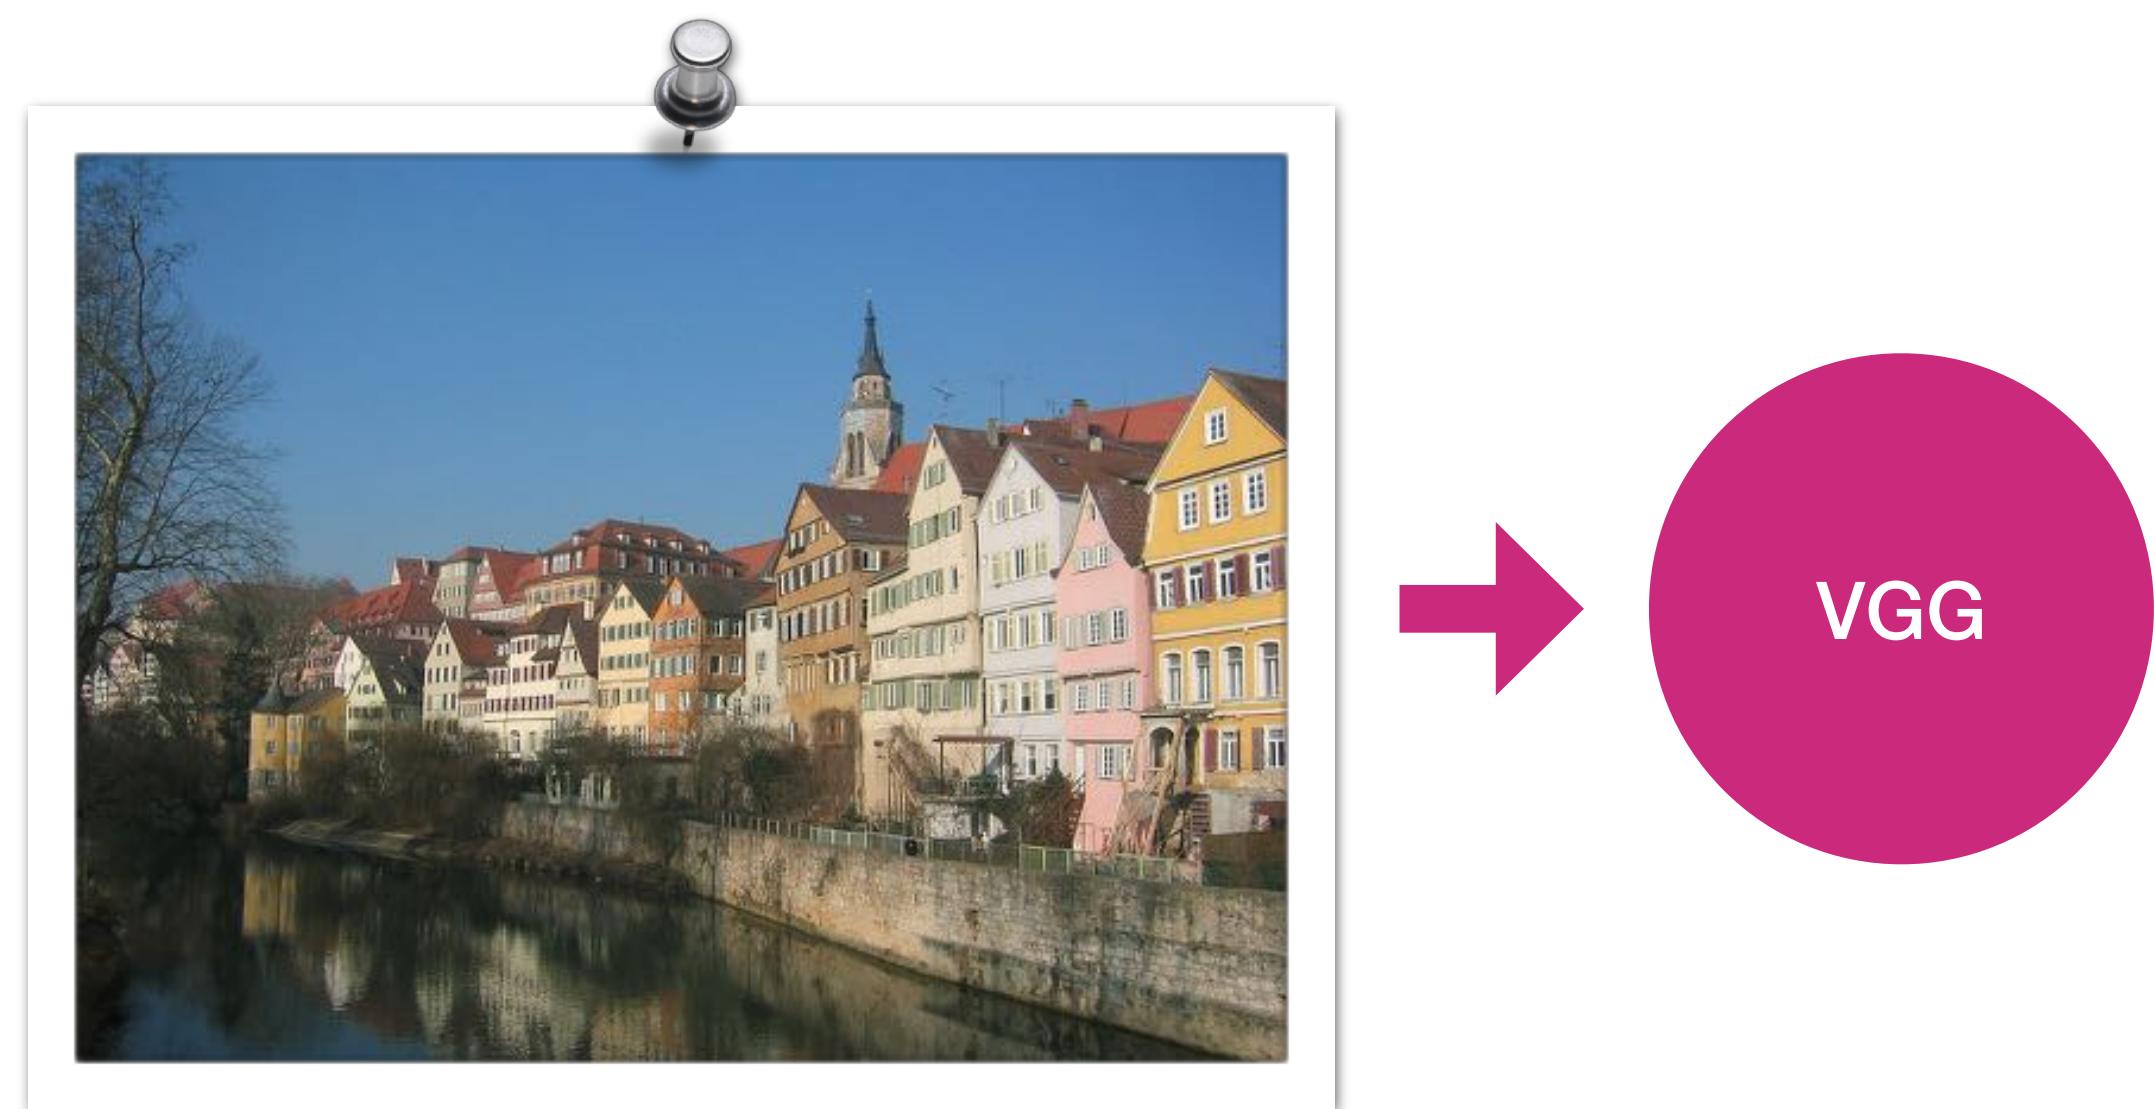
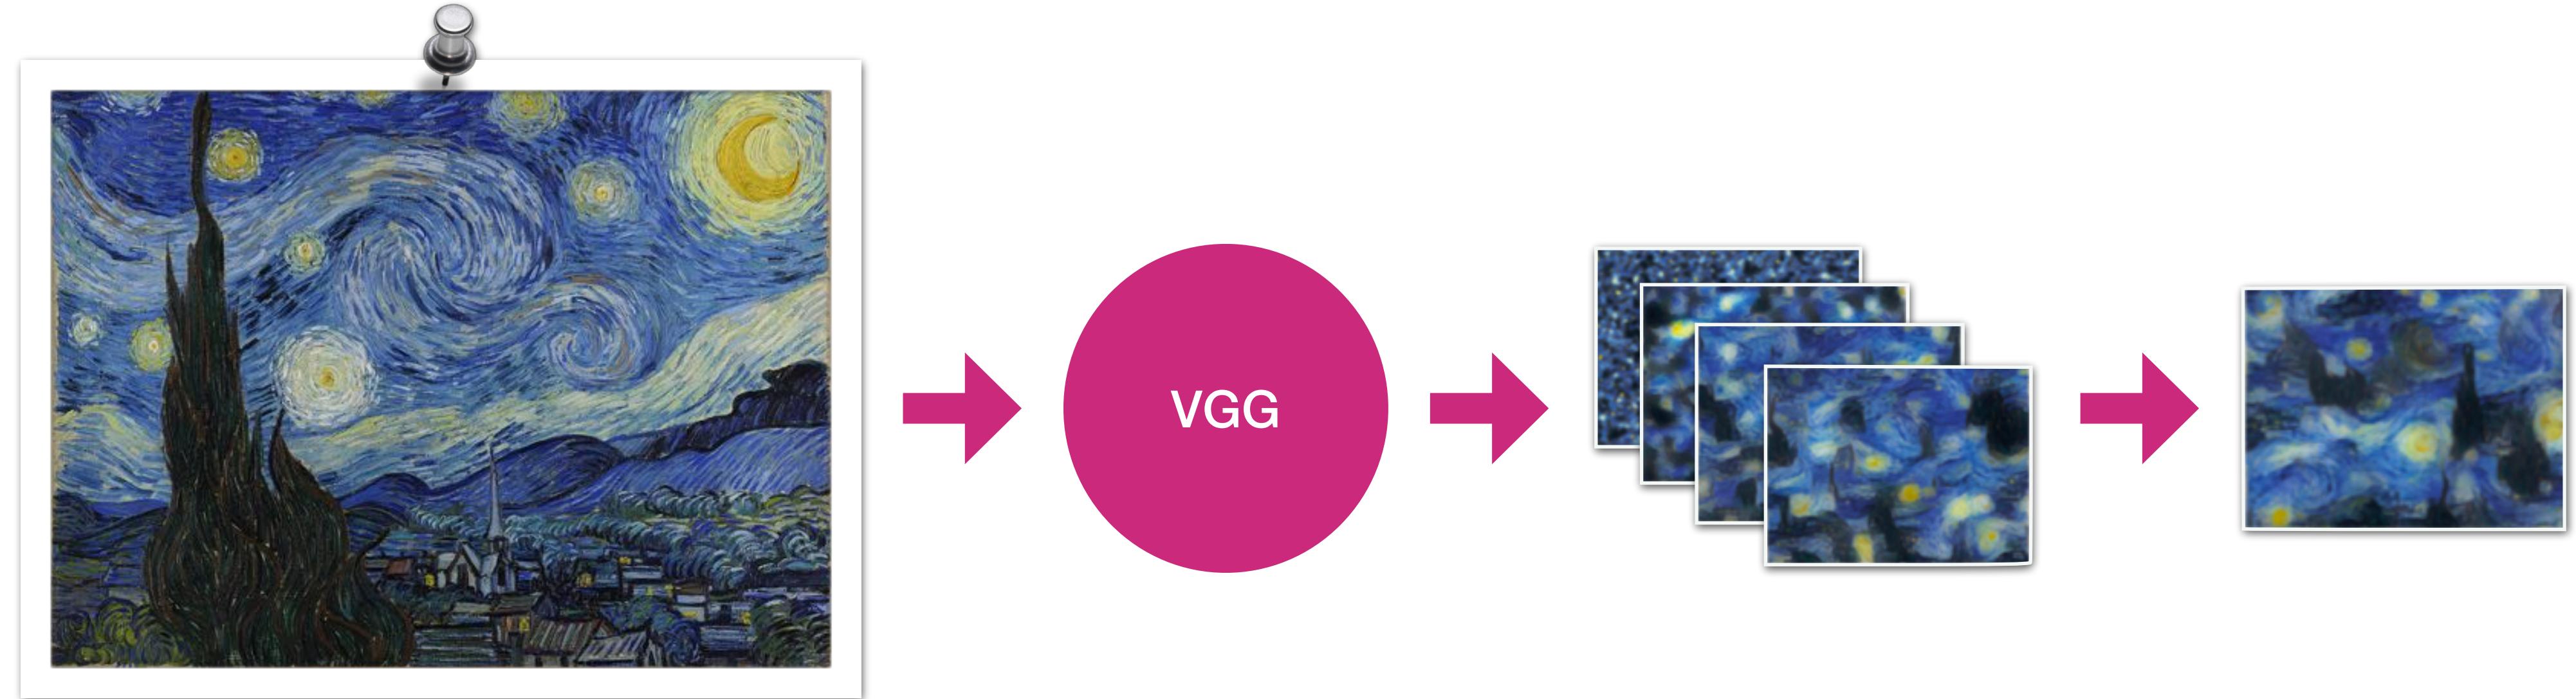


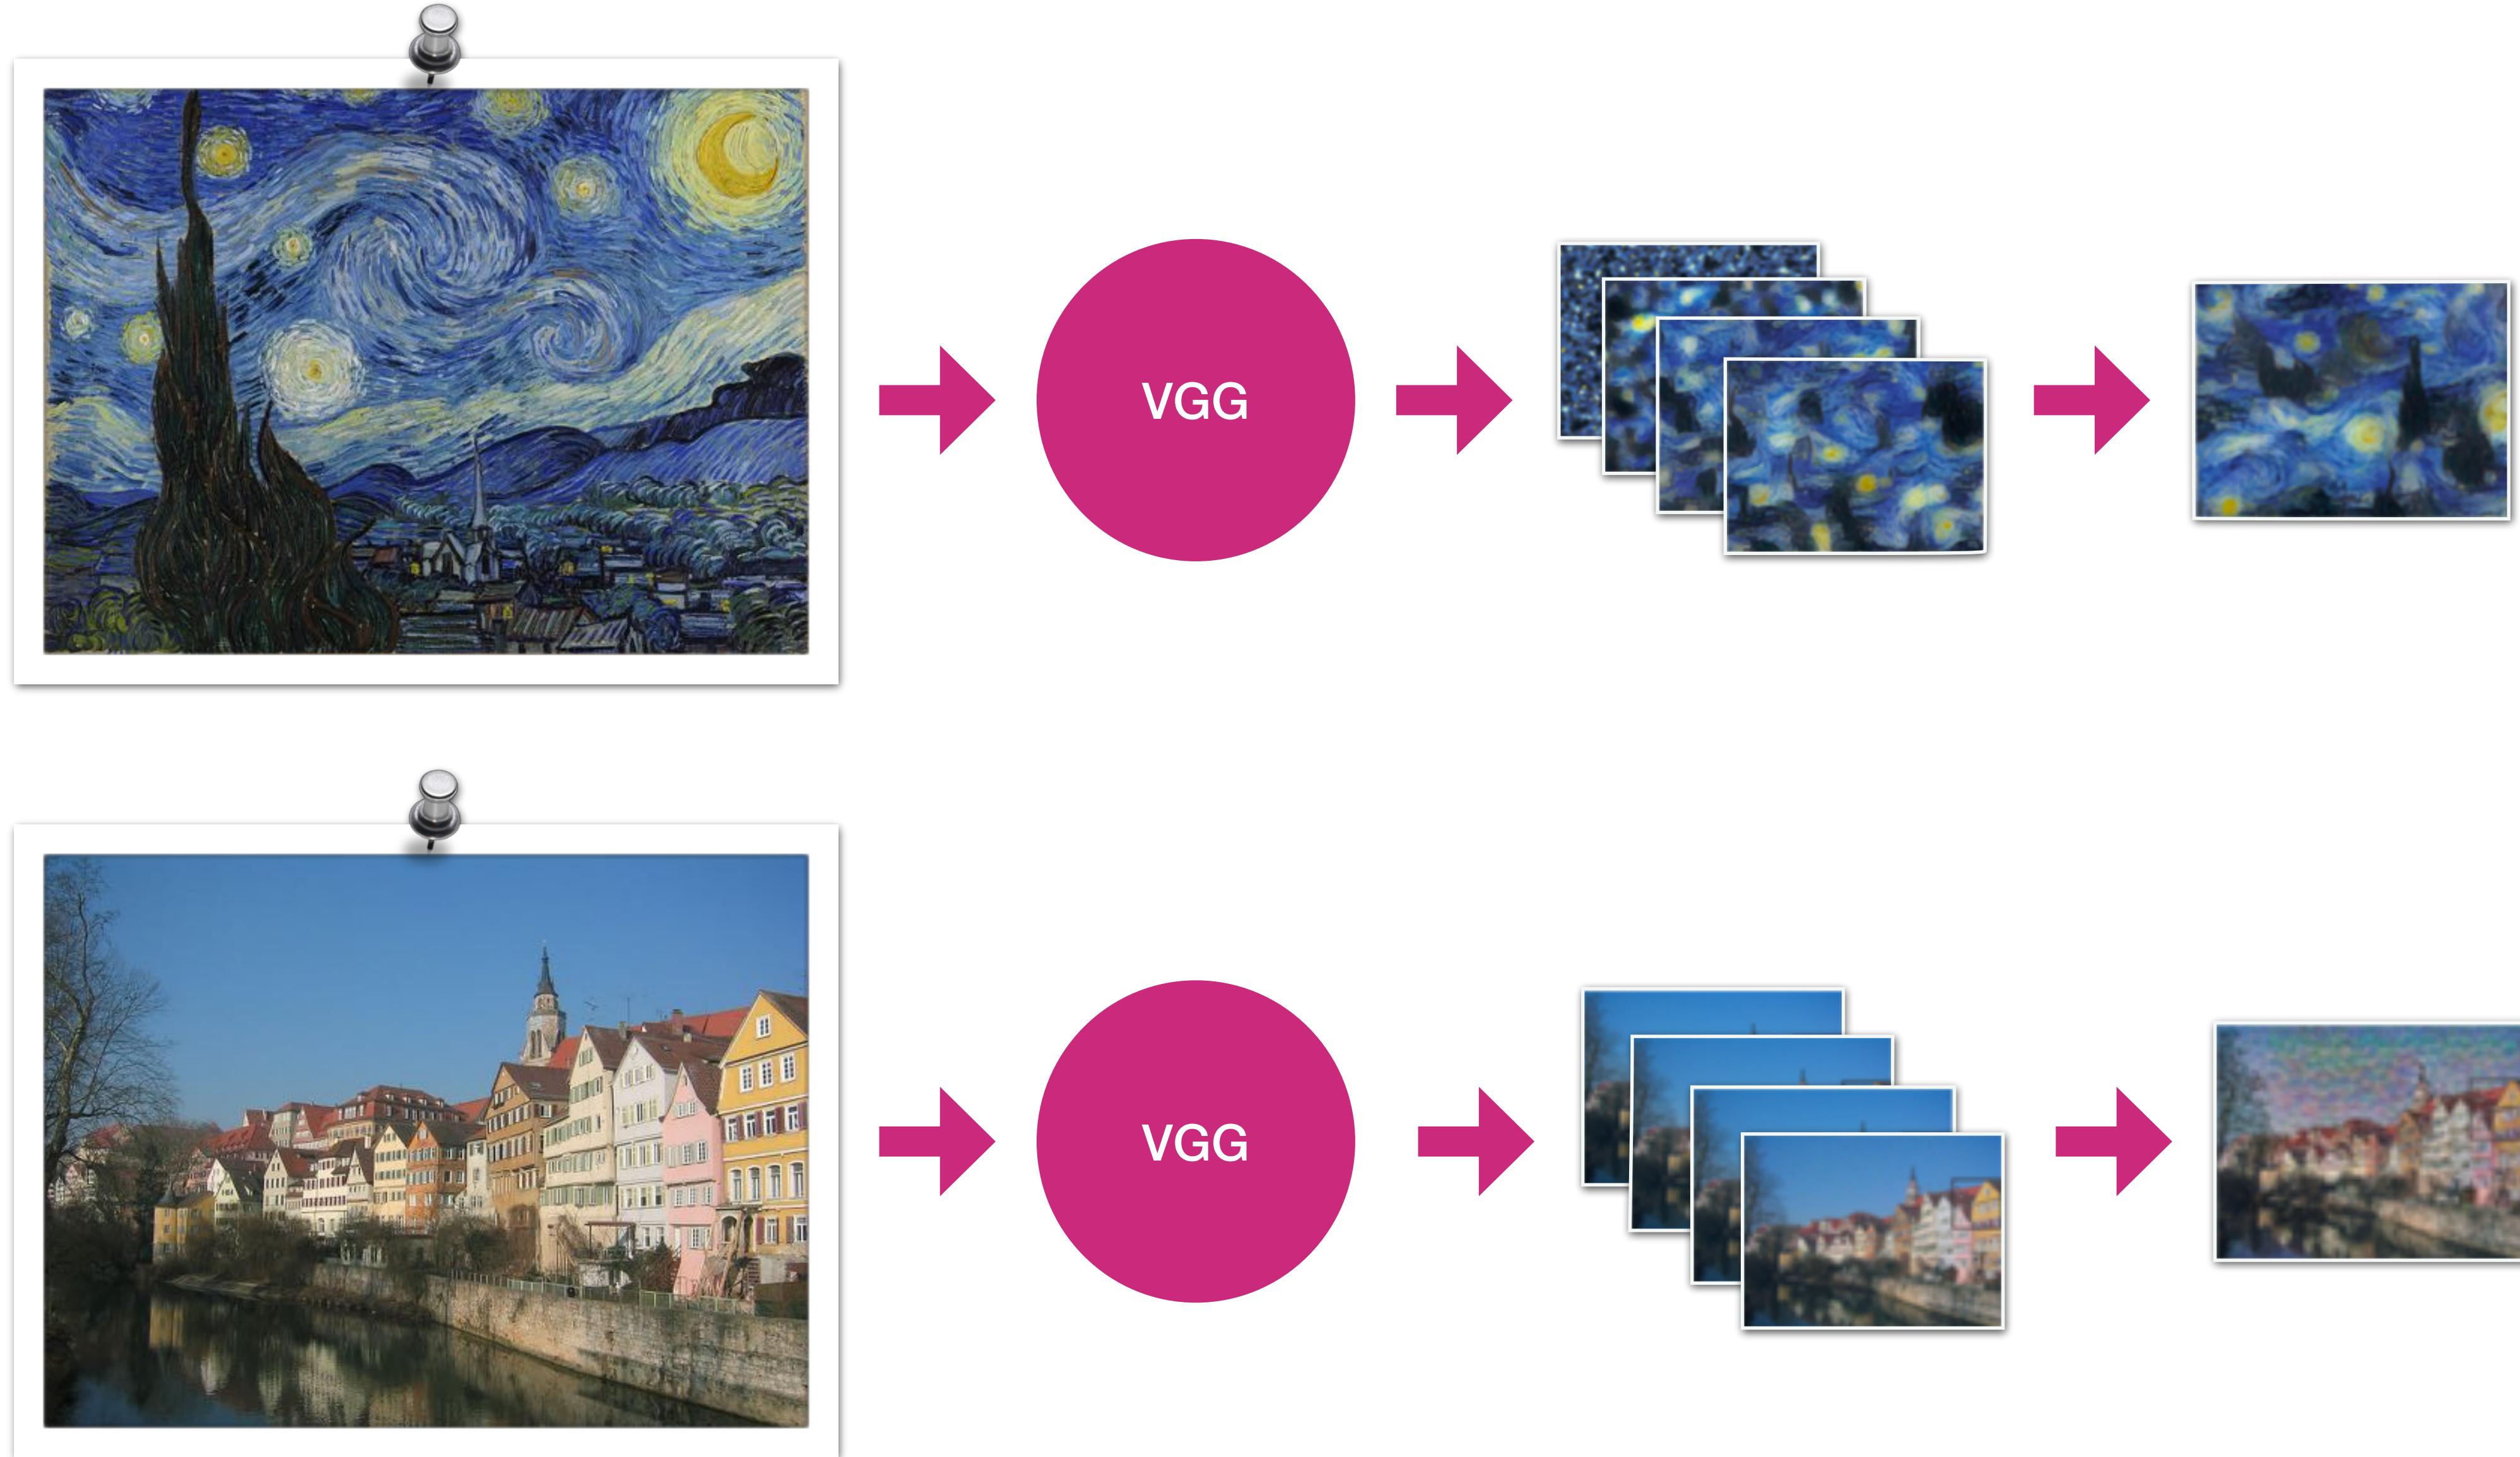
# Using a model to generate?

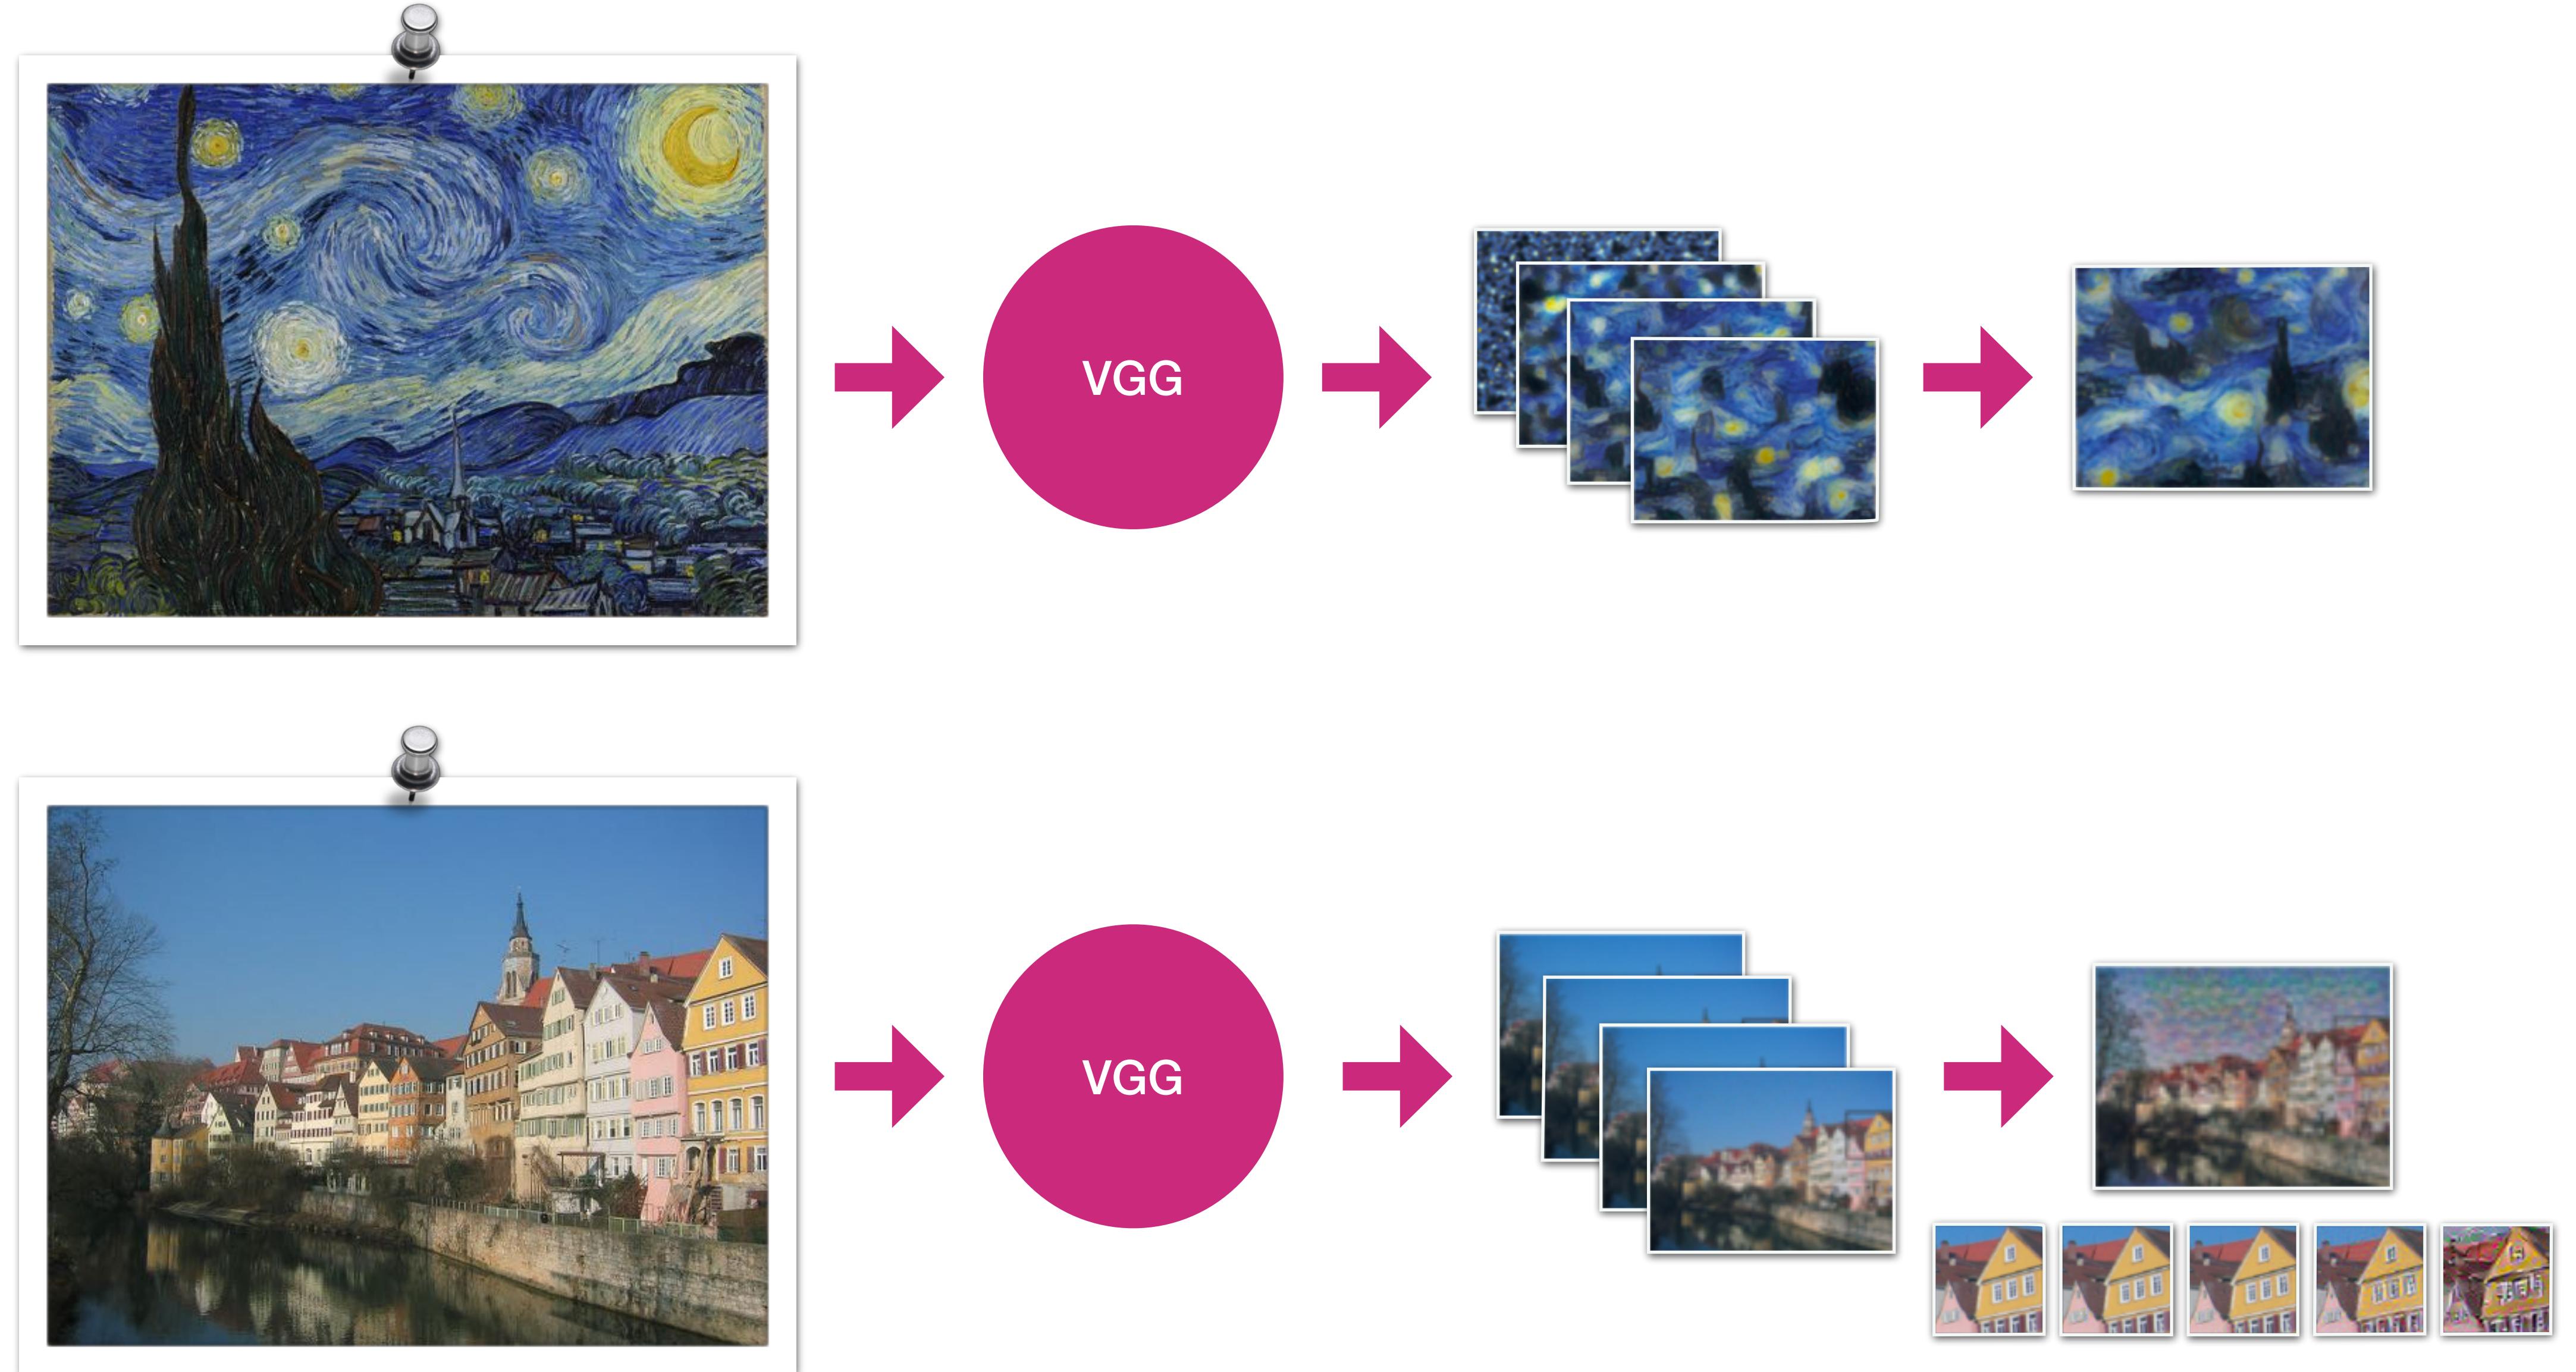


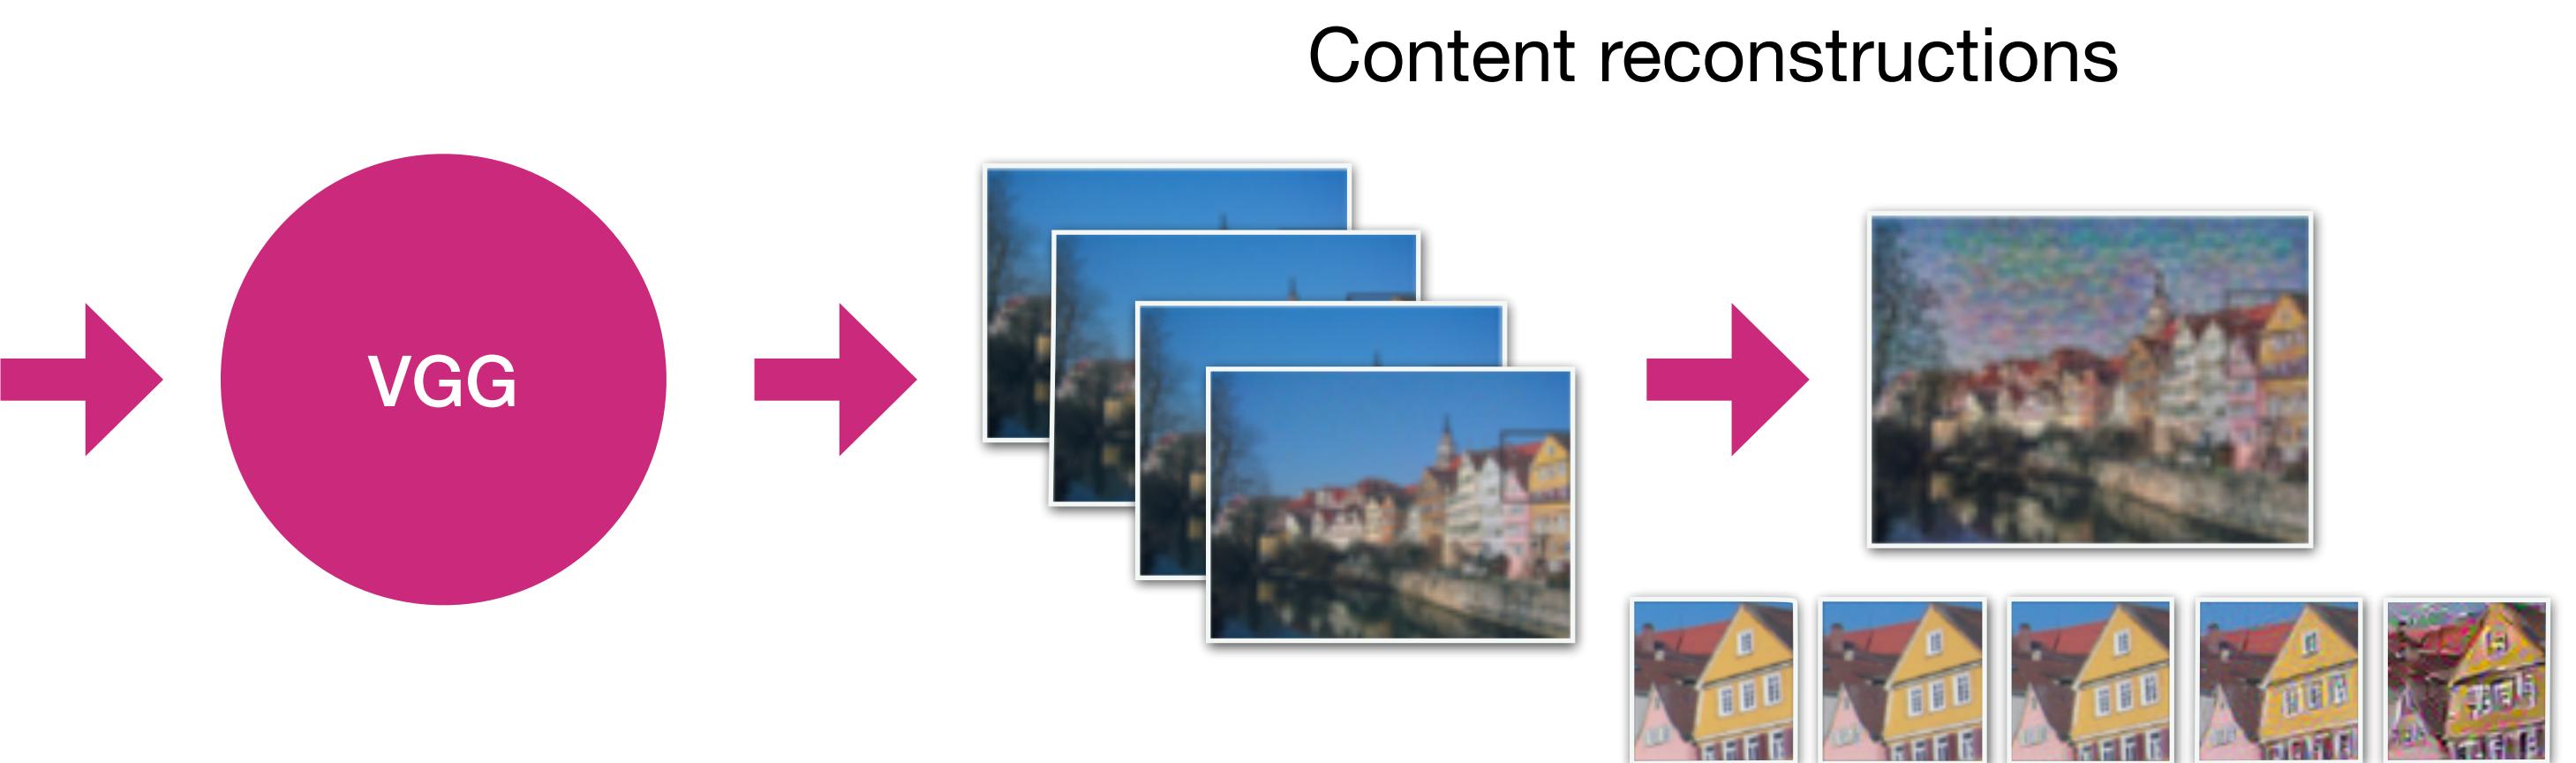
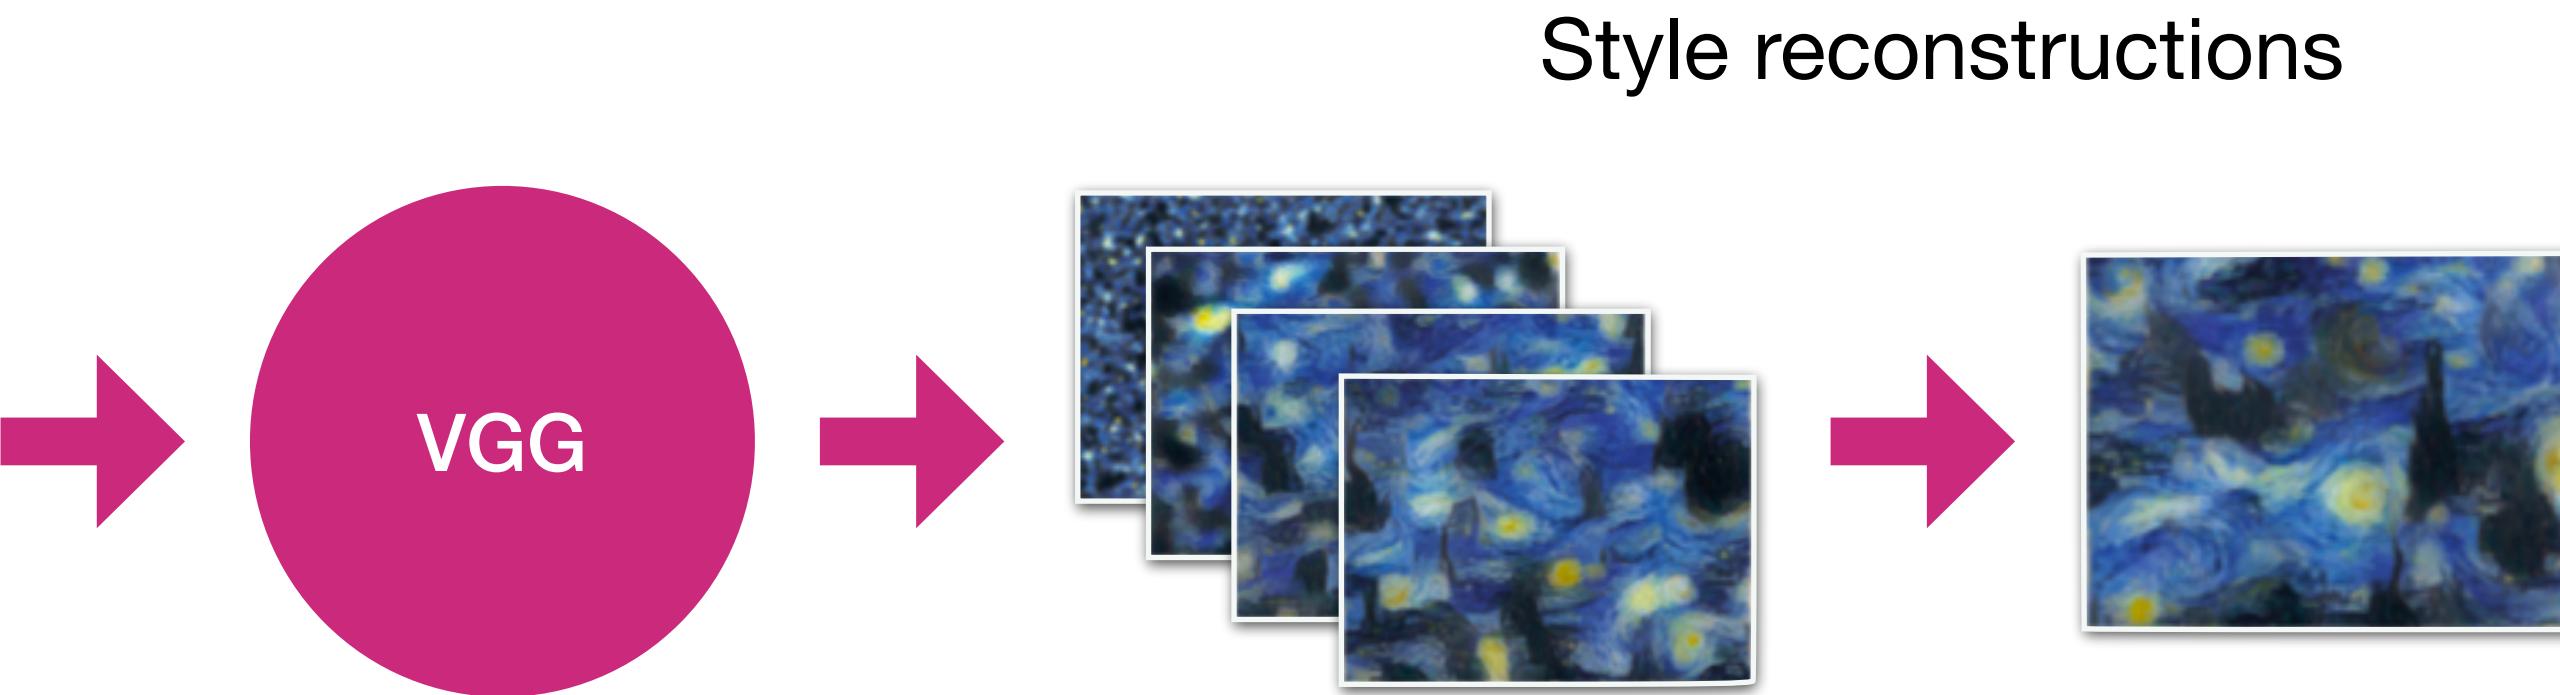




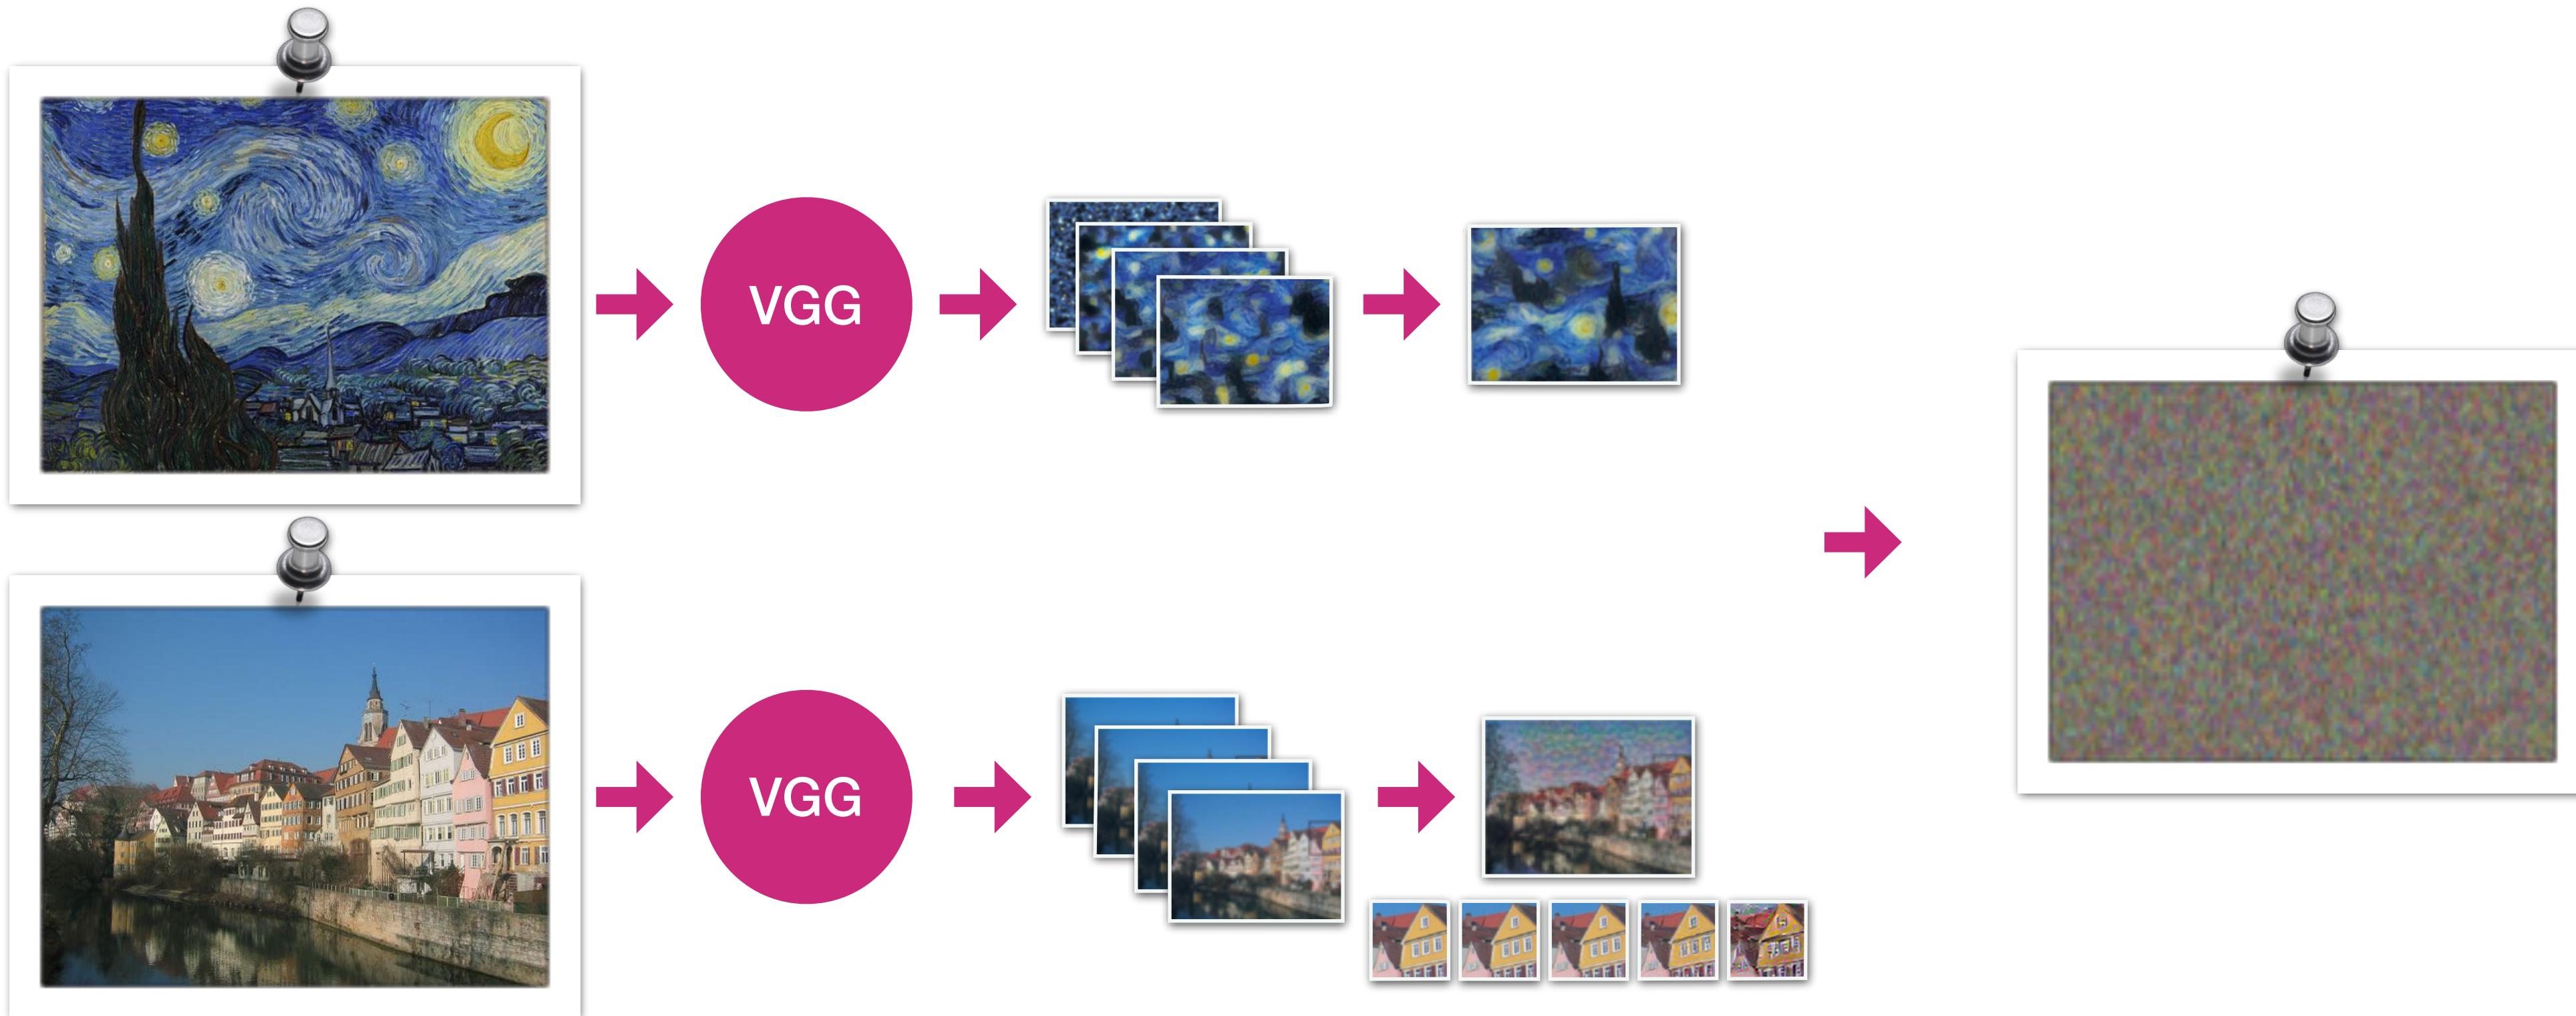


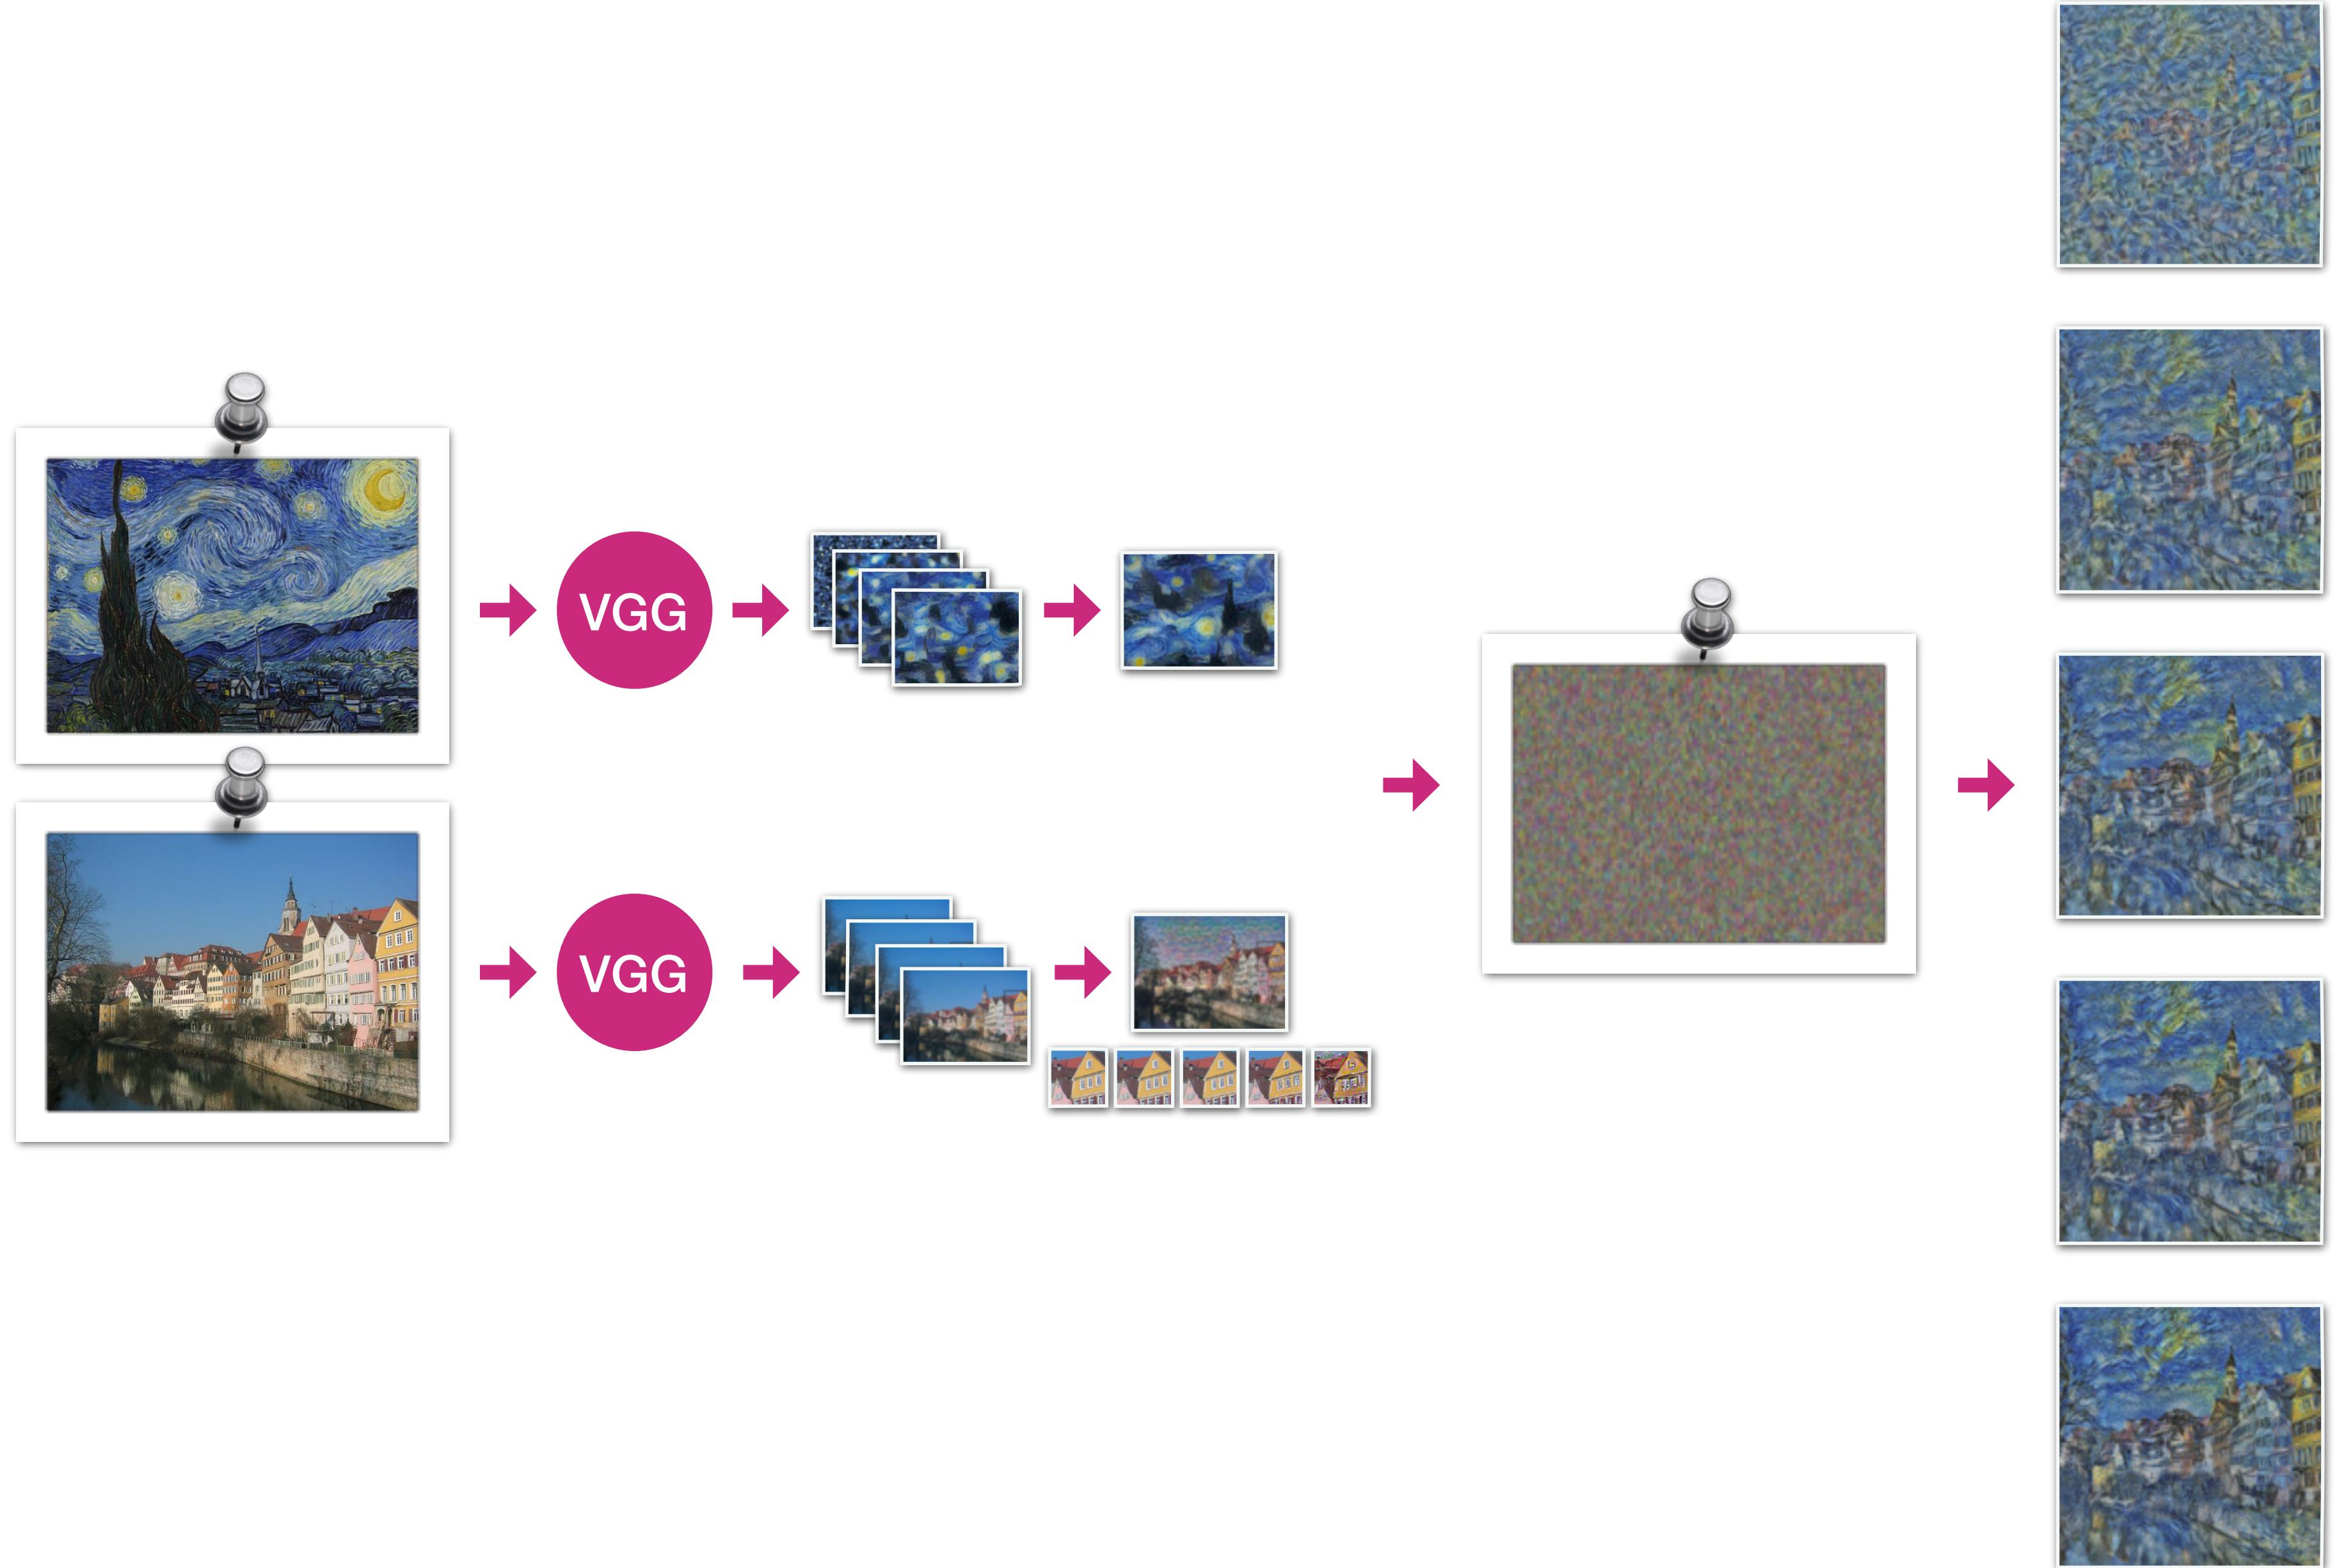


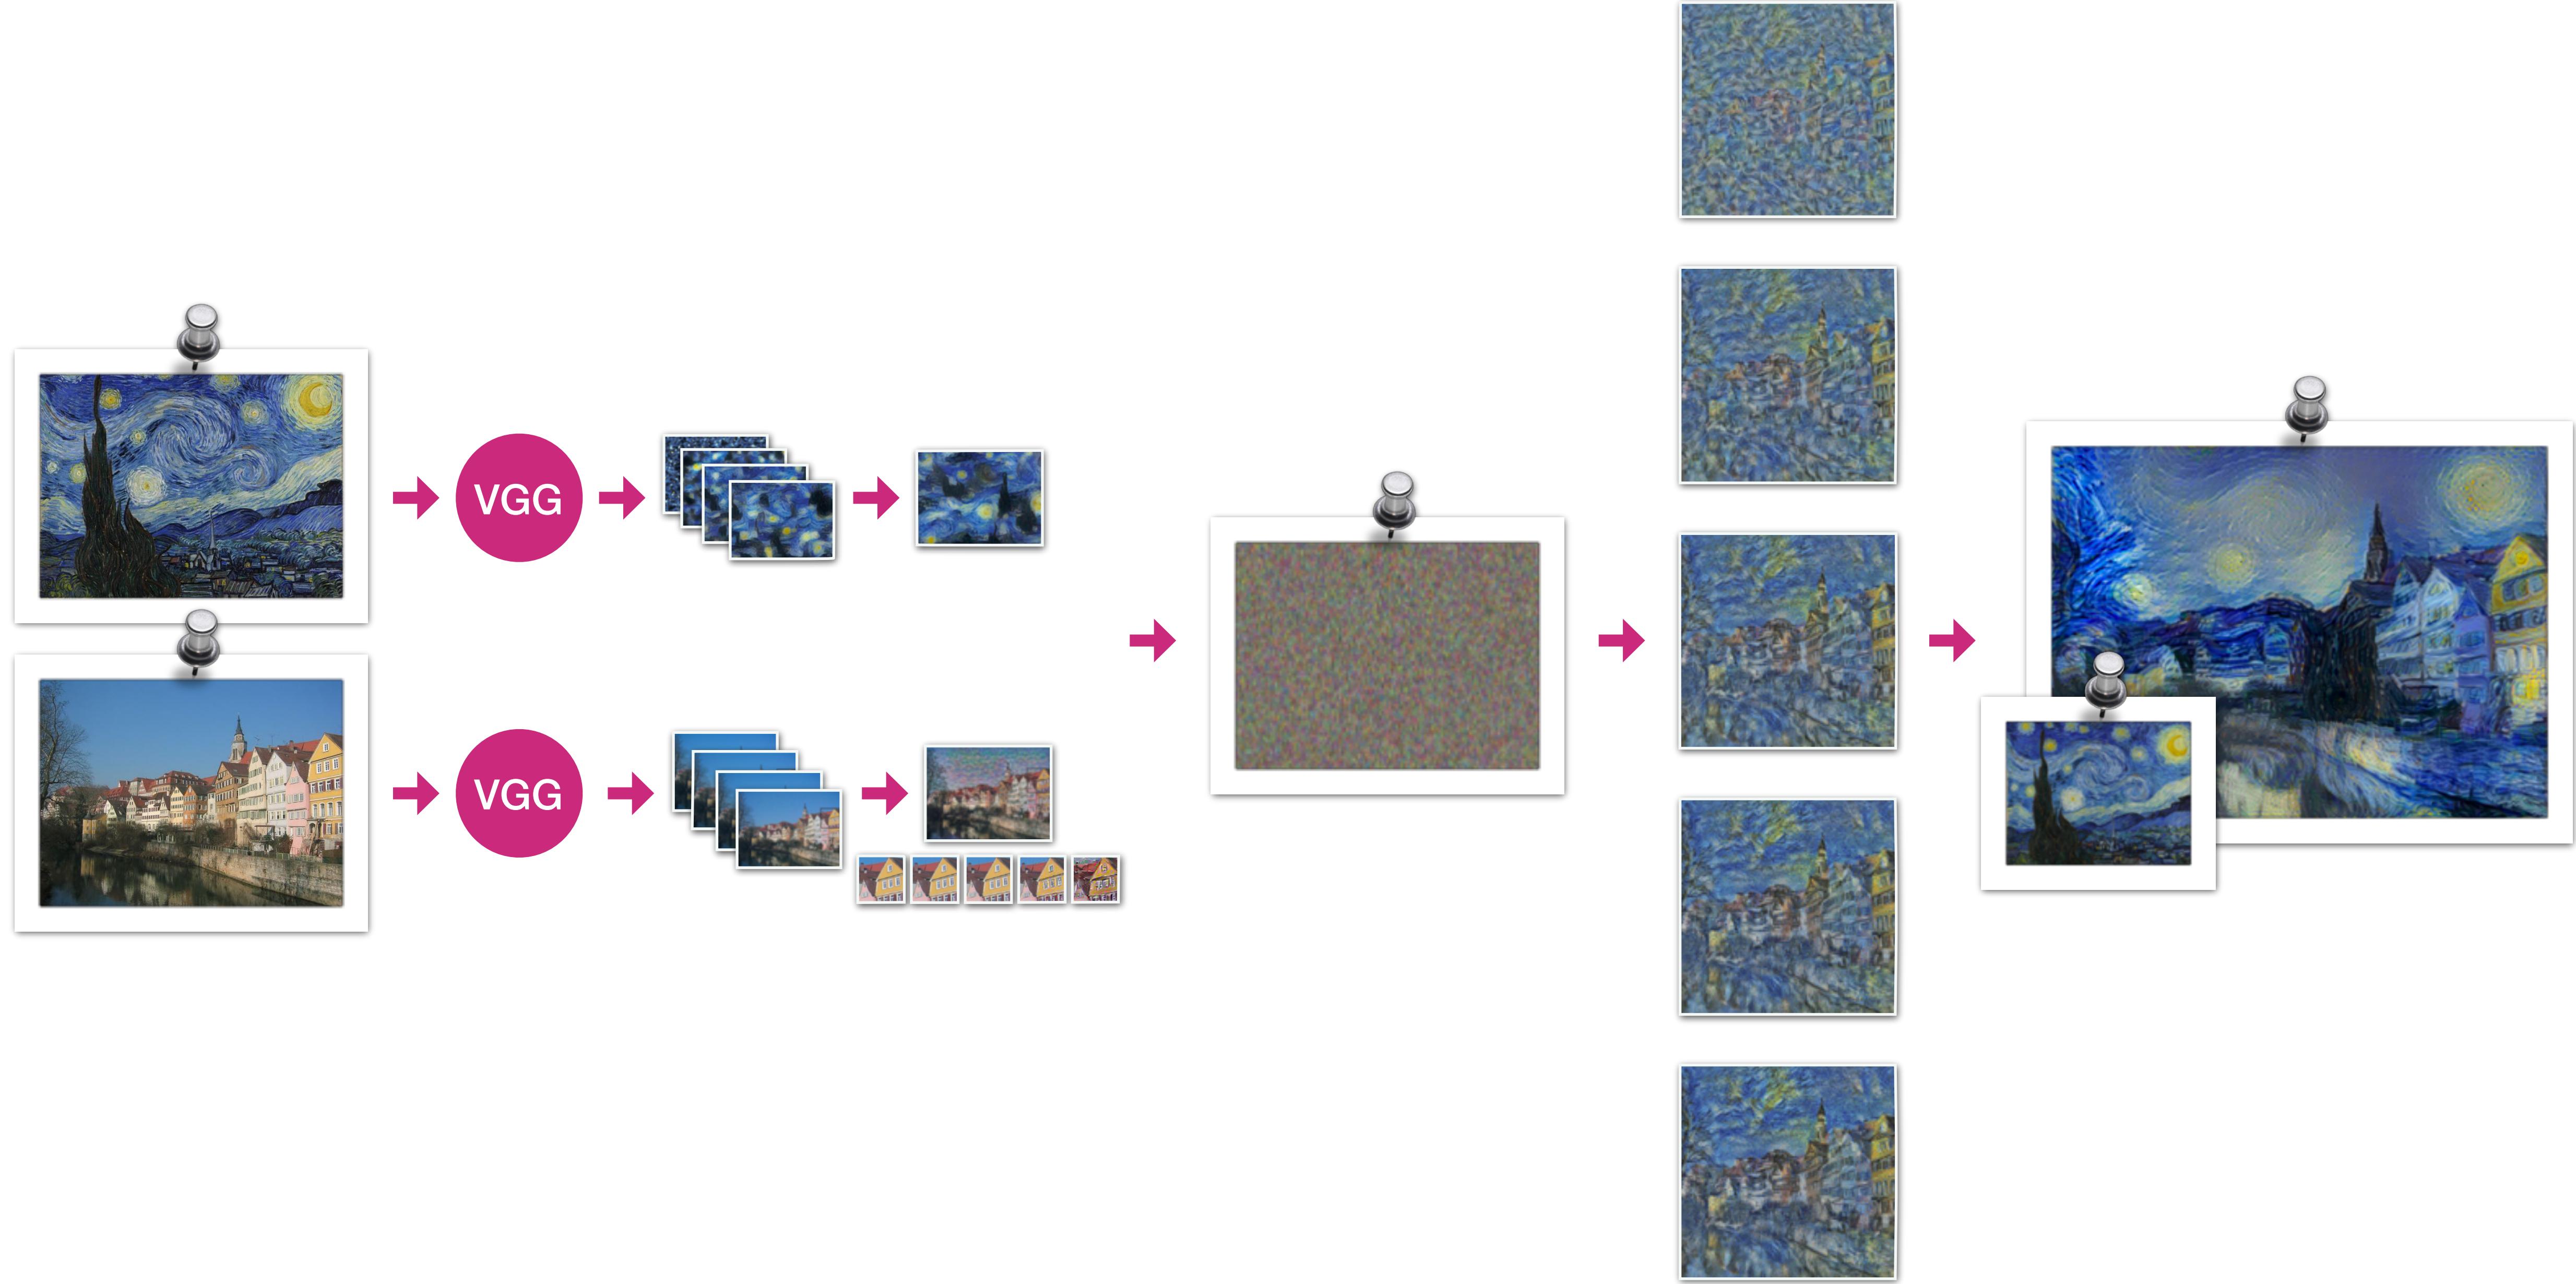












# A Neural Algorithm of Artistic Style

Leon A. Gatys,<sup>1,2,3\*</sup> Alexander S. Ecker,<sup>1,2,4,5</sup> Matthias Bethge<sup>1,2,4</sup>

<sup>1</sup>Werner Reichardt Centre for Integrative Neuroscience  
and Institute of Theoretical Physics, University of Tübingen, Germany

<sup>2</sup>Bernstein Center for Computational Neuroscience, Tübingen, Germany

<sup>3</sup>Graduate School for Neural Information Processing, Tübingen, Germany

<sup>4</sup>Max Planck Institute for Biological Cybernetics, Tübingen, Germany

<sup>5</sup>Department of Neuroscience, Baylor College of Medicine, Houston, TX, USA

\*To whom correspondence should be addressed; E-mail: leon.gatys@bethgelab.org

In fine art, especially painting, humans have mastered the skill to create unique visual experiences through composing a complex interplay between the content and style of an image. Thus far the algorithmic basis of this process is unknown and there exists no artificial system with similar capabilities. However, in other key areas of visual perception such as object and face recognition near-human performance was recently demonstrated by a class of biologically inspired vision models called Deep Neural Networks.<sup>1,2</sup> Here we introduce an artificial system based on a Deep Neural Network that creates artistic images of high perceptual quality. The system uses neural representations to separate and recombine content and style of arbitrary images, providing a neural algorithm for the creation of artistic images. Moreover, in light of the striking similarities between performance-optimised artificial neural networks and biological vision,<sup>3–7</sup> our work offers a path forward to an algorithmic understanding of how humans create and perceive artistic imagery.

# So what?

That's nice, but it makes sense. Why does this talk have 'on-device' in the title?  
What are these Australians ranting about?

# Quick and efficient

If it's going mobile, if it's going on-device, it needs to be better.

# The Old Way

Using a neural network... but not really training.

# An Old Way

- Trained classifier as a feature detector: texture synthesis
  - e.g. using a network to extract features from a texture image.
- Checking how close the synthesised image is to a content image with respect to the features extracted by the classifier
- Repeat.
- Feedforward network, going from content to new styled image in one pass, but only works on a single style.

# It's always the same problem...

As it turns out!

# A Neural Algorithm of Artistic Style

Leon A. Gatys,<sup>1,2,3\*</sup> Alexander S. Ecker,<sup>1,2,4,5</sup> Matthias Bethge<sup>1,2,4</sup>

<sup>1</sup>Werner Reichardt Centre for Integrative Neuroscience  
and Institute of Theoretical Physics, University of Tübingen, Germany

<sup>2</sup>Bernstein Center for Computational Neuroscience, Tübingen, Germany

<sup>3</sup>Graduate School for Neural Information Processing, Tübingen, Germany

<sup>4</sup>Max Planck Institute for Biological Cybernetics, Tübingen, Germany

<sup>5</sup>Department of Neuroscience, Baylor College of Medicine, Houston, TX, USA

\*To whom correspondence should be addressed; E-mail: leon.gatys@bethgelab.org

In fine art, especially painting, humans have mastered the skill to create unique visual experiences through composing a complex interplay between the content and style of an image. Thus far the algorithmic basis of this process is unknown and there exists no artificial system with similar capabilities. However, in other key areas of visual perception such as object and face recognition near-human performance was recently demonstrated by a class of biologically inspired vision models called Deep Neural Networks.<sup>1,2</sup> Here we introduce an artificial system based on a Deep Neural Network that creates artistic images of high perceptual quality. The system uses neural representations to separate and recombine content and style of arbitrary images, providing a neural algorithm for the creation of artistic images. Moreover, in light of the striking similarities between performance-optimised artificial neural networks and biological vision,<sup>3–7</sup> our work offers a path forward to an algorithmic understanding of how humans create and perceive artistic imagery.

# A Neural Algorithm of Artistic Style

Leon A. Gatys,<sup>1,2,3\*</sup> Alexander S. Ecker,<sup>1,2,4,5</sup> Matthias Bethge<sup>1,2</sup>

<sup>1</sup>Werner Reichardt Centre for Integrative Neuroscience  
and Institute of Theoretical Physics, University of Tübingen, Germany

<sup>2</sup>Bernstein Center for Computational Neuroscience, Tübingen, Germany

<sup>3</sup>Graduate School for Neural Information Processing, Tübingen, Germany

<sup>4</sup>Max Planck Institute for Biological Cybernetics, Tübingen, Germany

<sup>5</sup>Department of Neuroscience, Baylor College of Medicine, Houston, TX, USA

\*To whom correspondence should be addressed; E-mail: leon.gatys@bethge

In fine art, especially painting, humans have mastered the skill to create visual experiences through composing a complex interplay between content and style of an image. Thus far the algorithmic basis of this process is unknown and there exists no artificial system with similar capabilities. However, in other key areas of visual perception such as object and face recognition, near-human performance was recently demonstrated by a class of biologically inspired vision models called Deep Neural Networks.<sup>1,2</sup> Here we introduce an artificial system based on a Deep Neural Network that creates art of high perceptual quality. The system uses neural representations to separate and recombine content and style of arbitrary images, providing a general algorithm for the creation of artistic images. Moreover, in light of the striking similarities between performance-optimised artificial neural networks and biological vision,<sup>3–7</sup> our work offers a path forward to an algorithmic understanding of how humans create and perceive artistic imagery.

## A LEARNED REPRESENTATION FOR ARTISTIC STYLE

Vincent Dumoulin & Jonathon Shlens & Manjunath Koduri  
Google Brain, Mountain View, CA  
vi.dumoulin@gmail.com, shlens@google.com, keveman@google.com

### ABSTRACT

The diversity of painting styles represents a rich visual vocabulary for the construction of an image. The degree to which one may learn and parsimoniously capture this visual vocabulary measures our understanding of the higher level features of paintings, if not images in general. In this work we investigate the construction of a single, scalable deep network that can parsimoniously capture the artistic style of a diversity of paintings. We demonstrate that such a network generalizes across a diversity of artistic styles by reducing a painting to a point in an embedding space. Importantly, this model permits a user to explore new painting styles by arbitrarily combining the styles learned from individual paintings. We hope that this work provides a useful step towards building rich models of paintings and offers a window on to the structure of the learned representation of artistic style.

### 1 INTRODUCTION

A *pastiche* is an artistic work that imitates the style of another one. Computer vision and more recently machine learning have a history of trying to automate pastiche, that is, render an image in the style of another one. This task is called *style transfer*, and is closely related to the texture synthesis task. While the latter tries to capture the statistical relationship between the pixels of a source image which is assumed to have a stationary distribution at some scale, the former does so while also attempting to preserve some notion of content.

On the computer vision side, Efros & Leung (1999) and Wei & Levoy (2000) attempt to “grow” textures one pixel at a time using non-parametric sampling of pixels in an exemplar image. Efros & Freeman (2001) and Liang et al. (2001) extend this idea to “growing” textures one patch at a time, and Efros & Freeman (2001) uses the approach to implement “texture transfer”, i.e. transferring the texture of an object onto another one. Kwatra et al. (2005) approaches the texture synthesis problem from an energy minimization perspective, progressively refining the texture using an EM-like algorithm. Hertzmann et al. (2001) introduces the concept of “image analogies”: given a pair of “unfiltered” and “filtered” versions of an exemplar image, a target image is processed to create an analogous “filtered” result. More recently, Frigo et al. (2016) treats style transfer as a local texture transfer (using an adaptive patch partition) followed by a global color transfer, and Milasfar (2016) extends Kwatra’s energy-based method into a style transfer algorithm by taking content similarity into account.

On the machine learning side, it has been shown that a trained classifier can be used as a feature extractor to drive texture synthesis and style transfer. Gatys et al. (2015a) uses the VGG-19 network (Simonyan & Zisserman, 2014) to extract features from a texture image and a synthesized texture. The two sets of features are compared and the synthesized texture is modified by gradient descent so that the two sets of features are as close as possible. Gatys et al. (2015b) extends this idea to style transfer by adding the constraint that the synthesized image also be close to a content image with respect to another set of features extracted by the trained VGG-19 classifier.

While very flexible, this algorithm is expensive to run due to the optimization loop being carried. Ulyanov et al. (2016a), Li & Wand (2016) and Johnson et al. (2016) tackle this problem by introducing a *feedforward style transfer network*, which is trained to go from content to pastiche image in one pass. However, in doing so some of the flexibility of the original algorithm is lost: the style transfer network is tied to a single style, which means that separate networks have to be trained

## A LEARNED REPRESENTATION FOR ARTISTIC STYLE

Vincent Dumoulin & Jonathon Shlens & Manjunath Kudlur

Google Brain, Mountain View, CA

vi.dumoulin@gmail.com, shlens@google.com, keveman@google.com

### ABSTRACT

The diversity of painting styles represents a rich visual vocabulary for the construction of an image. The degree to which one may learn and parsimoniously capture this visual vocabulary measures our understanding of the higher level features of paintings, if not images in general. In this work we investigate the construction of a single, scalable deep network that can parsimoniously capture the artistic style of a diversity of paintings. We demonstrate that such a network generalizes across a diversity of artistic styles by reducing a painting to a point in an embedding space. Importantly, this model permits a user to explore new painting styles by arbitrarily combining the styles learned from individual paintings. We hope that this work provides a useful step towards building rich models of paintings and offers a window on to the structure of the learned representation of artistic style.

### 1 INTRODUCTION

A *pastiche* is an artistic work that imitates the style of another one. Computer vision and more recently machine learning have a history of trying to automate *pastiche*, that is, render an image in the style of another one. This task is called *style transfer*, and is closely related to the texture synthesis task. While the latter tries to capture the statistical relationship between the pixels of a source image which is assumed to have a stationary distribution at some scale, the former does so while also attempting to preserve some notion of content.

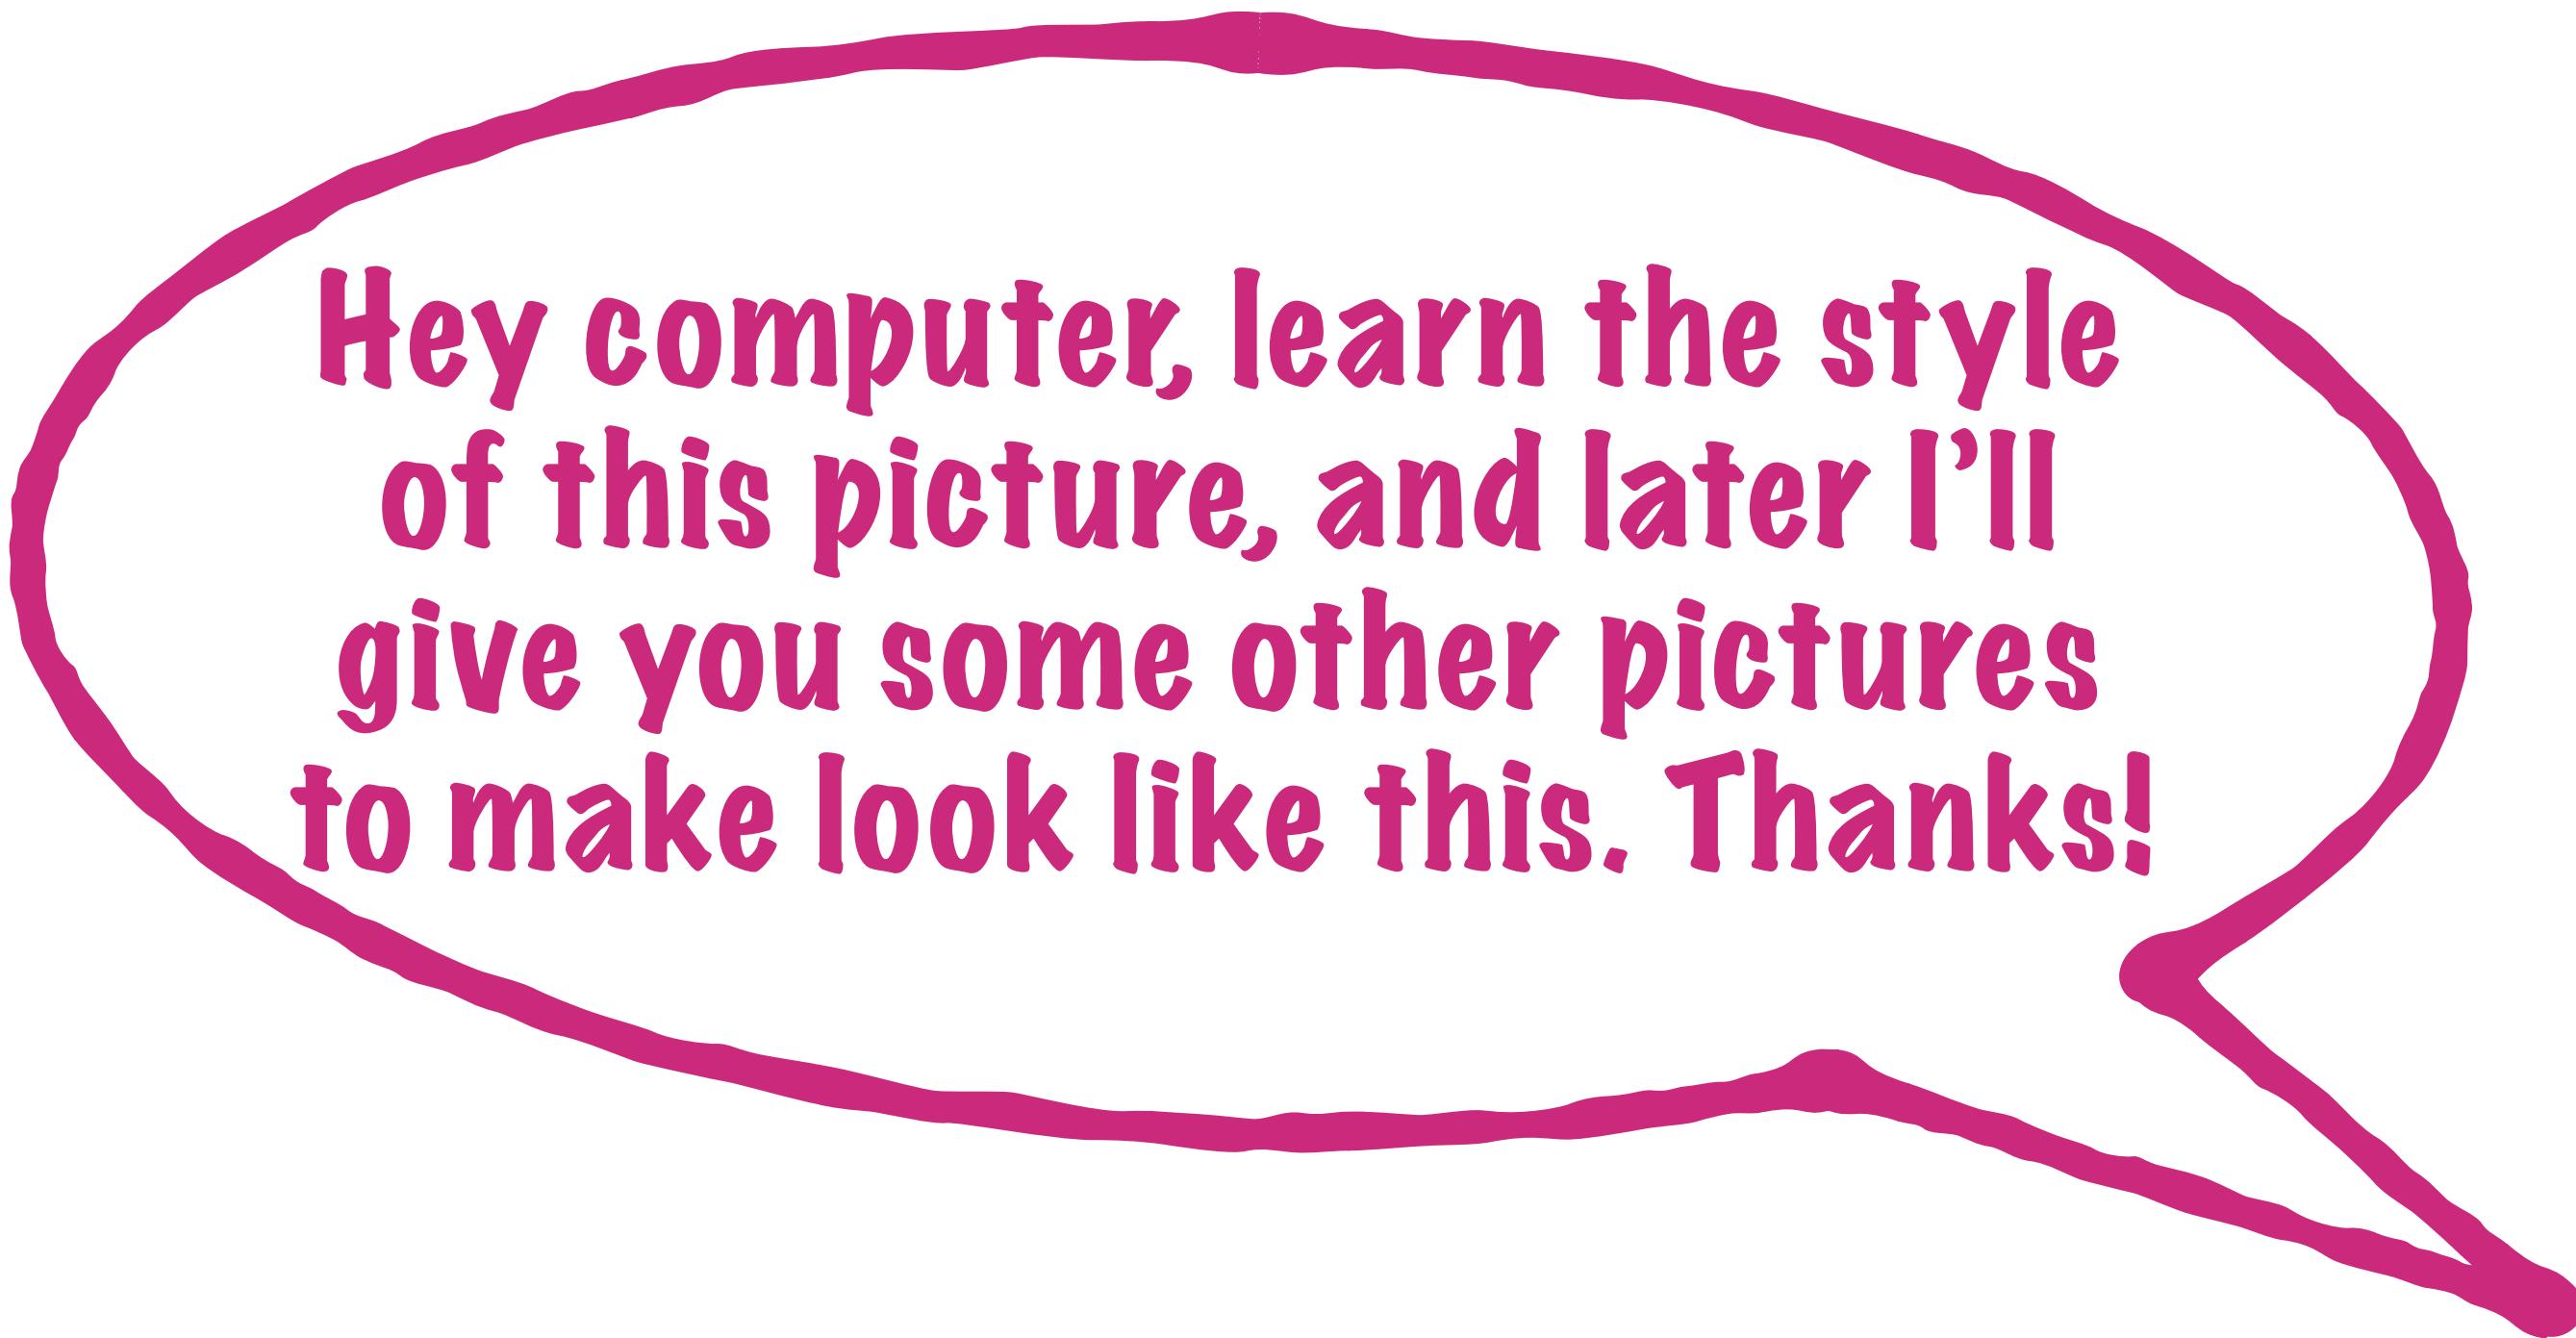
On the computer vision side, Efros & Leung (1999) and Wei & Levoy (2000) attempt to “grow” textures one pixel at a time using non-parametric sampling of pixels in an exemplar image. Efros & Freeman (2001) and Liang et al. (2001) extend this idea to “growing” textures one patch at a time, and Efros & Freeman (2001) uses the approach to implement “texture transfer”, i.e. transferring the texture of an object onto another one. Kwatra et al. (2005) approaches the texture synthesis problem from an energy minimization perspective, progressively refining the texture using an EM-like algorithm. Hertzmann et al. (2001) introduces the concept of “image analogies”: given a pair of “unfiltered” and “filtered” versions of an exemplar image, a target image is processed to create an analogous “filtered” result. More recently, Frigo et al. (2016) treats style transfer as a local texture transfer (using an adaptive patch partition) followed by a global color transfer, and Elad & Milanfar (2016) extends Kwatra’s energy-based method into a style transfer algorithm by taking content similarity into account.

On the machine learning side, it has been shown that a trained classifier can be used as a feature extractor to drive texture synthesis and style transfer. Gatys et al. (2015a) uses the VGG-19 network (Simonyan & Zisserman, 2014) to extract features from a texture image and a synthesized texture. The two sets of features are compared and the synthesized texture is modified by gradient descent so that the two sets of features are as close as possible. Gatys et al. (2015b) extends this idea to style transfer by adding the constraint that the synthesized image also be close to a content image with respect to another set of features extracted by the trained VGG-19 classifier.

While very flexible, this algorithm is expensive to run due to the optimization loop being carried. Ulyanov et al. (2016a), Li & Wand (2016) and Johnson et al. (2016) tackle this problem by introducing a *feedforward style transfer network*, which is trained to go from content to *pastiche* image in one pass. However, in doing so some of the flexibility of the original algorithm is lost: the style transfer network is tied to a single style, which means that separate networks have to be trained

# The New Way

# The New Way



Hey computer, learn the style  
of this picture, and later I'll  
give you some other pictures  
to make look like this. Thanks!

# The New Image

Activations in the neural network that are similar  
to both the style and to the content activations.

# A million approaches

- Turi Create + Core ML + Swift iOS app
- Keras + Core ML Tools + Core ML + Swift iOS app
- PyTorch + Core ML Tools + Core ML + Swift iOS app
- TensorFlow + Core ML Tools + Core ML + Swift iOS app
- .... and many others!



*Part 1*

# On-device Machine Learning

*Part 2*

# Neural Style Transfer

*Part 3*

# Demonstration

*Part 4*

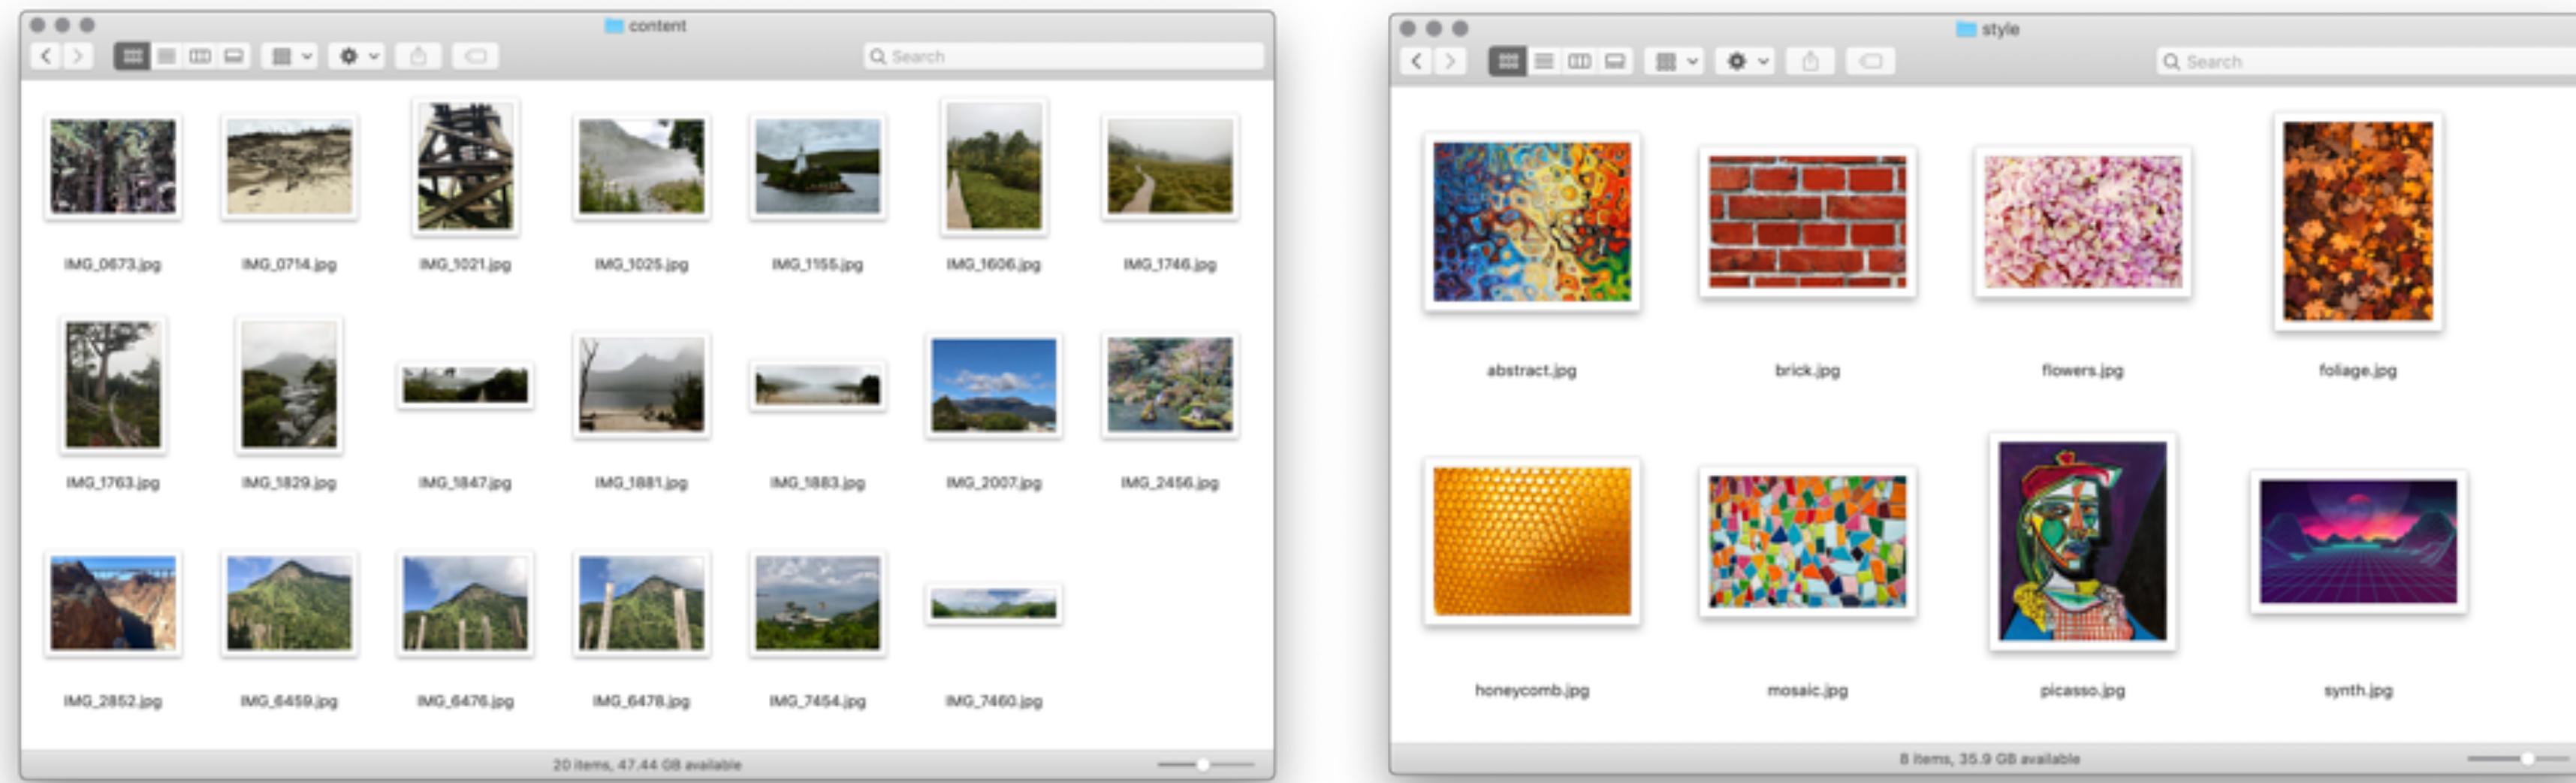
# The Future

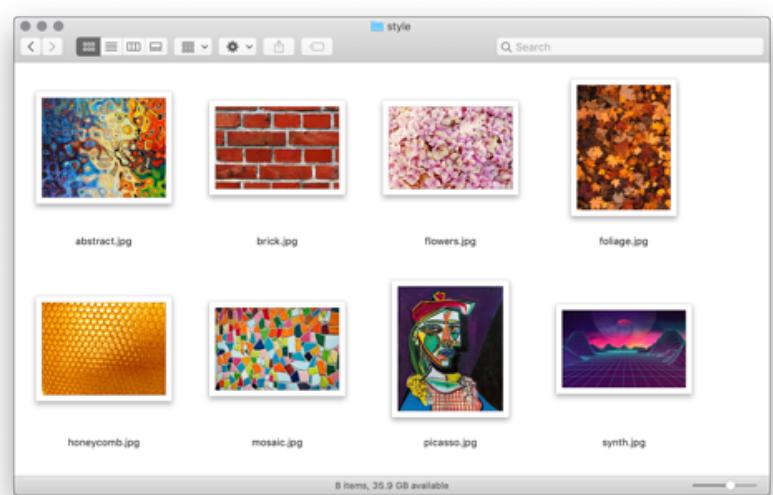
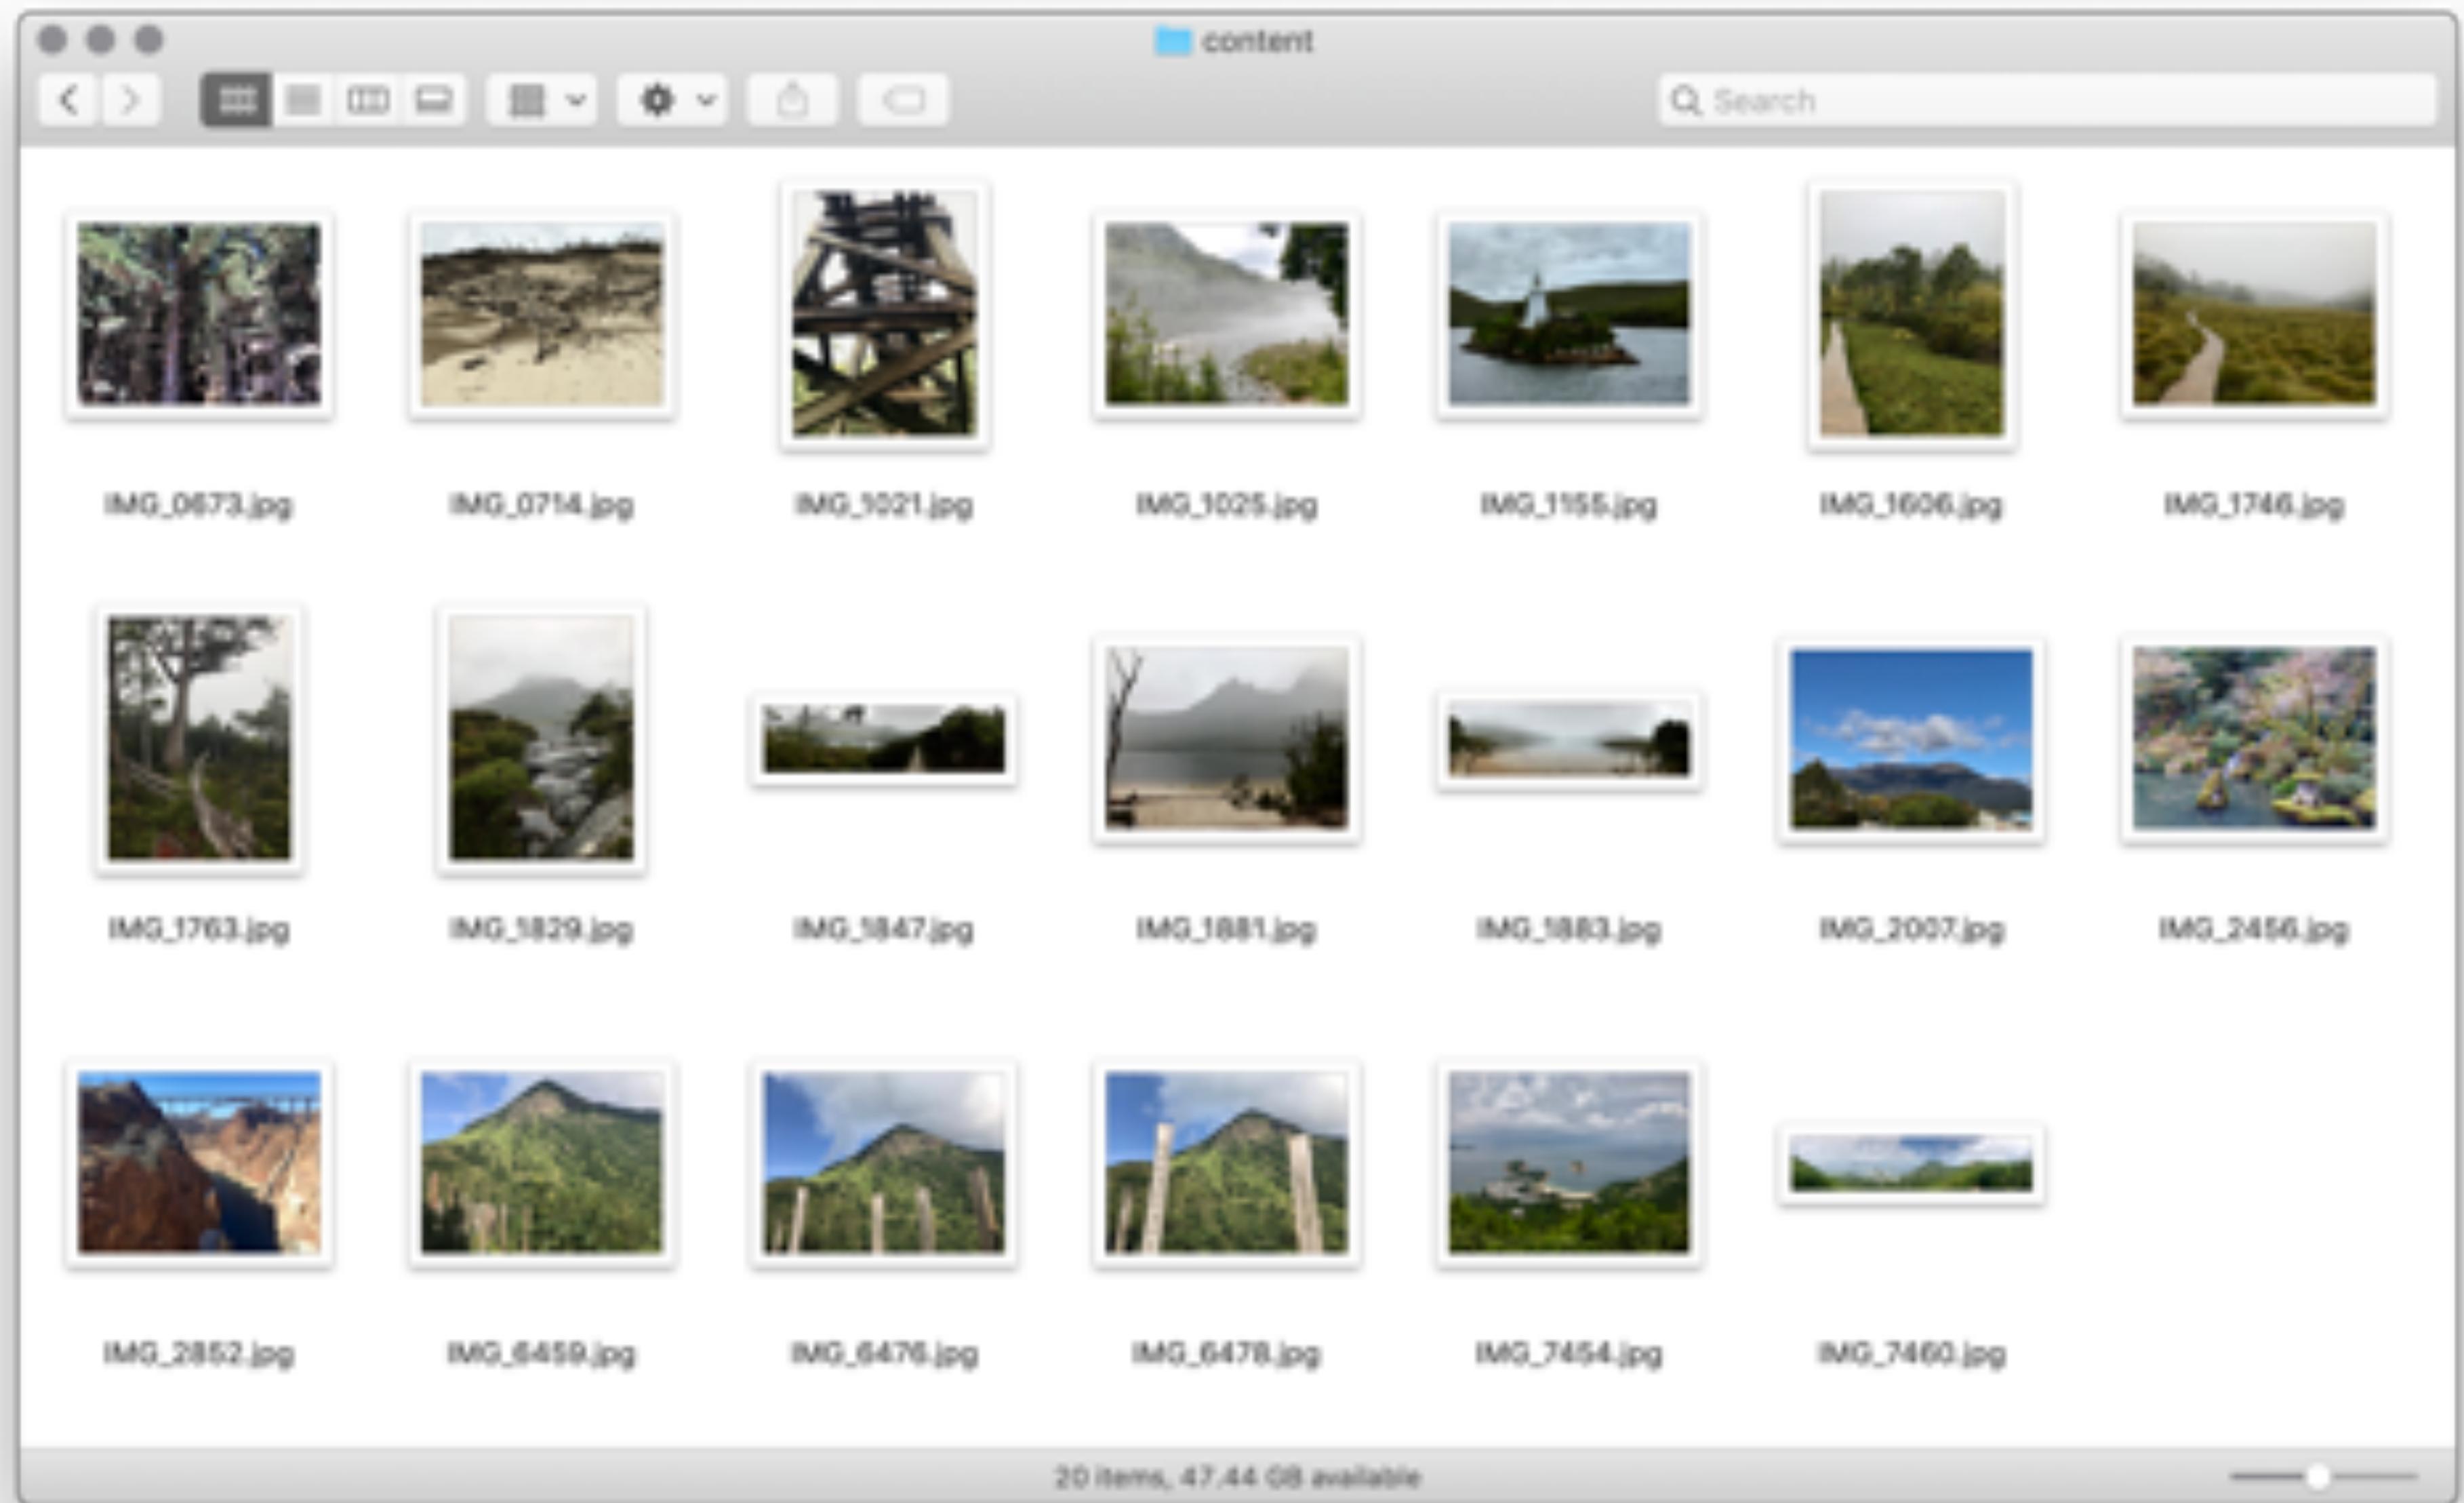
*Part 3*

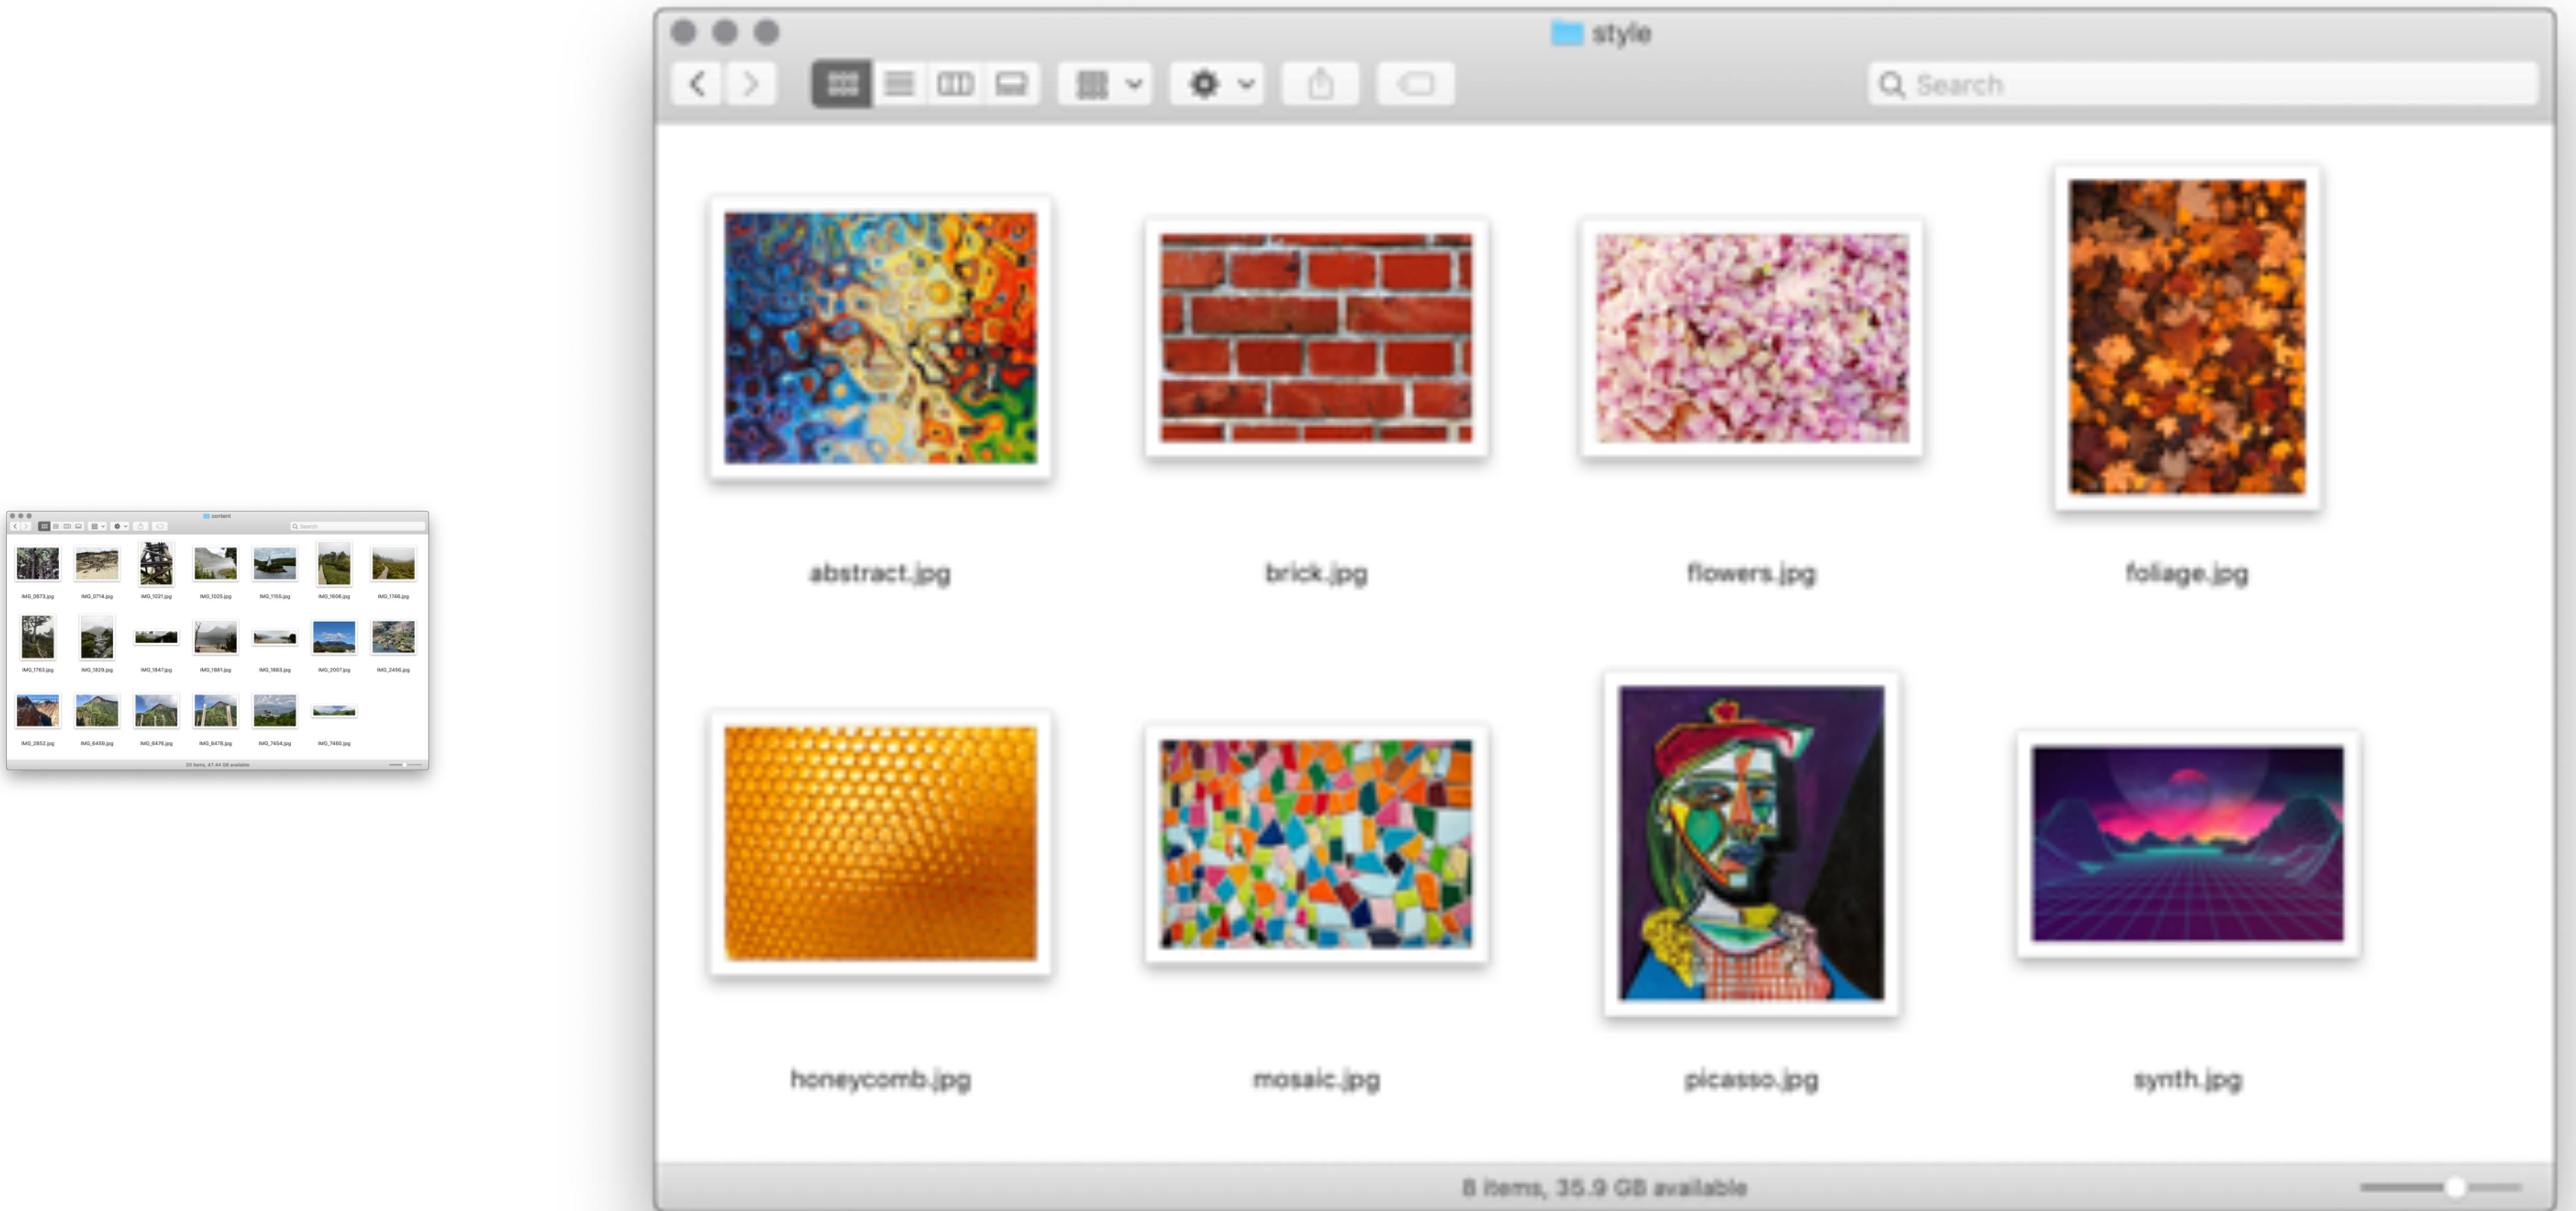
# Demonstration

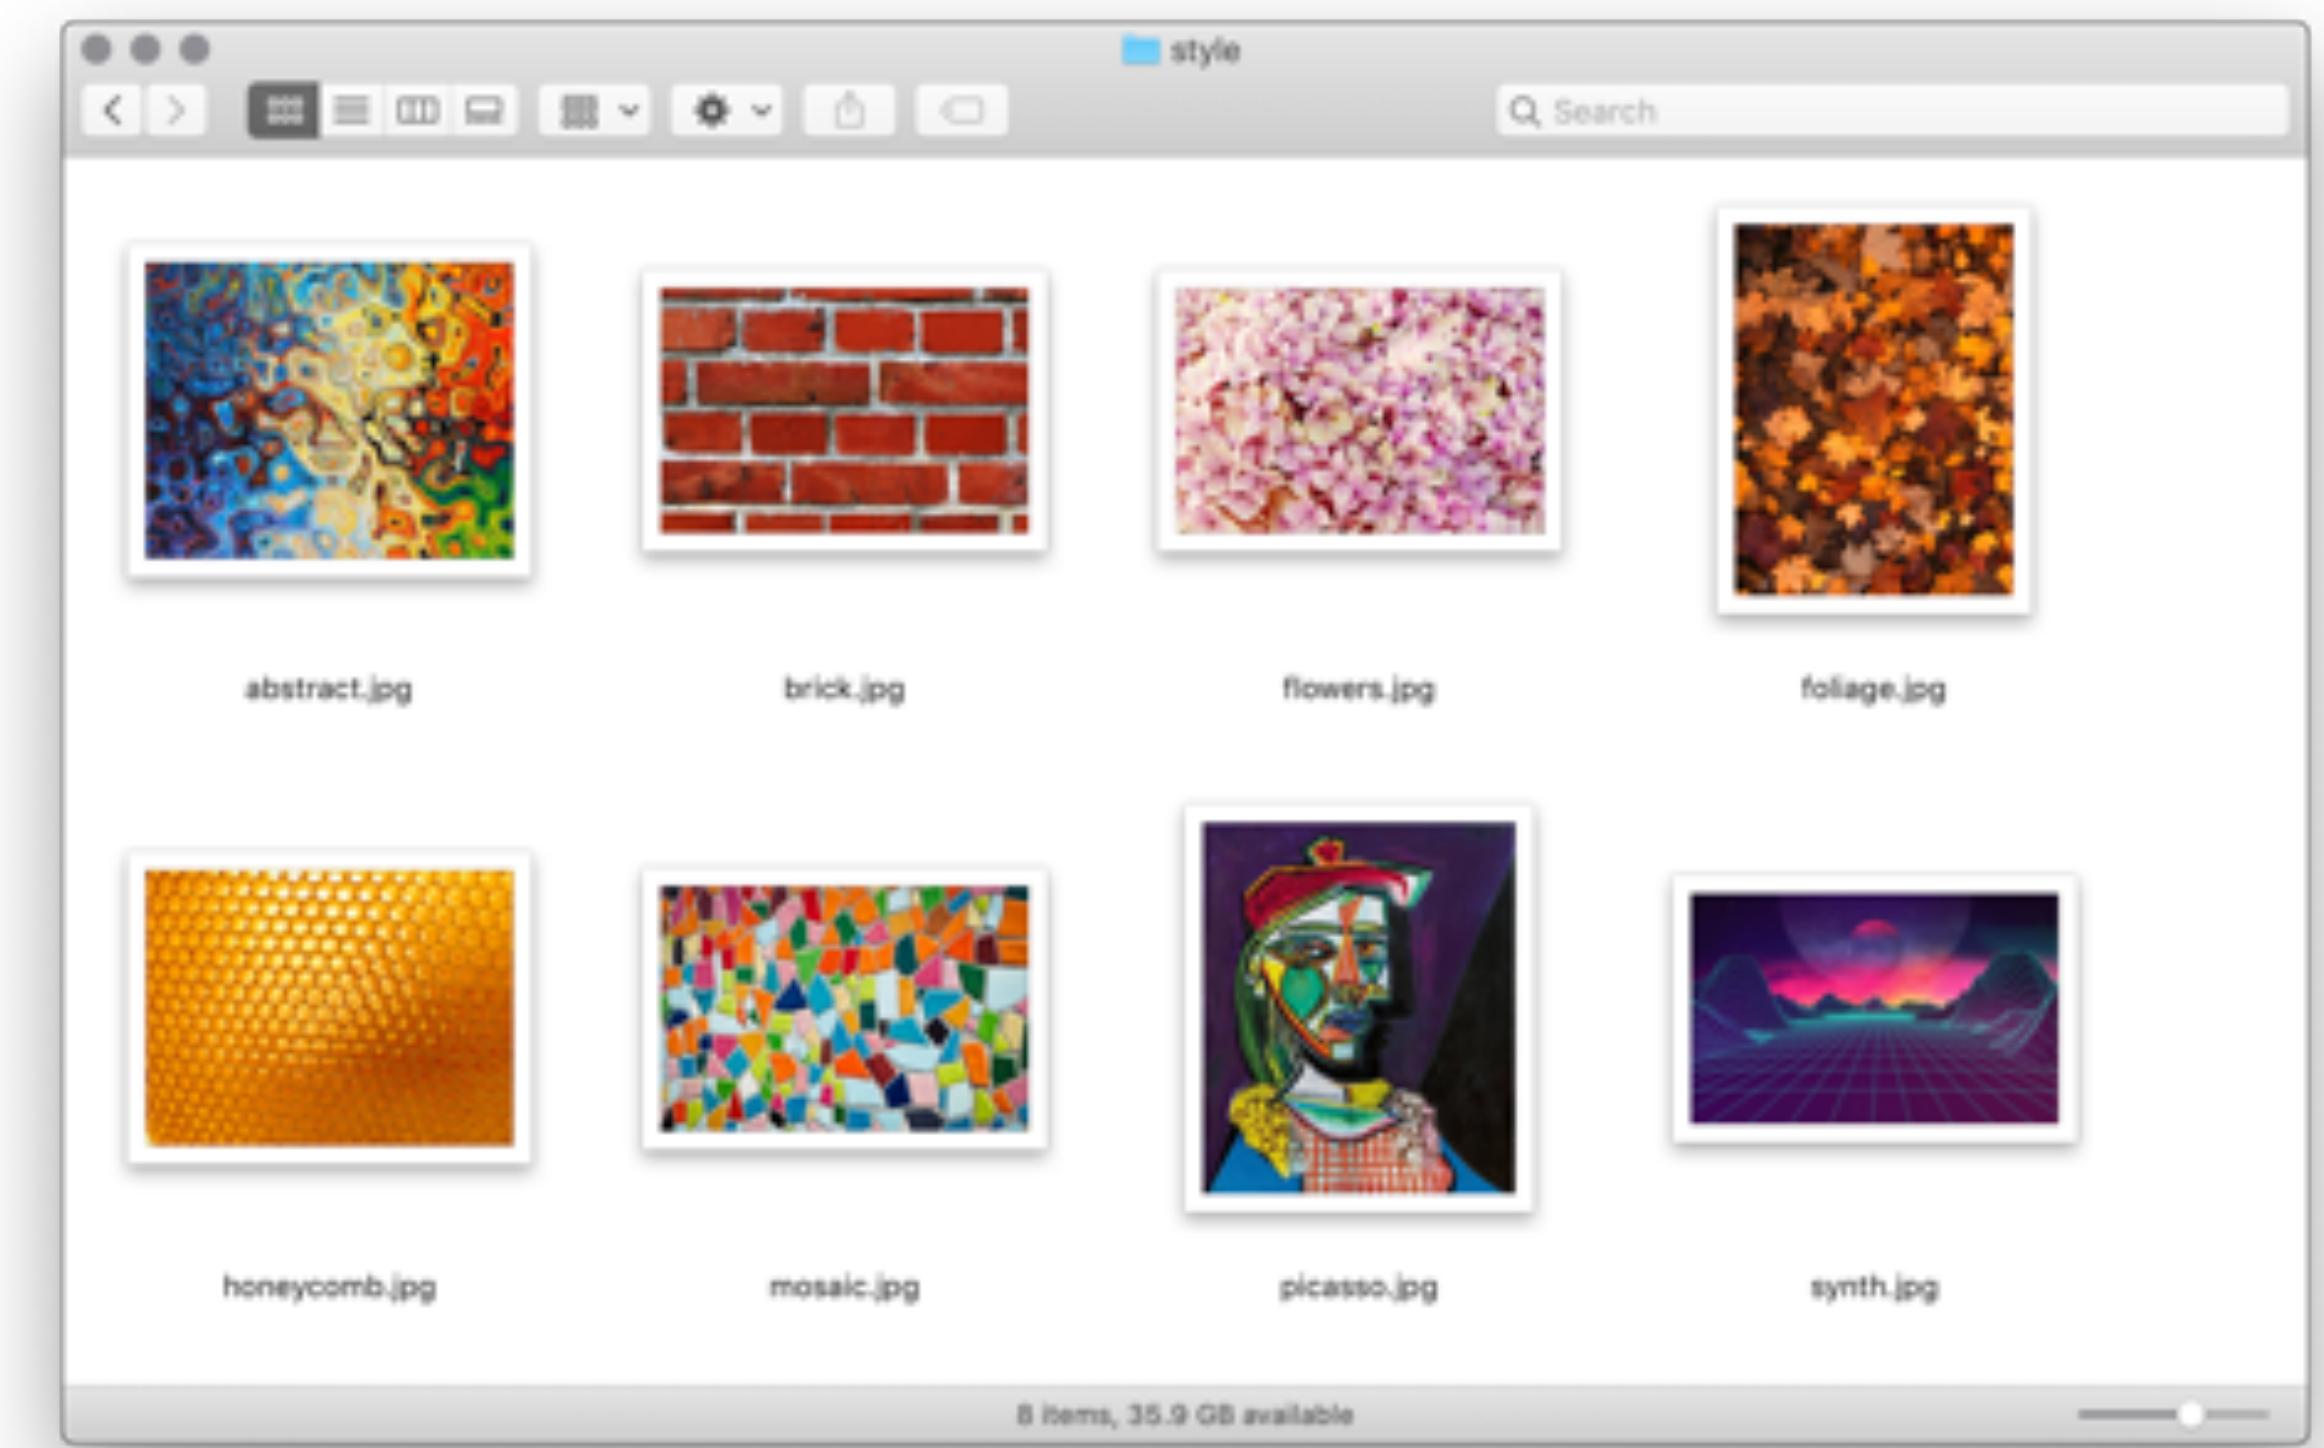
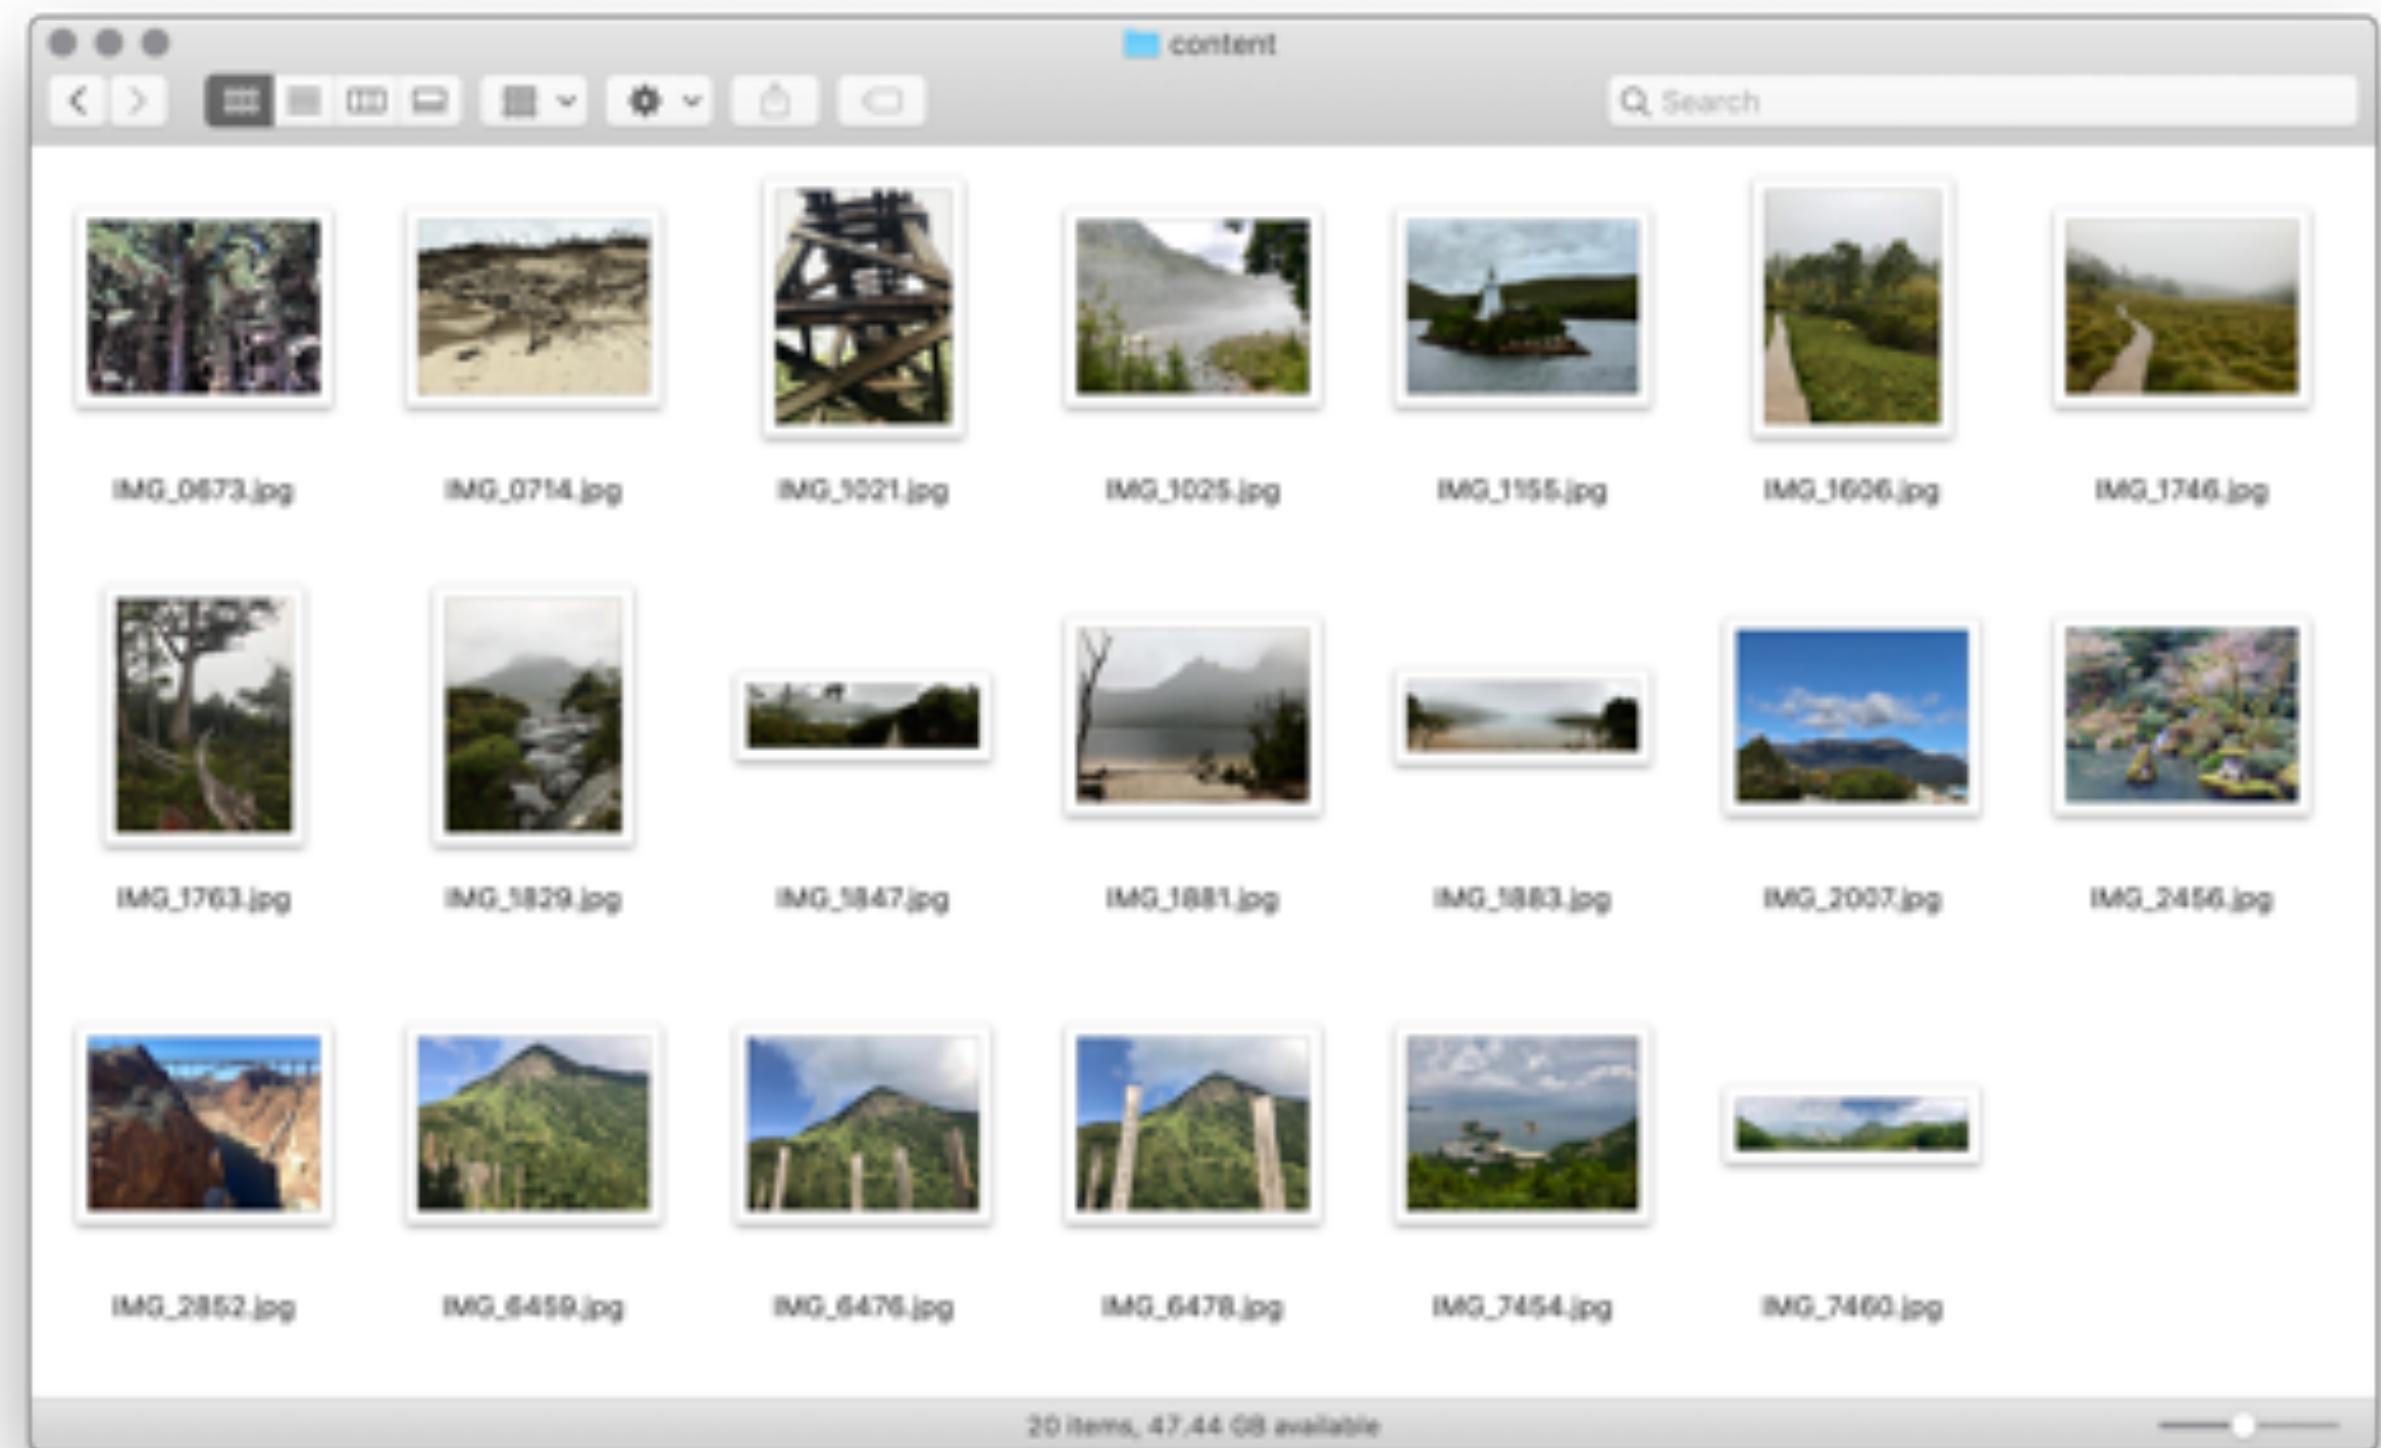
# Data Preparation

Step 1









# Training

## Step 2

**Turi Create**

ML Task	Description
Recommender	Personalize choices for users
Image Classification	Label images
Object Detection	Recognize objects within images
Style Transfer	Stylize images
Activity Classification	Detect an activity using sensors
Image Similarity	Find similar images
Classifiers	Predict a label
Regression	Predict numeric values
Clustering	Group similar datapoints together
Text Classifier	Analyze sentiment of messages

```
git clone https://github.com/apple/turicreate.git
cd turicreate
# build
cd examples/activity_recognition
# run
./activity_recognition.py
```

```
import turicreate as tc

# Load the style and content images
styles = tc.load_images('style/')
content = tc.load_images('content/')

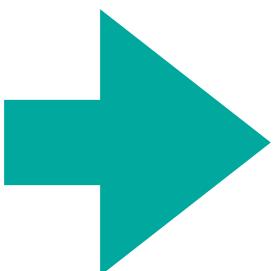
# Create a StyleTransfer model
model = tc.style_transfer.create(styles, content, max_iterations=2000)

# Load some test images
test_images = tc.load_images('test/')

# Stylize the test images
stylized_images = model.stylize(test_images)

# Save the model for later use in Turi Create
model.save('DemoNST.model')

# Export for use in Core ML
model.export_coreml('DemoNST.mlmodel')
```



```
import turicreate as tc

# Load the style and content images
styles = tc.load_images('style/')
content = tc.load_images('content/')

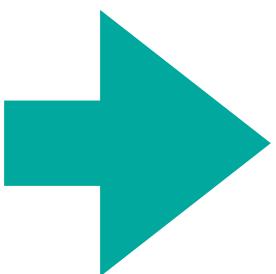
# Create a StyleTransfer model
model = tc.style_transfer.create(styles, content, max_iterations=2000)

# Load some test images
test_images = tc.load_images('test/')

# Stylize the test images
stylized_images = model.stylize(test_images)

# Save the model for later use in Turi Create
model.save('DemoNST.model')

# Export for use in Core ML
model.export_coreml('DemoNST.mlmodel')
```



```
import turicreate as tc

# Load the style and content images
styles = tc.load_images('style/')
content = tc.load_images('content/')

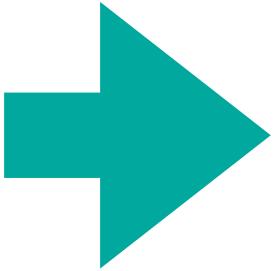
# Create a StyleTransfer model
model = tc.style_transfer.create(styles, content, max_iterations=2000)

# Load some test images
test_images = tc.load_images('test/')

# Stylize the test images
stylized_images = model.stylize(test_images)

# Save the model for later use in Turi Create
model.save('DemoNST.model')

# Export for use in Core ML
model.export_coreml('DemoNST.mlmodel')
```



```
import turicreate as tc

# Load the style and content images
styles = tc.load_images('style/')
content = tc.load_images('content/')

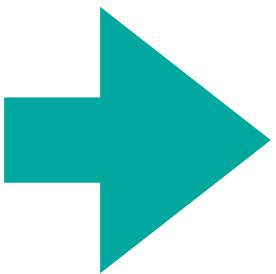
# Create a StyleTransfer model
model = tc.style_transfer.create(styles, content, max_iterations=2000)

# Load some test images
test_images = tc.load_images('test/')

# Stylize the test images
stylized_images = model.stylize(test_images)

# Save the model for later use in Turi Create
model.save('DemoNST.model')

# Export for use in Core ML
model.export_coreml('DemoNST.mlmodel')
```



```
import turicreate as tc

# Load the style and content images
styles = tc.load_images('style/')
content = tc.load_images('content/')

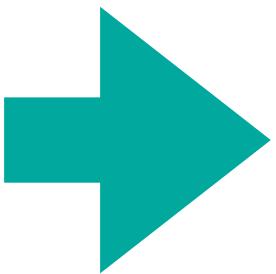
# Create a StyleTransfer model
model = tc.style_transfer.create(styles, content, max_iterations=2000)

# Load some test images
test_images = tc.load_images('test/')

# Stylize the test images
stylized_images = model.stylize(test_images)

# Save the model for later use in Turi Create
model.save('DemoNST.model')

# Export for use in Core ML
model.export_coreml('DemoNST.mlmodel')
```



```
import turicreate as tc

# Load the style and content images
styles = tc.load_images('style/')
content = tc.load_images('content/')

# Create a StyleTransfer model
model = tc.style_transfer.create(styles, content, max_iterations=2000)

# Load some test images
test_images = tc.load_images('test/')

# Stylize the test images
stylized_images = model.stylize(test_images)

# Save the model for later use in Turi Create
model.save('DemoNST.model')

# Export for use in Core ML
model.export_coreml('DemoNST.mlmodel')
```

```
import turicreate as tc

# Load the style and content images
styles = tc.load_images('style/')
content = tc.load_images('content/')

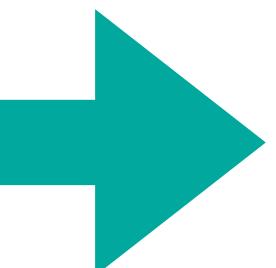
# Create a StyleTransfer model
model = tc.style_transfer.create(styles, content, max_iterations=2000)

# Load some test images
test_images = tc.load_images('test/')

# Stylize the test images
stylized_images = model.stylize(test_images)

# Save the model for later use in Turi Create
model.save('DemoNST.model')

# Export for use in Core ML
model.export_coreml('DemoNST.mlmodel')
```



```
import turicreate as tc

# Load the style and content images
styles = tc.load_images('style/')
content = tc.load_images('content/')

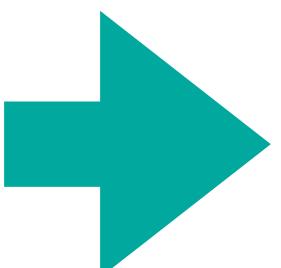
# Create a StyleTransfer model
model = tc.style_transfer.create(styles, content, max_iterations=2000)

# Load some test images
test_images = tc.load_images('test/')

# Stylize the test images
stylized_images = model.stylize(test_images)

# Save the model for later use in Turi Create
model.save('DemoNST.model')

# Export for use in Core ML
model.export_coreml('DemoNST.mlmodel')
```





A screenshot of a terminal window with a light gray background. The title bar at the top reads "turicreate-training — train.py — 80x24". The command line shows "(venv) Hyperion:turicreate-training desplesda\$ ./train.py". The terminal is mostly empty below the command line.

```
NST — python train_nst.py — python — python — python train_nst.py — 119x30
python train_nst.py
Using 'image' in style_dataset as feature column and using 'image' in content_dataset as feature column

Downloading https://docs-assets.developer.apple.com/turicreate/models/resnet-16.params
Download completed: /var/folders/sz/9p_l7vsn4zn0f4051by7yh9m0000gn/T/model_cache/resnet-16.params
Downloading https://docs-assets.developer.apple.com/turicreate/models/vgg16-conv1_1-4_3.params
Download completed: /var/folders/sz/9p_l7vsn4zn0f4051by7yh9m0000gn/T/model_cache/vgg16-conv1_1-4_3.params

| 1 | 28.058 | 7.5
| 2 | 19.391 | 29.7
| 3 | 22.495 | 52.2
| 4 | 25.156 | 74.5
| 5 | 24.388 | 97.4
| 6 | 23.689 | 126.4
| 7 | 23.828 | 153.1
| 8 | 25.688 | 179.0
| 9 | 24.263 | 203.8
| 10 | 22.451 | 230.1
| 11 | 21.359 | 260.5
| 12 | 20.486 | 292.7
| 13 | 19.633 | 319.2
| 14 | 18.834 | 342.8
| 15 | 21.639 | 367.2
| 16 | 20.758 | 392.9
| 17 | 23.566 | 418.9
| 18 | 22.874 | 444.5
```

# Training

VGG-16

ResNet-16

# Training

# Super fast

# Training in the cloud

Many options!

- <https://notebooks.azure.com>
- <https://colab.research.google.com>
- <https://www.paperspace.com>
- <https://www.floydhub.com/>

# Demo

# Quick & Easy to Make & Deploy

Maybe not so quick to run...

# Get our code!

<https://github.com/thesecretlab/ReinforceConf2019>

<https://lab.to/NSTDemoJupyter>

*Part 1*

# On-device Machine Learning

*Part 2*

# Neural Style Transfer

*Part 3*

# Demonstration

*Part 4*

# The Future

*Part 4*

# The Future

# What's next for on-device AI?

- Privacy focus!
- Bringing AI-driven solutions to disconnected areas
- Better tools
- Task focused tools
- Smooth mobile experiences

# Privacy

**“I have done nothing wrong, so I  
have nothing to hide.”**

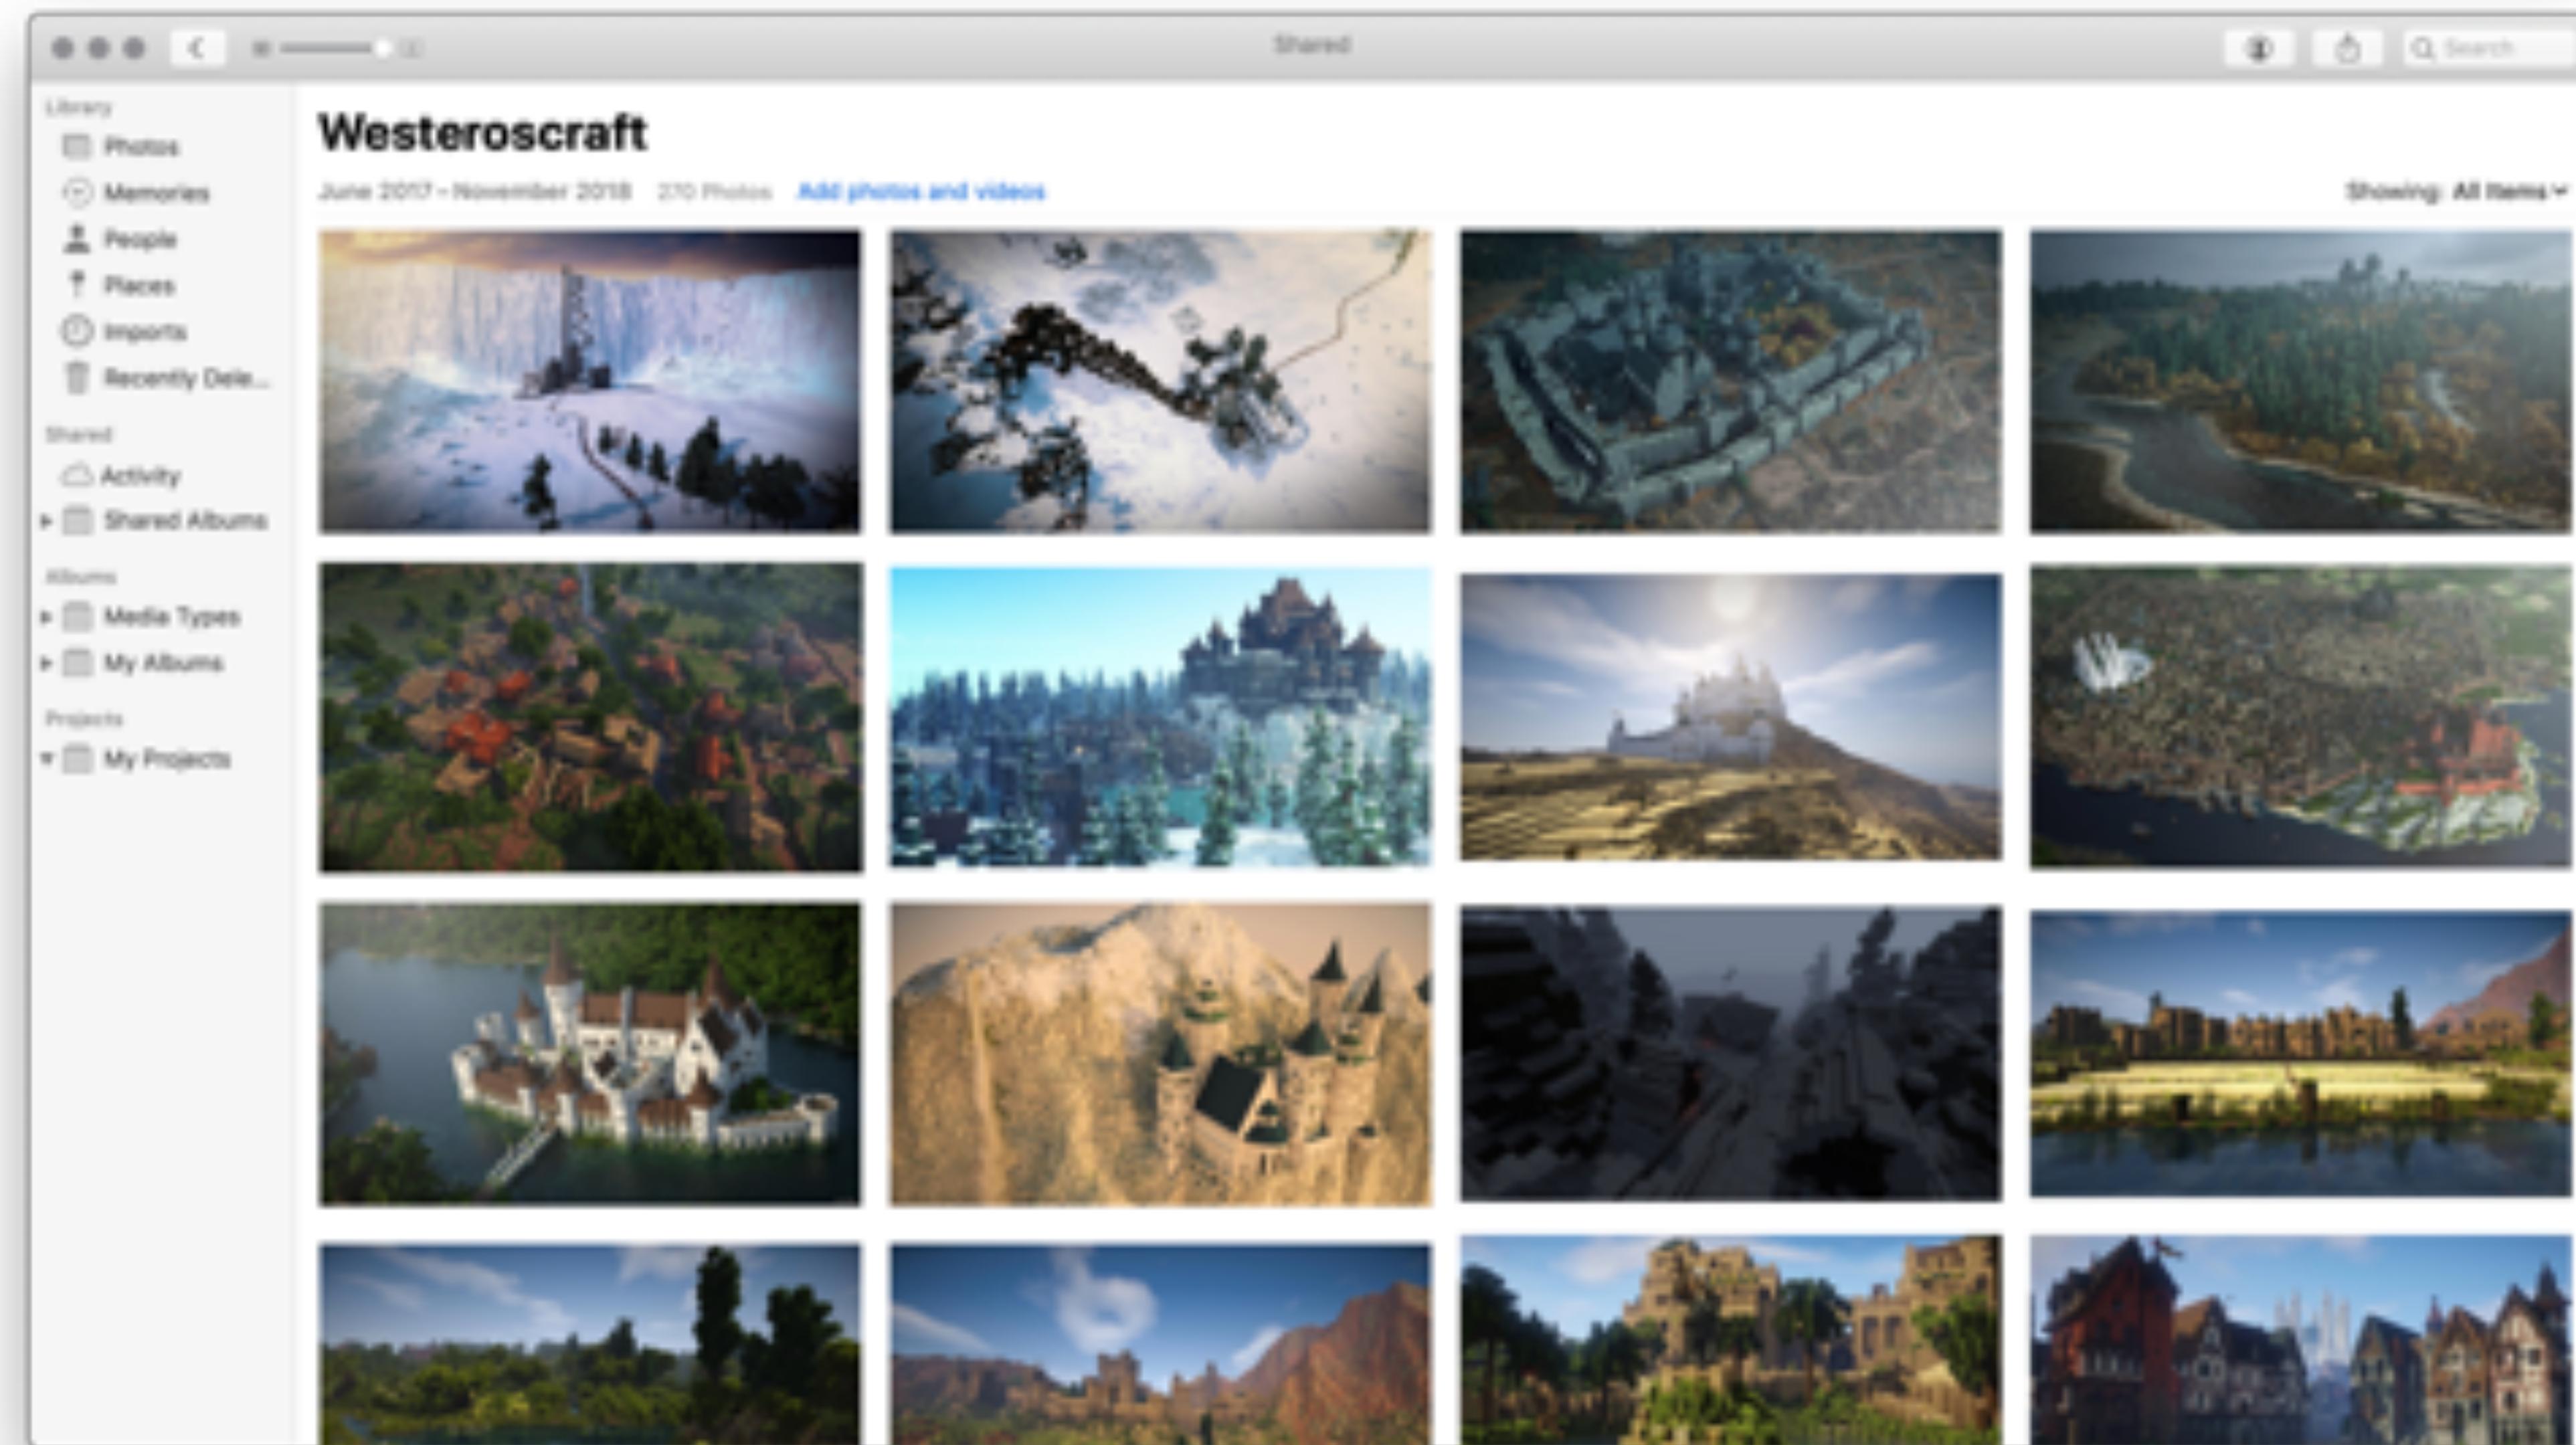
*–Almost everyone we have ever spoken to under  
the age of 25, on personal privacy*

**“My problem with [statements like these] is  
that they accept the premise that privacy is  
about hiding a wrong. It's not.**

**Privacy is an inherent human right, and a  
requirement for maintaining the human  
condition with dignity and respect.”**

*–Bruce Schneier, computer security expert and cryptographer*

# Let's look at a (bad) example



# Let's pretend this is **SUPER** embarrassing

(even it's not; it's obviously really cool 😊)



**Schmapple**

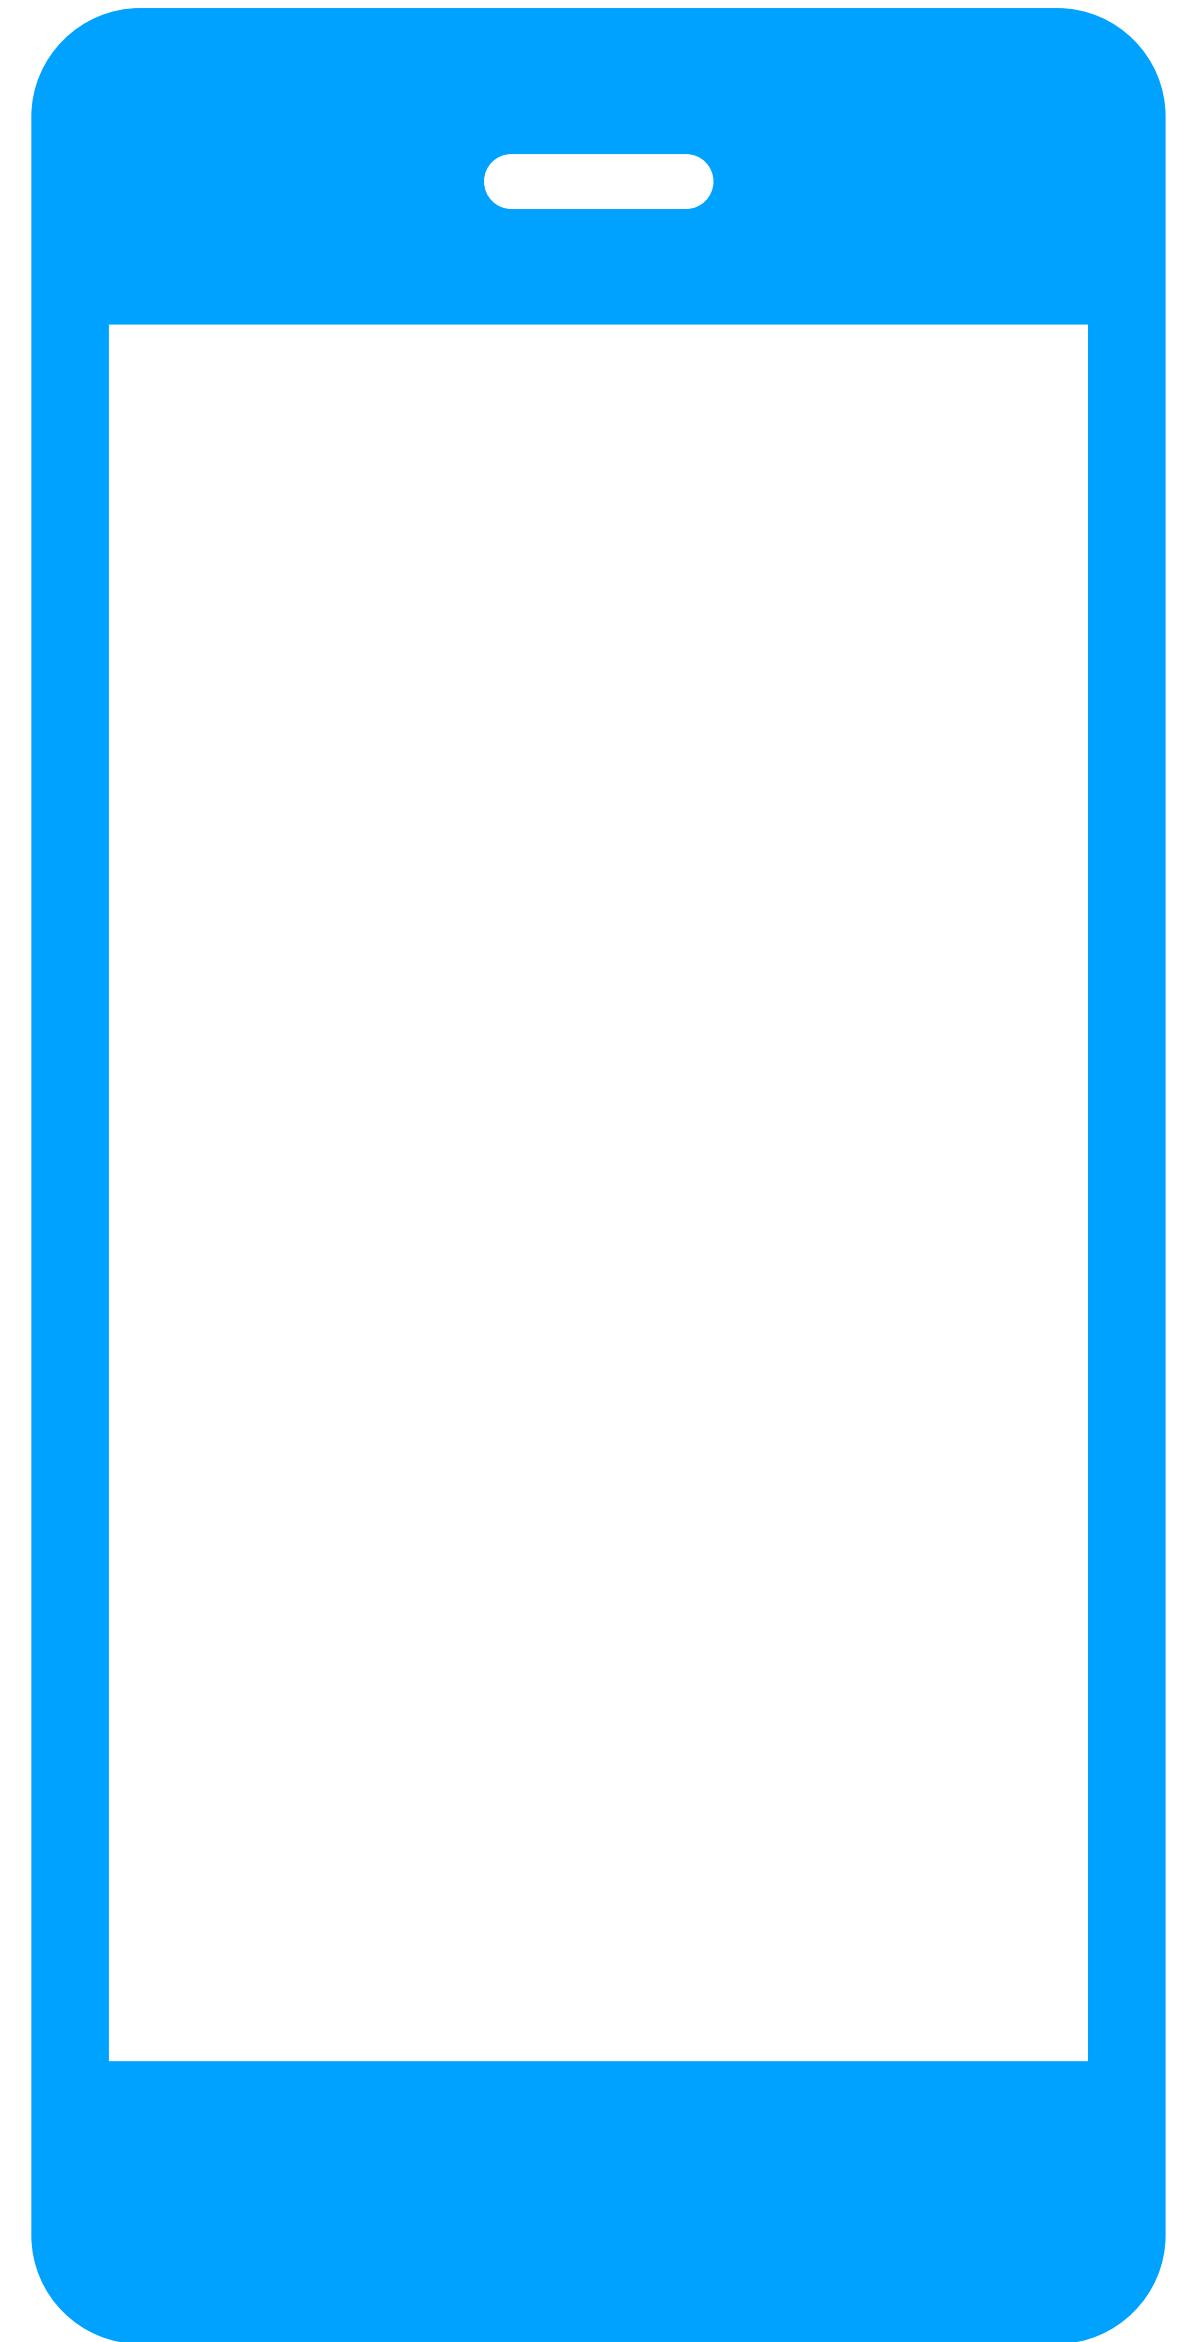


**Schmandroid**

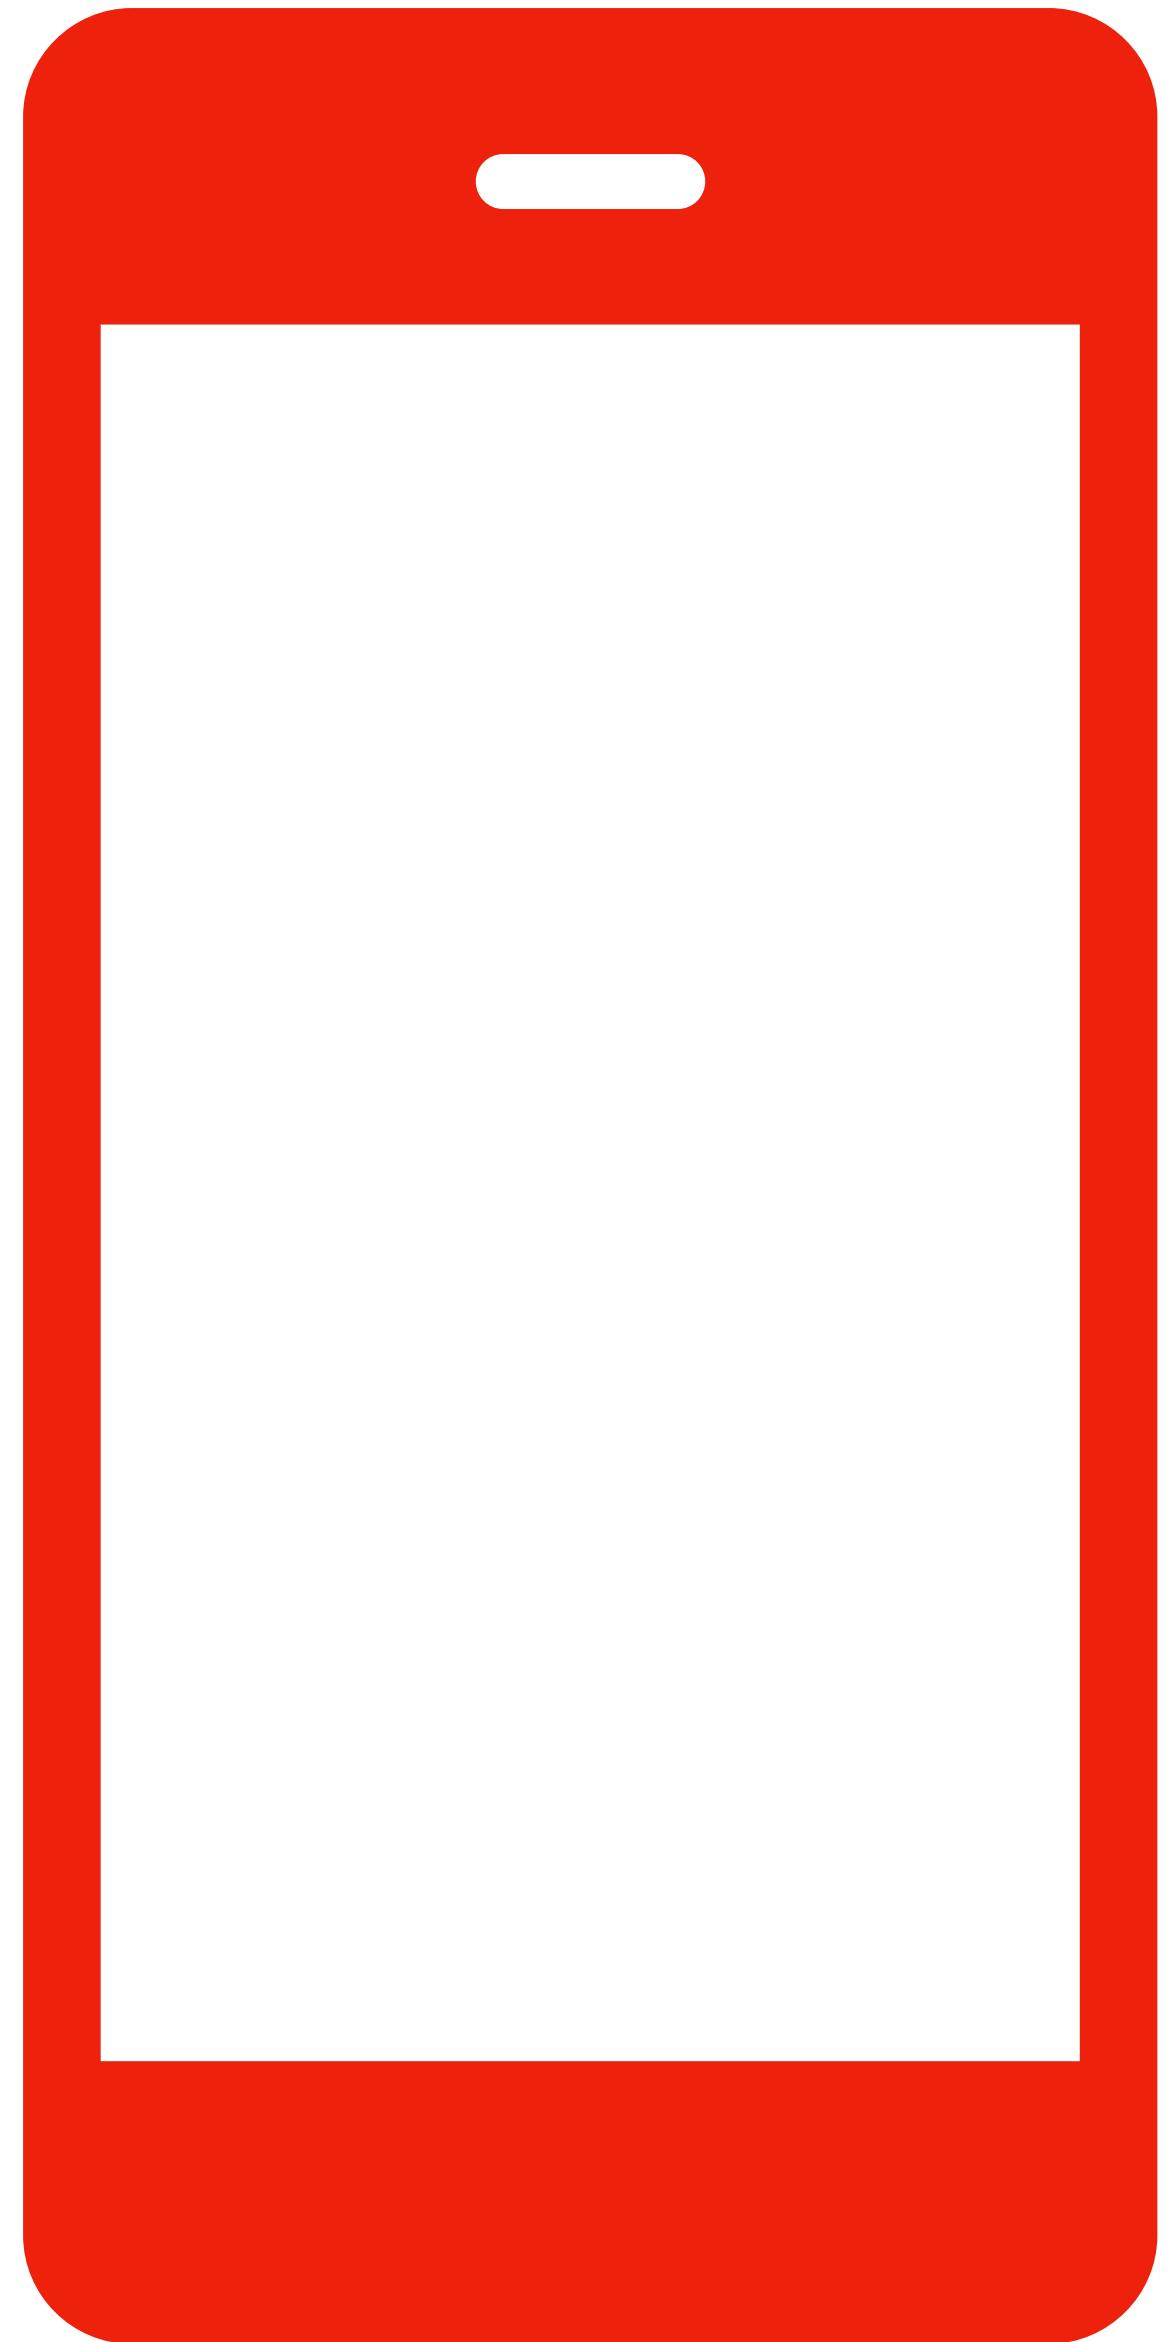


# **DISCLAIMER**

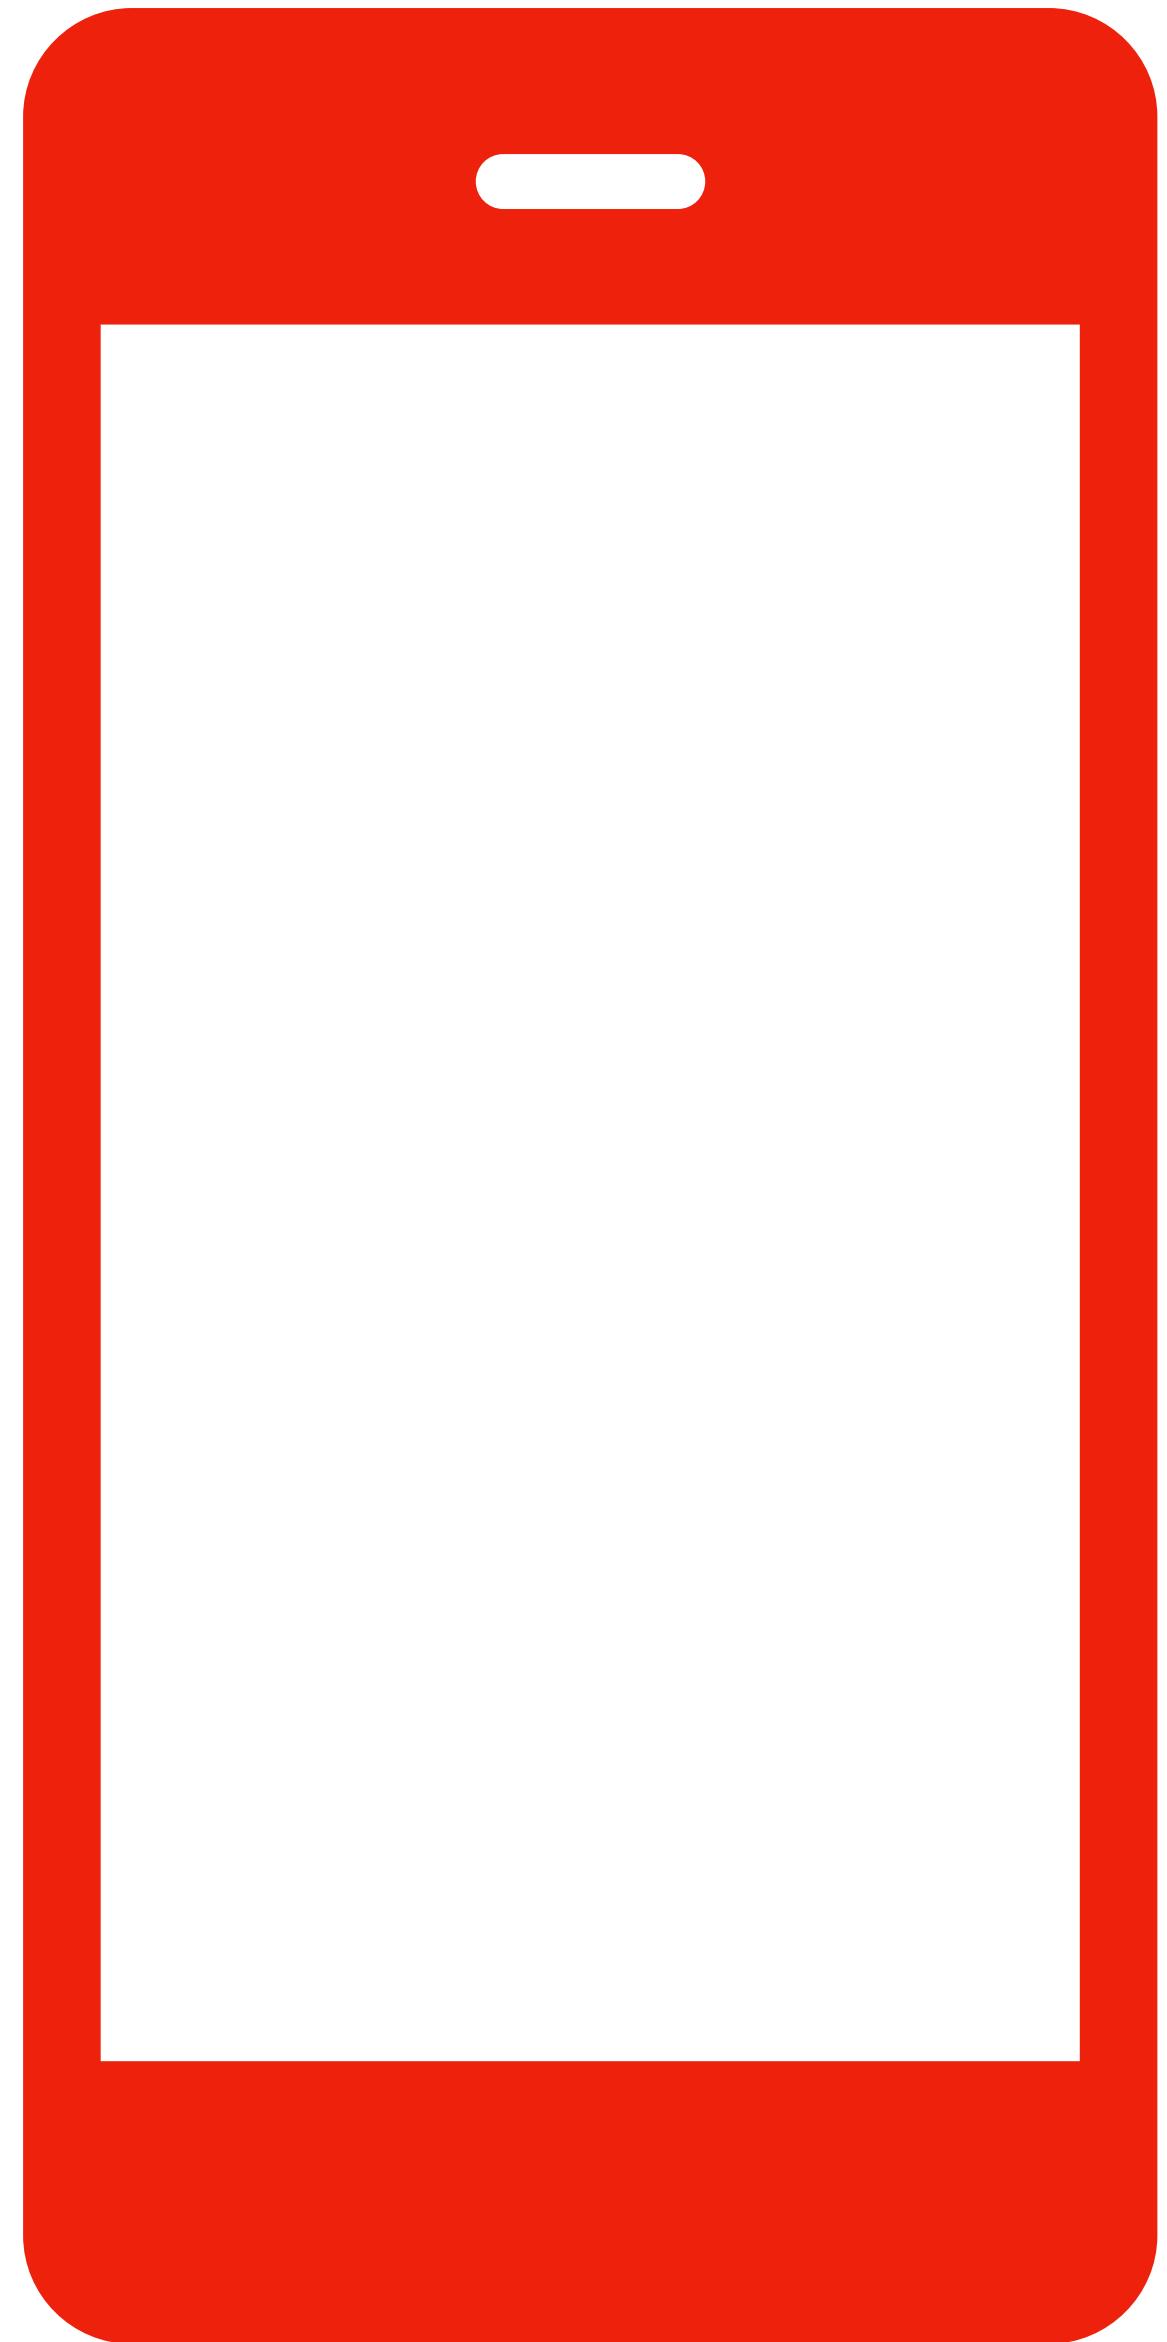
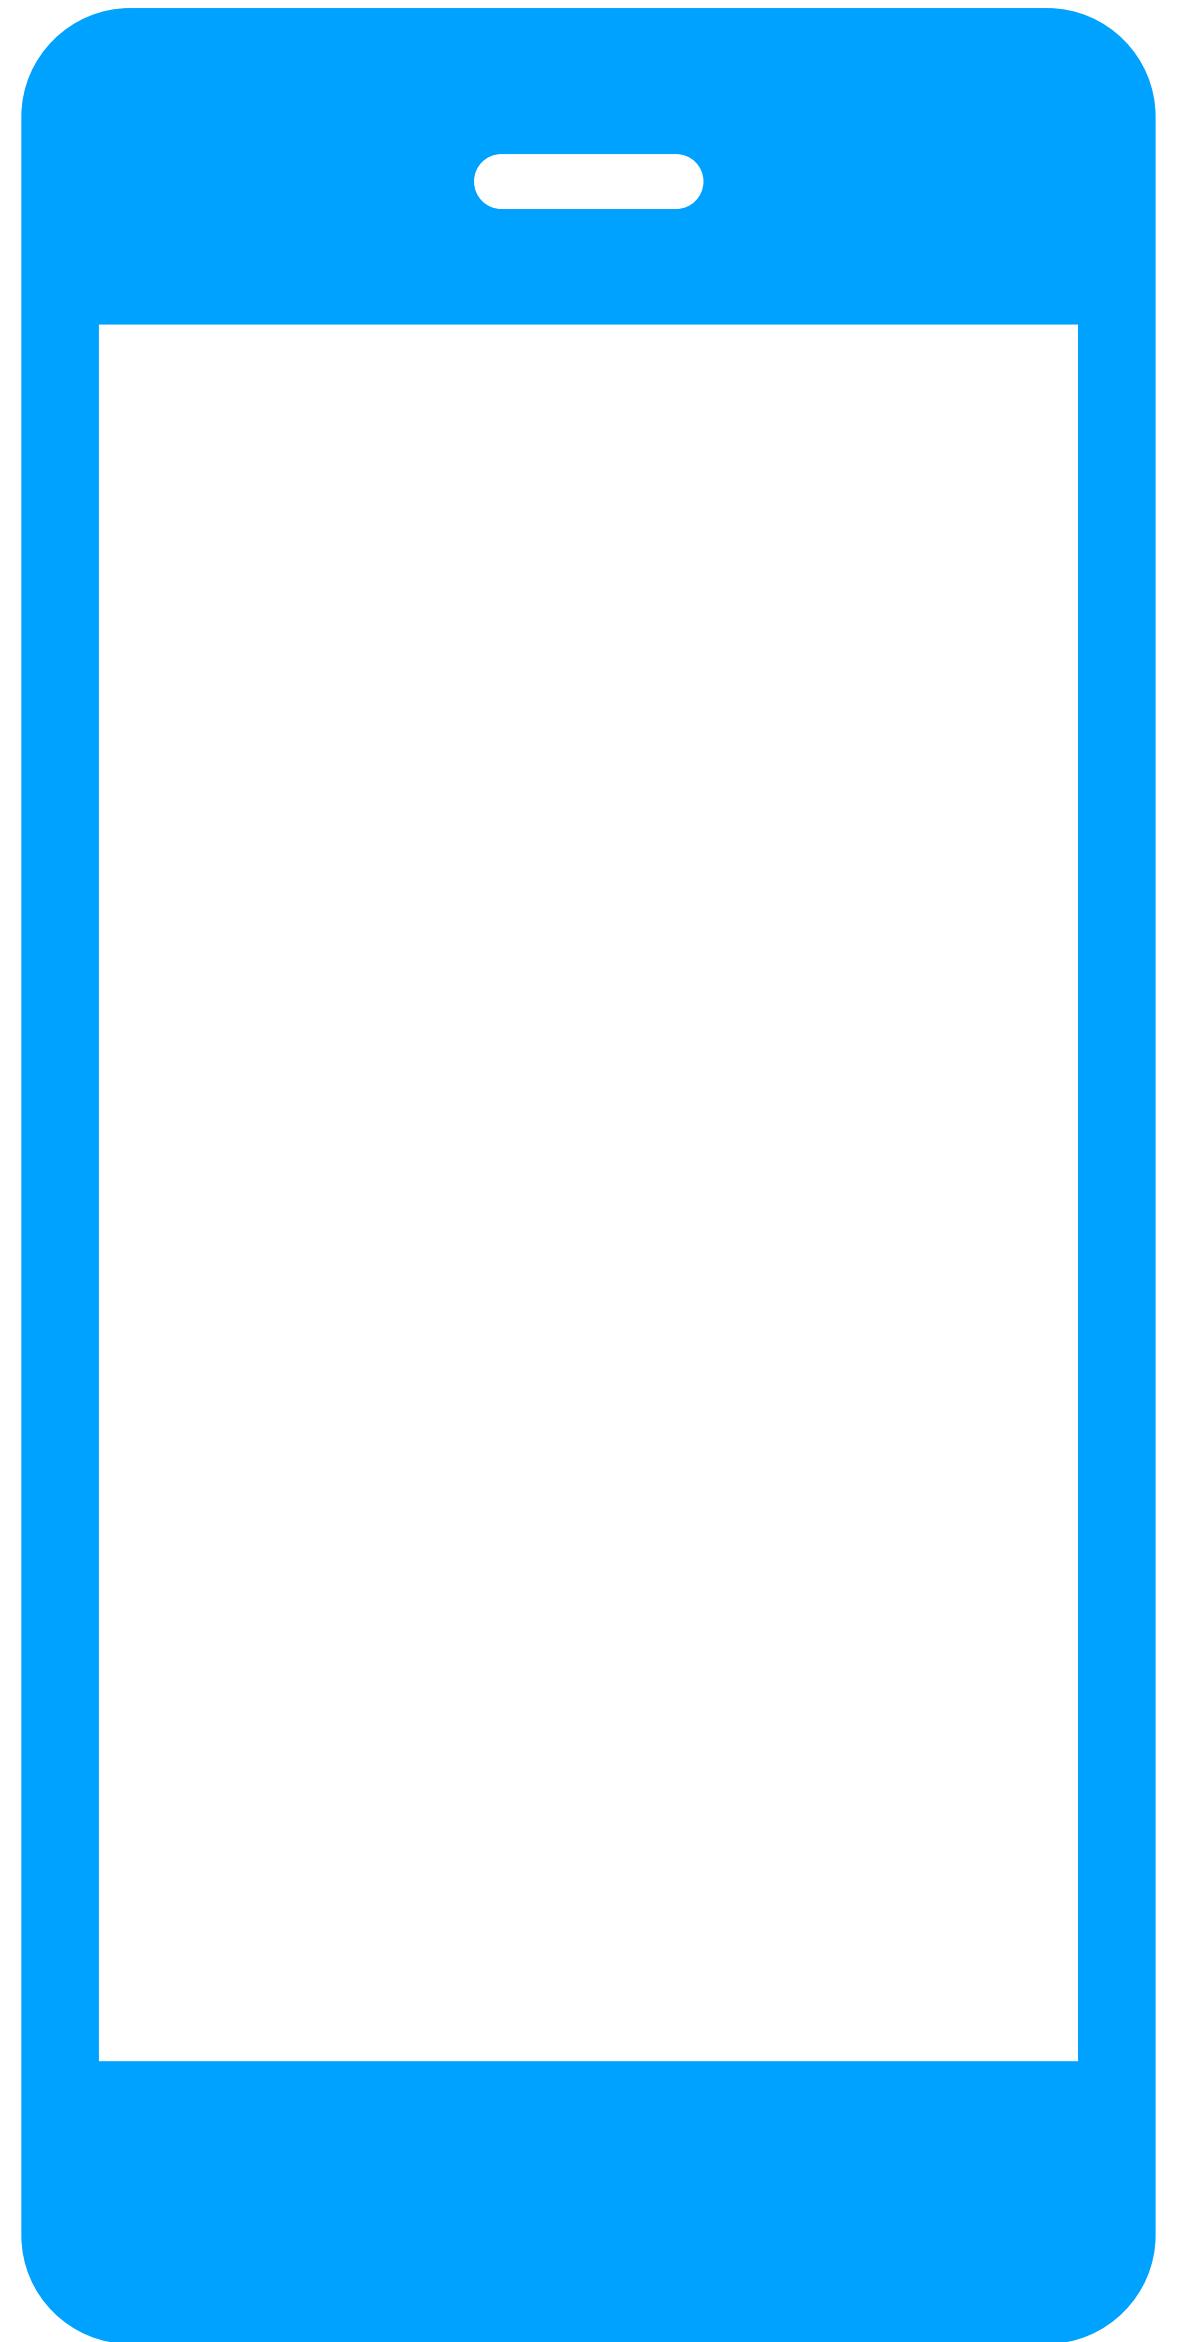
**Schmapple**



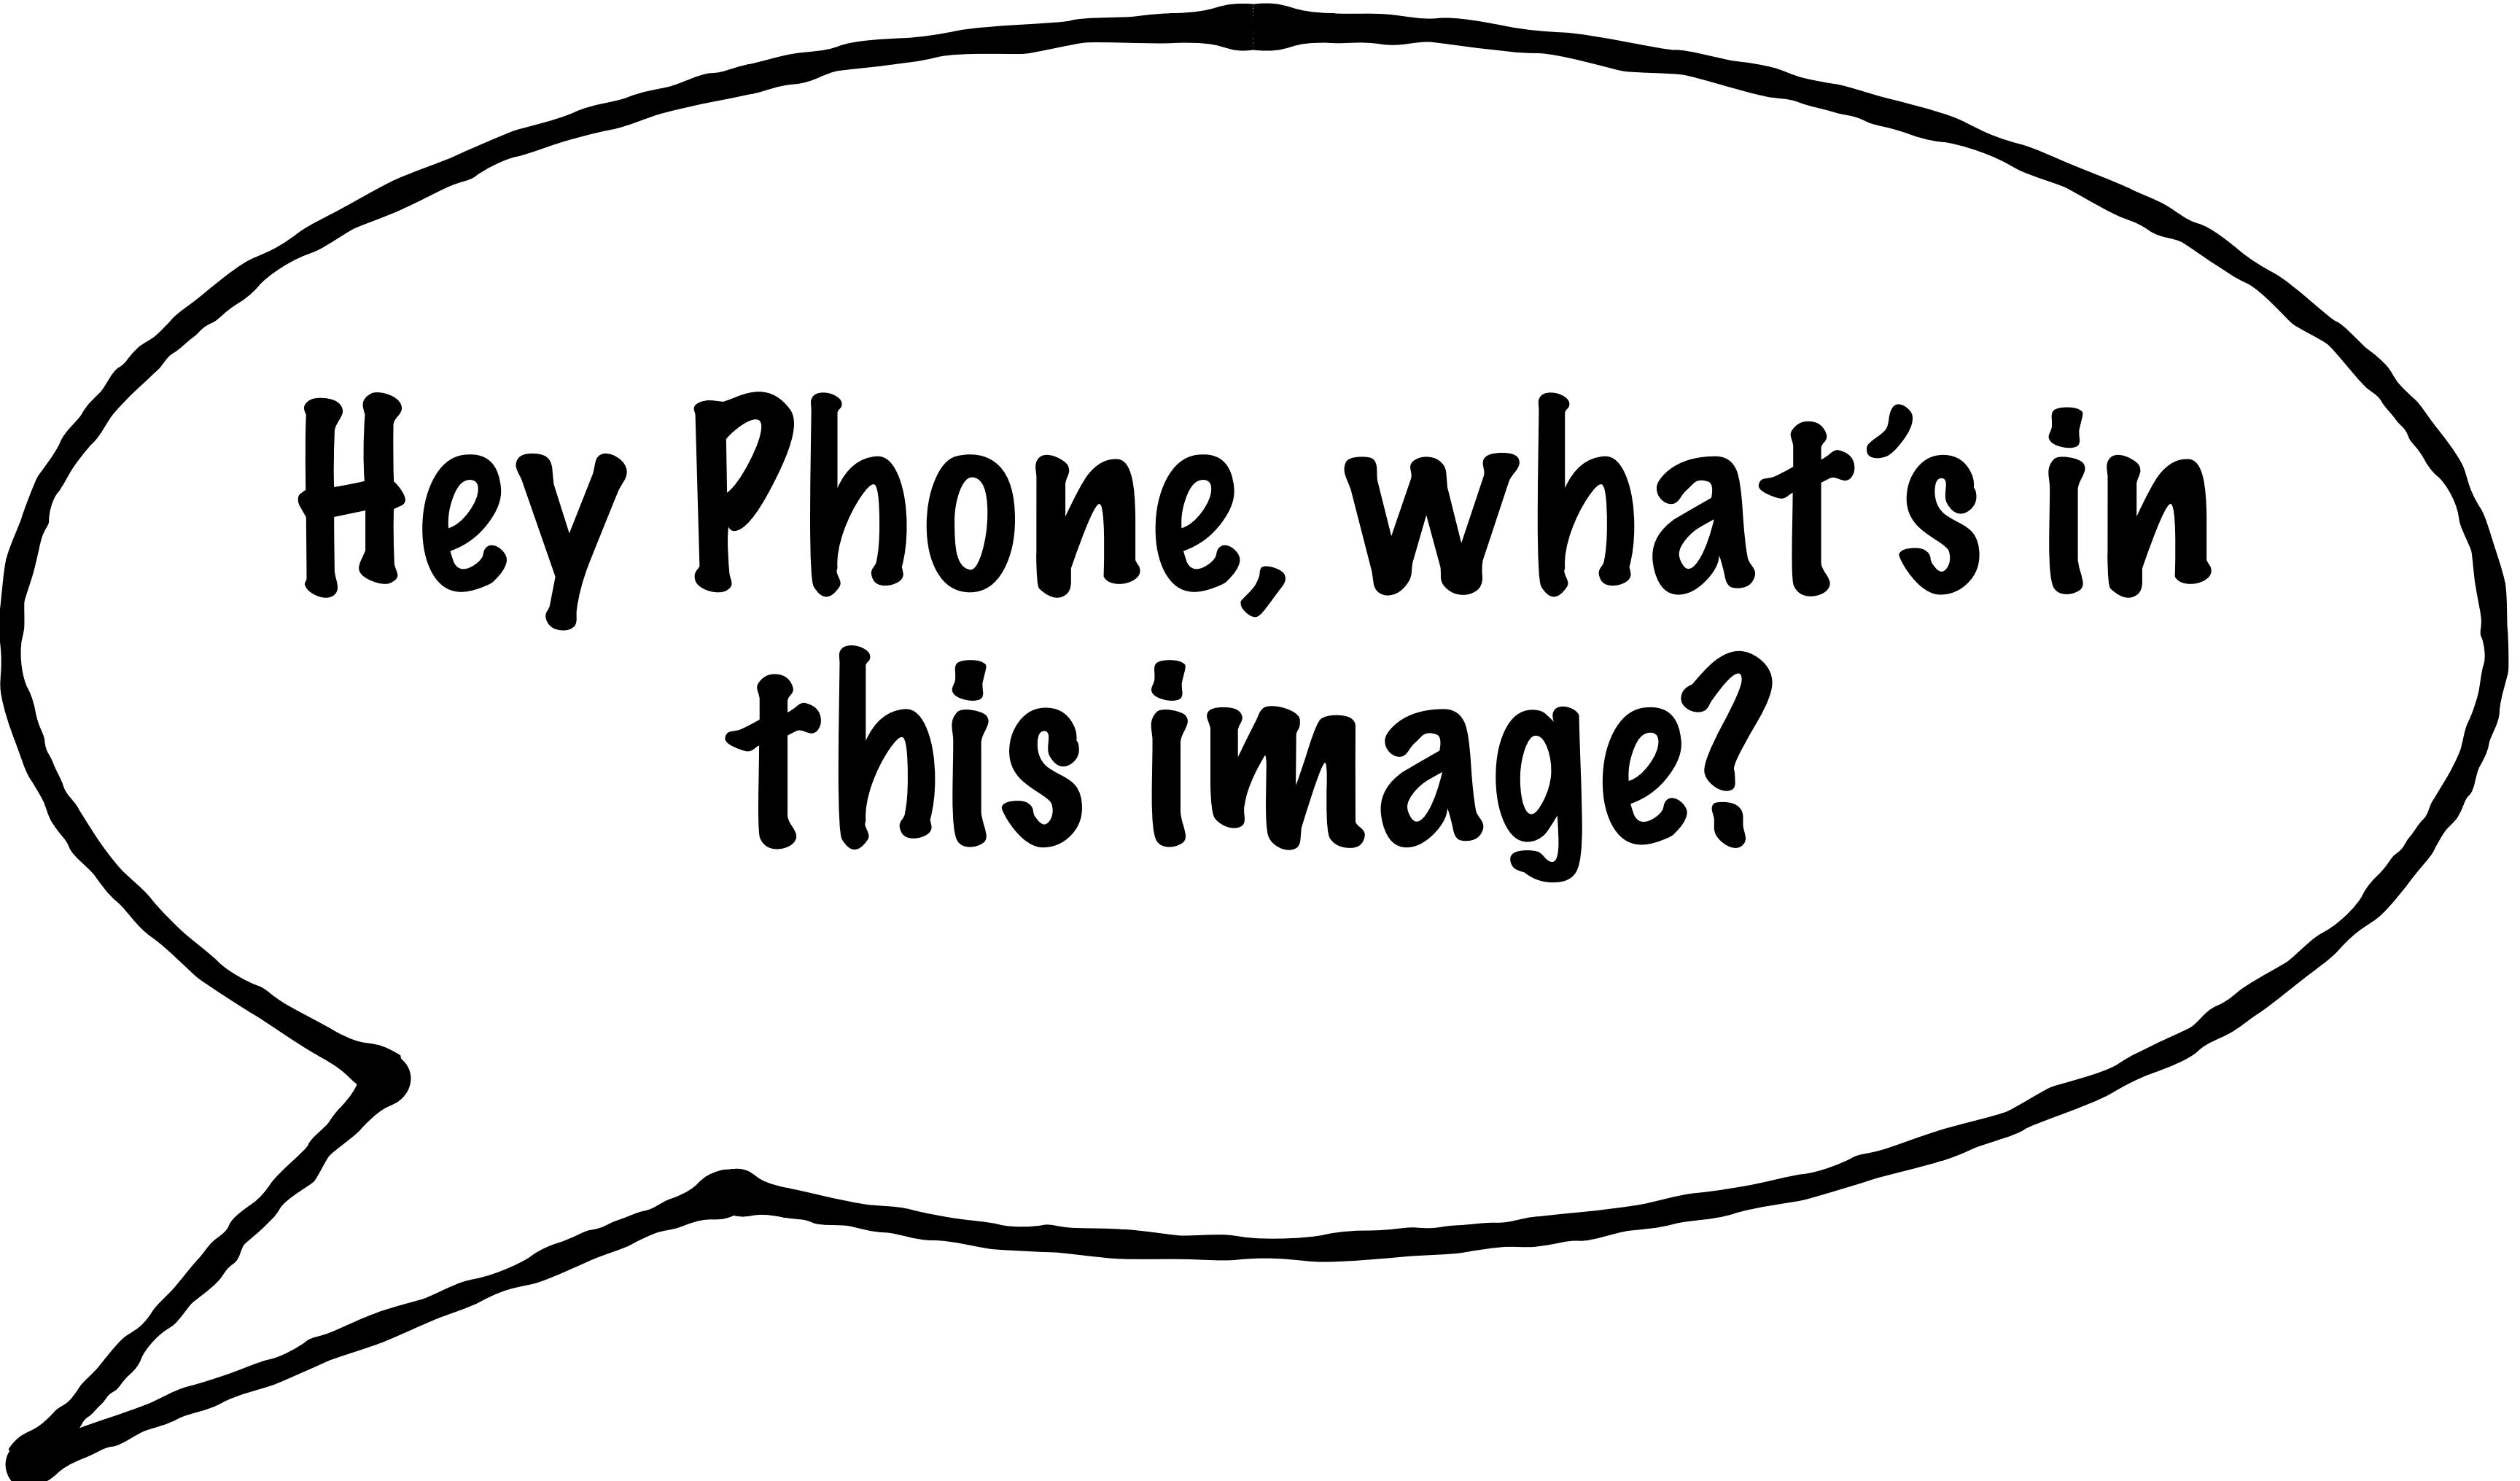
**Schmandroid**



**On-device**

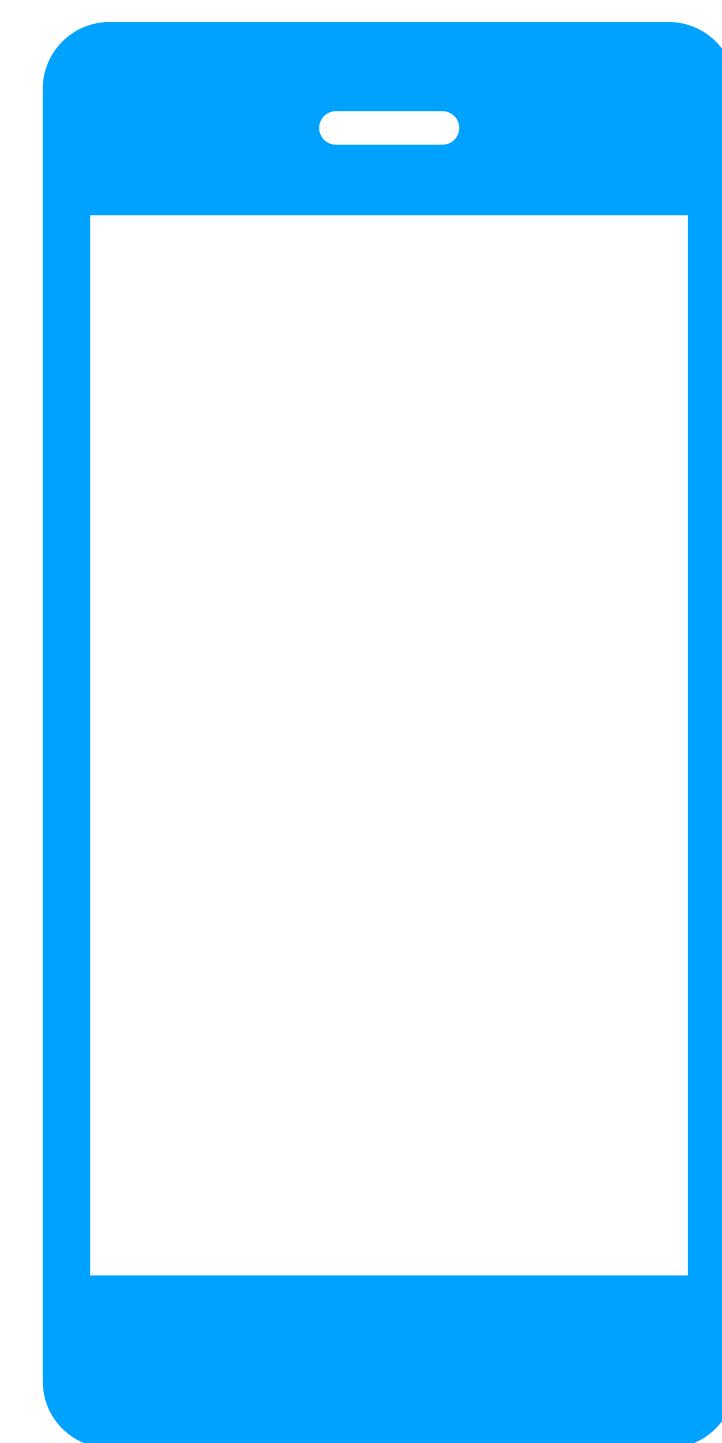


**Not**



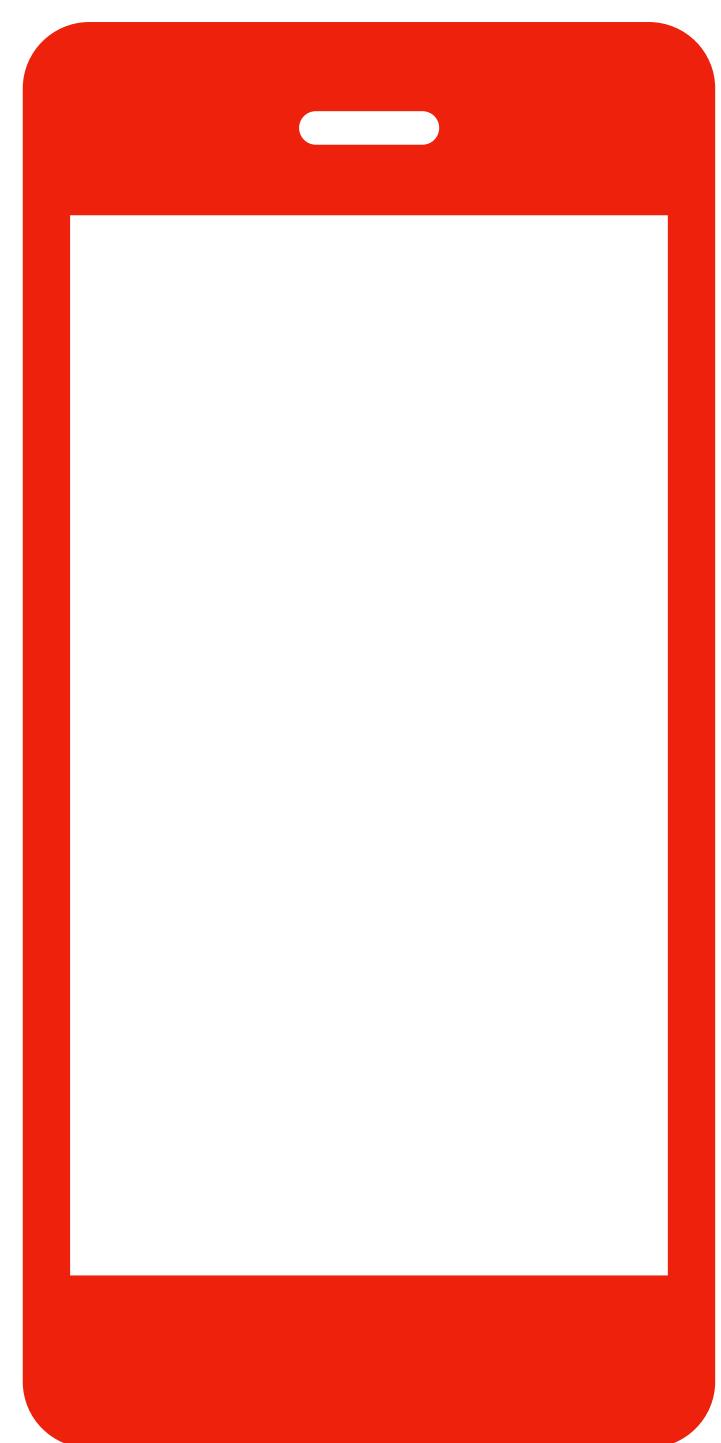
Hey Phone, what's in  
this image?

Hey Model, what's  
in this image?



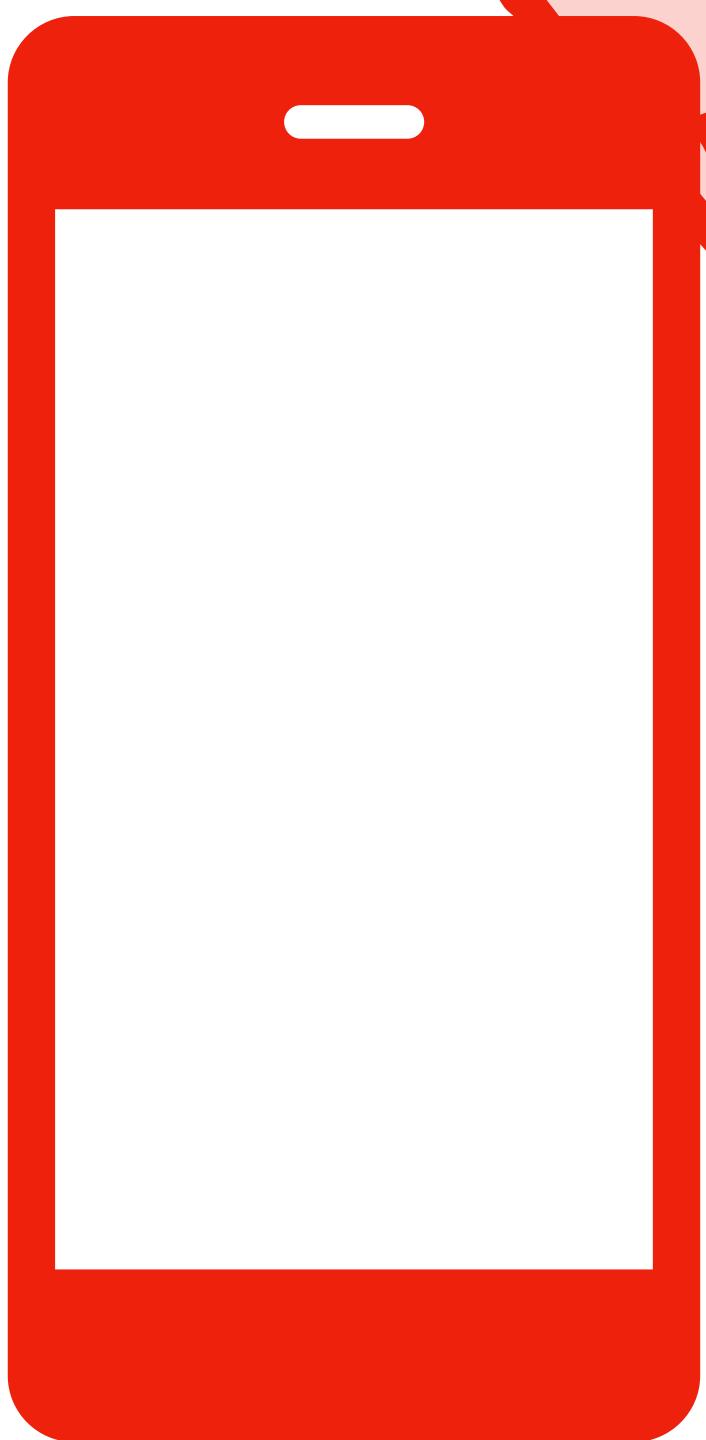
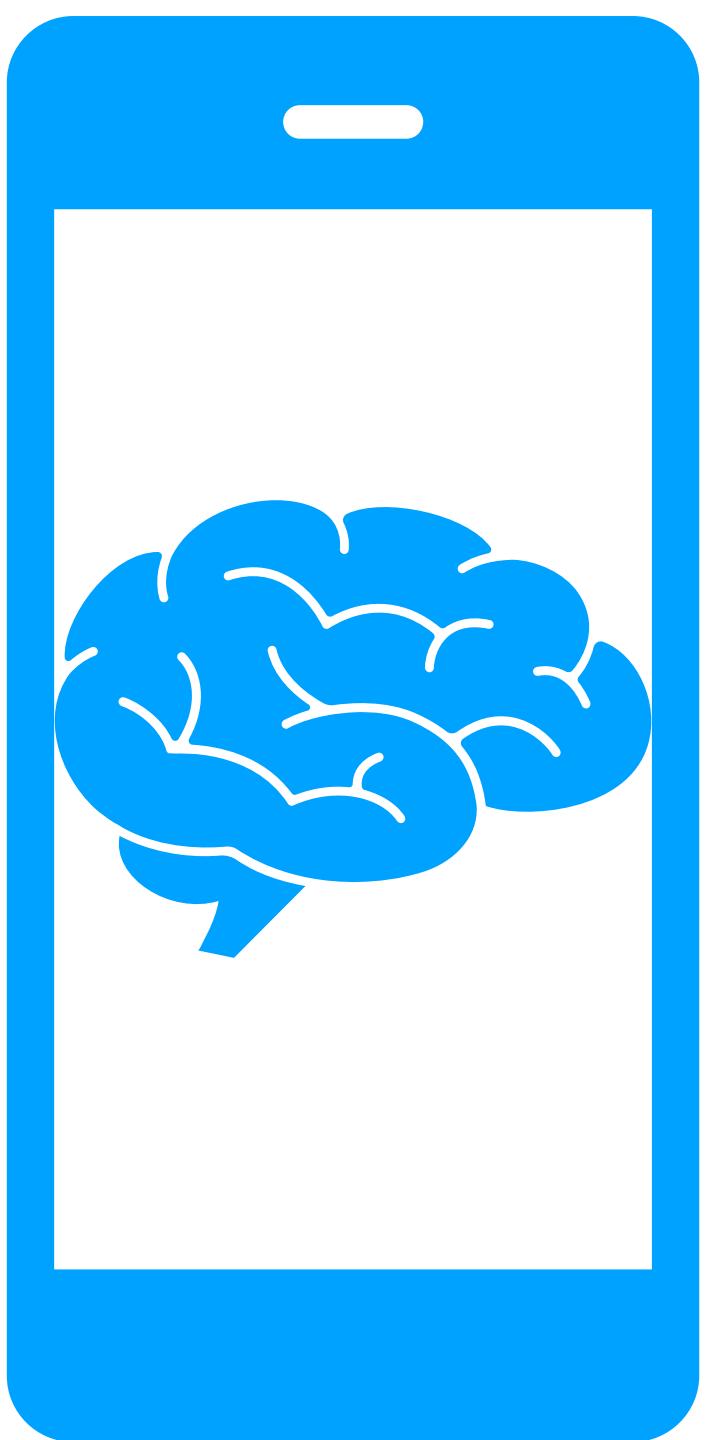
Schmapple

Hey Model, what's  
in this image?



Schmandroid

**Schmapple**

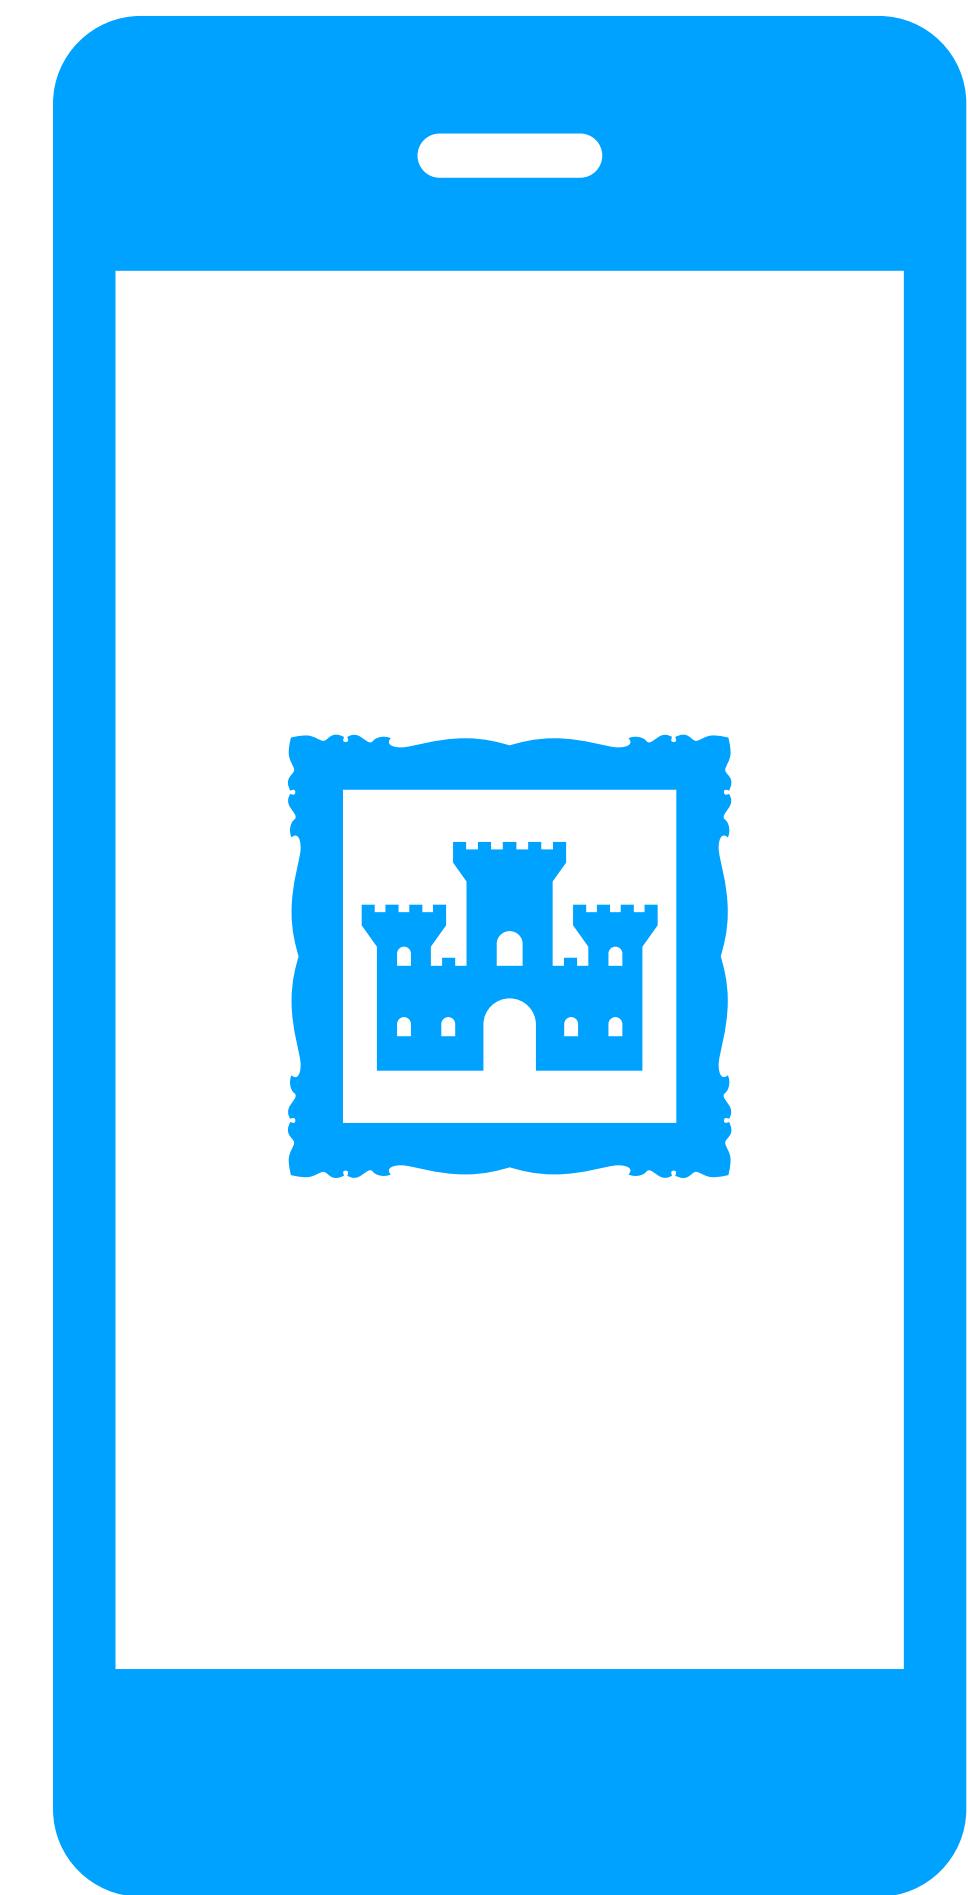


**Schmandroid**

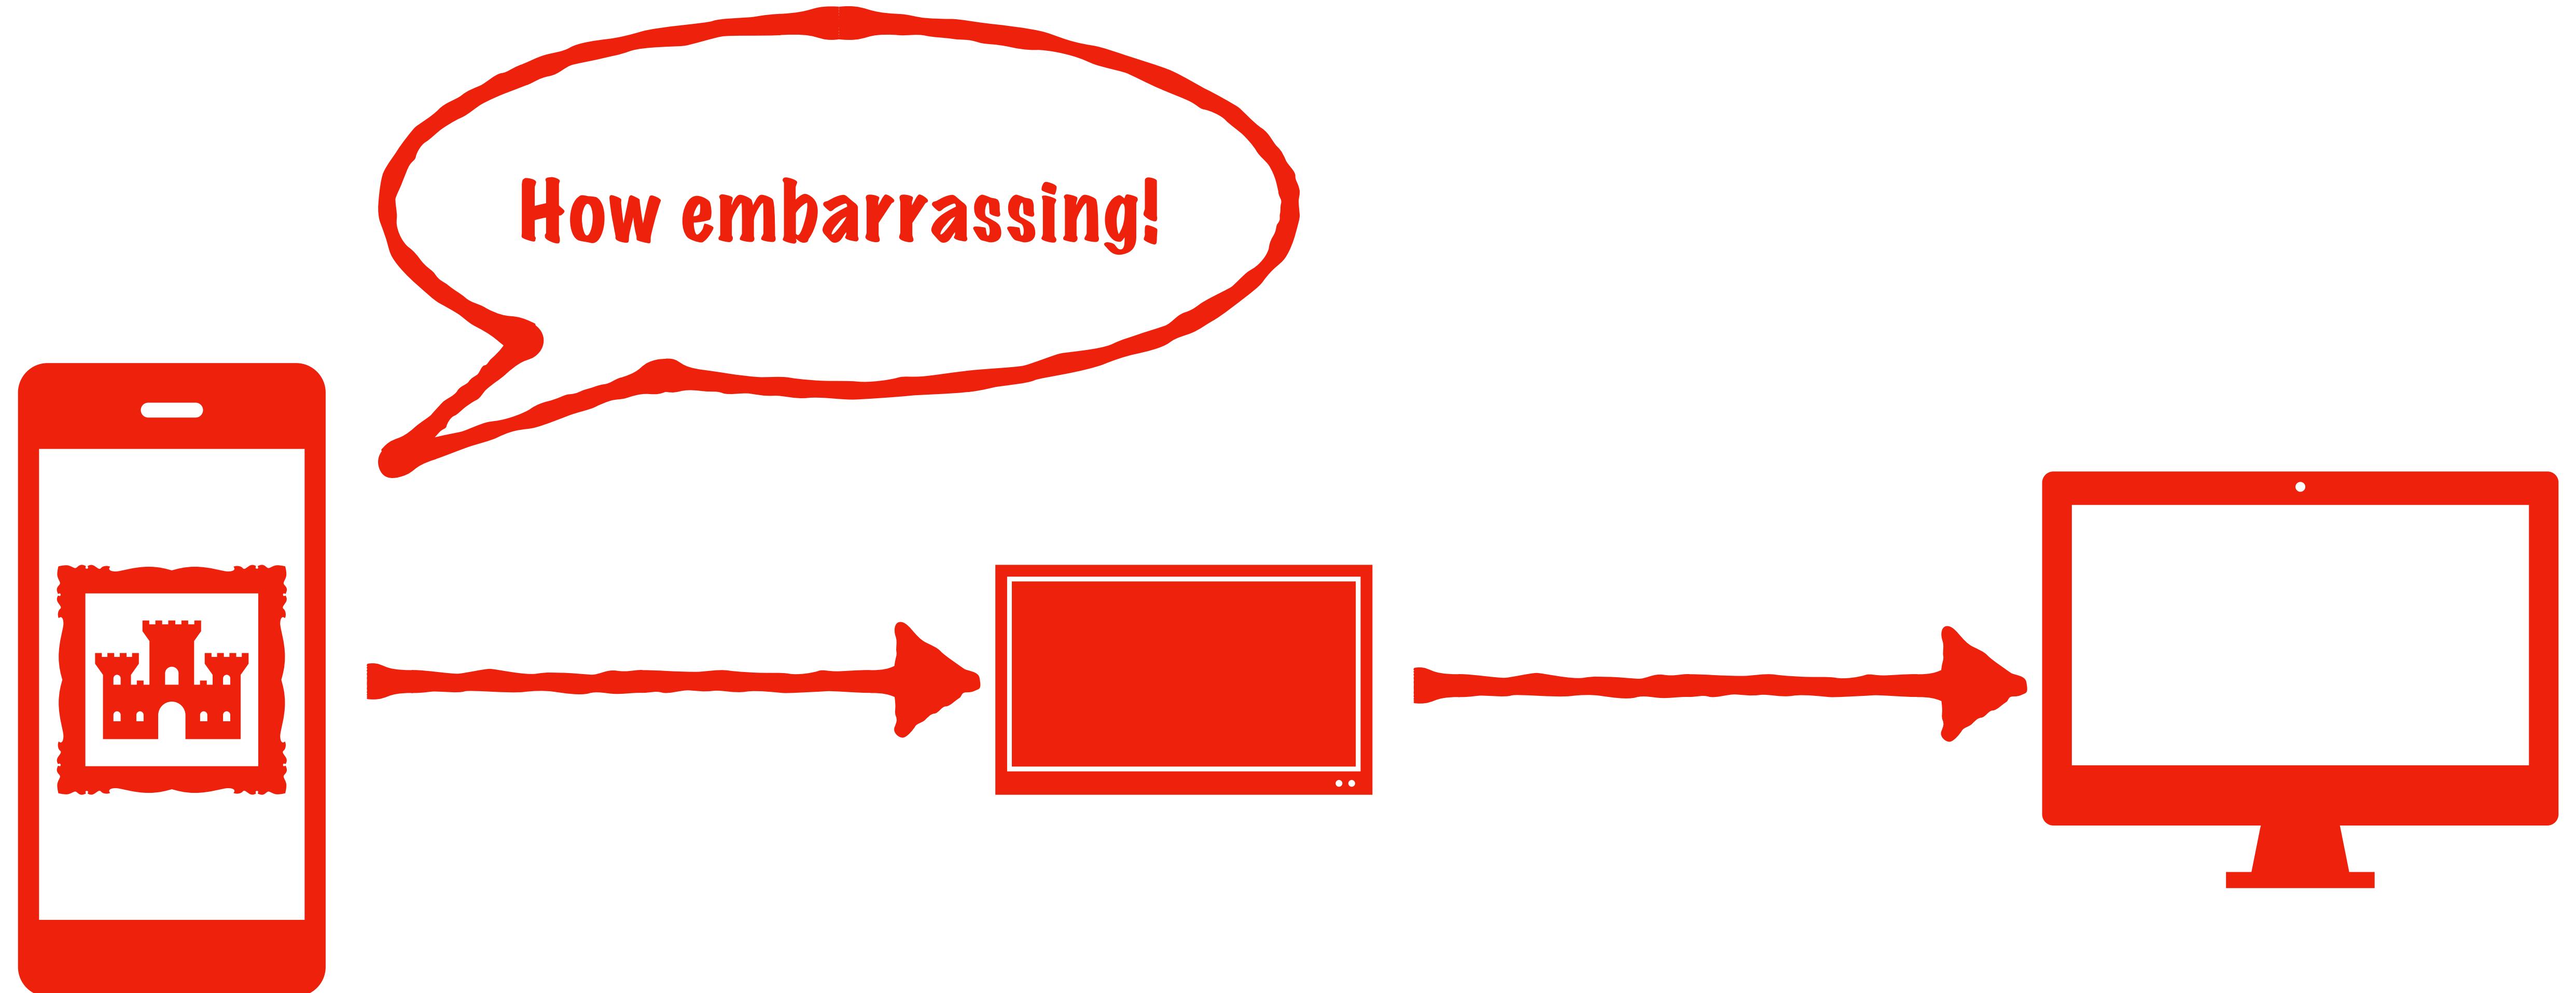


**Somebody else's  
hardware**

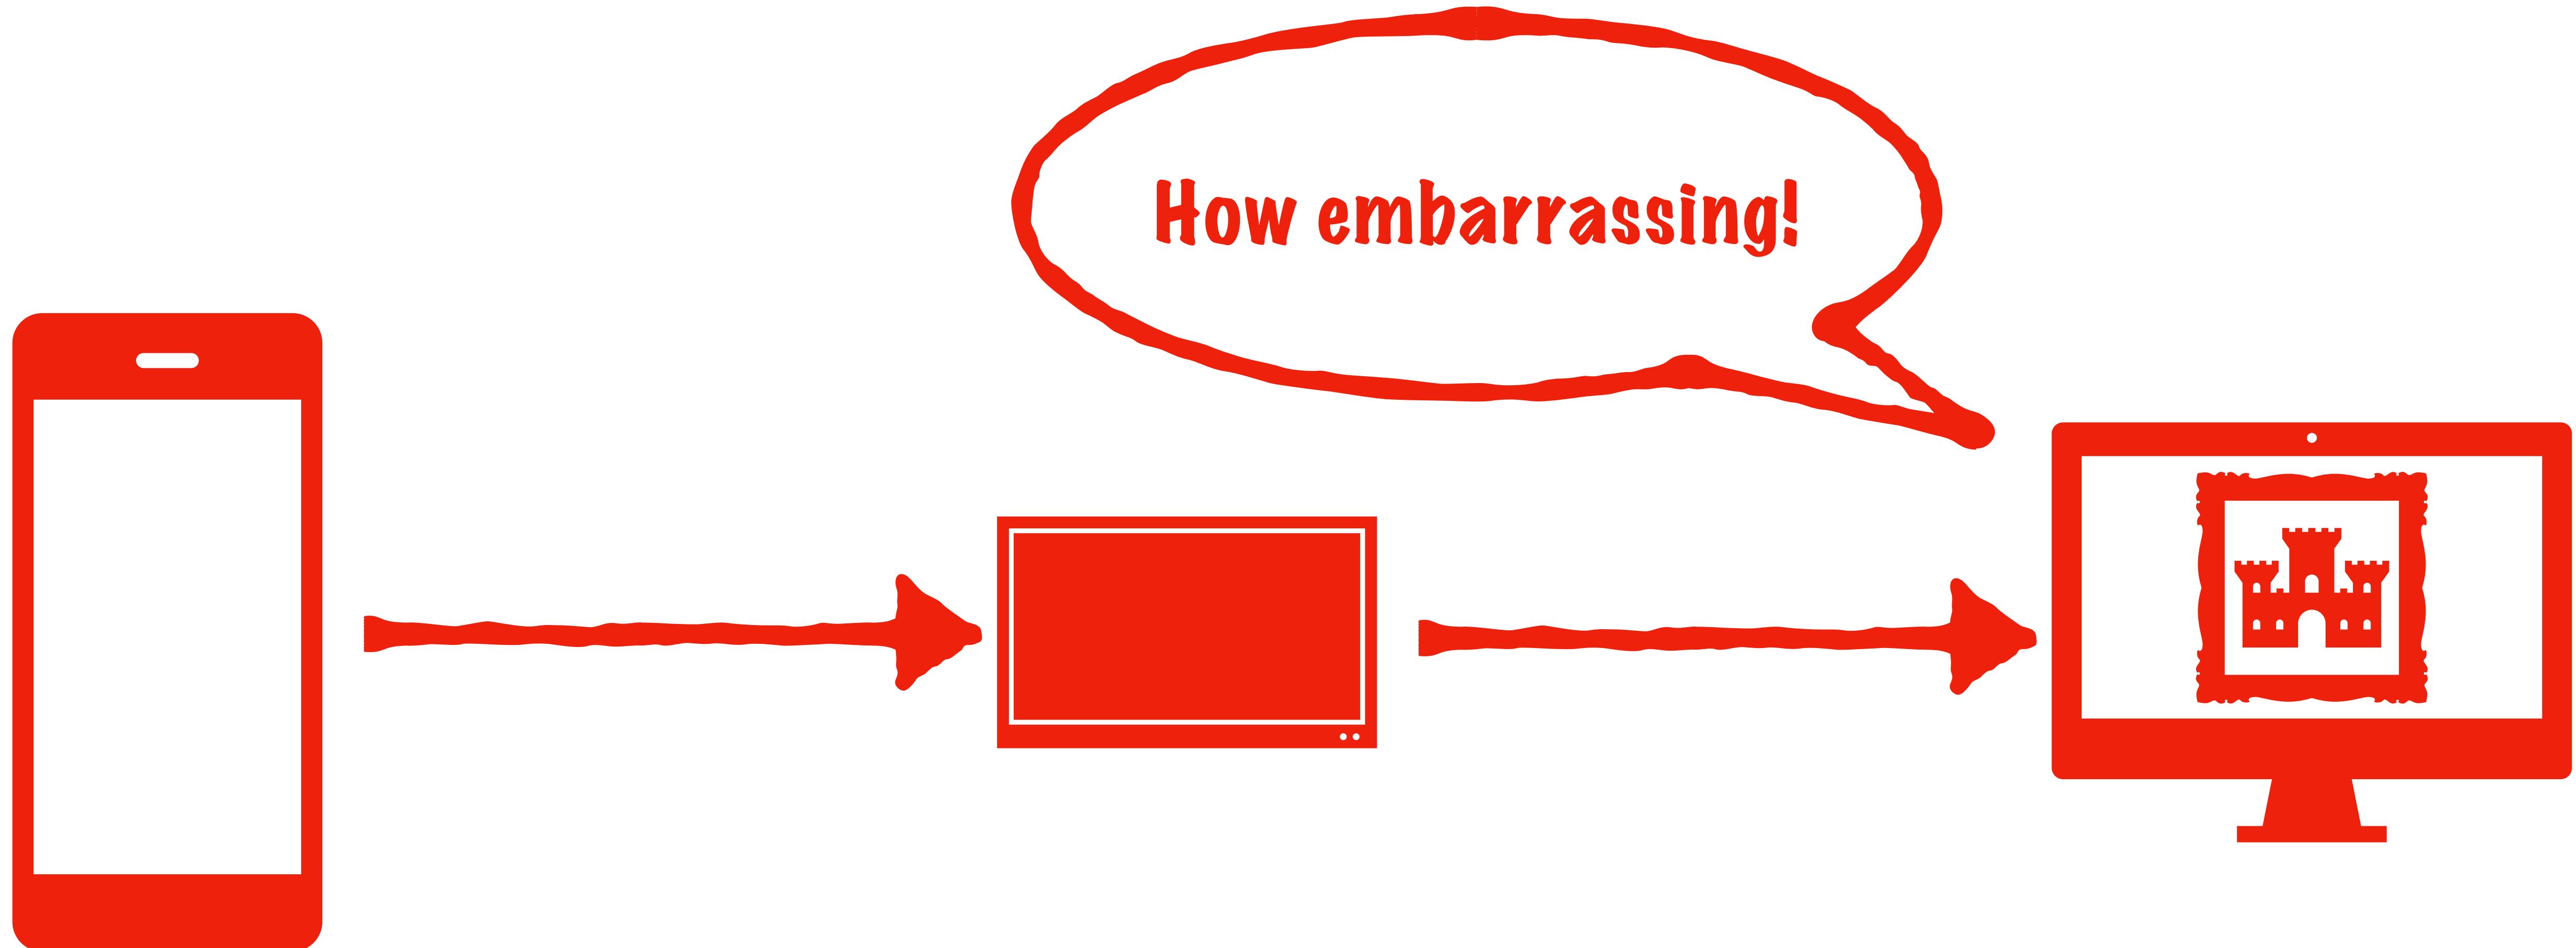
**The Internet**

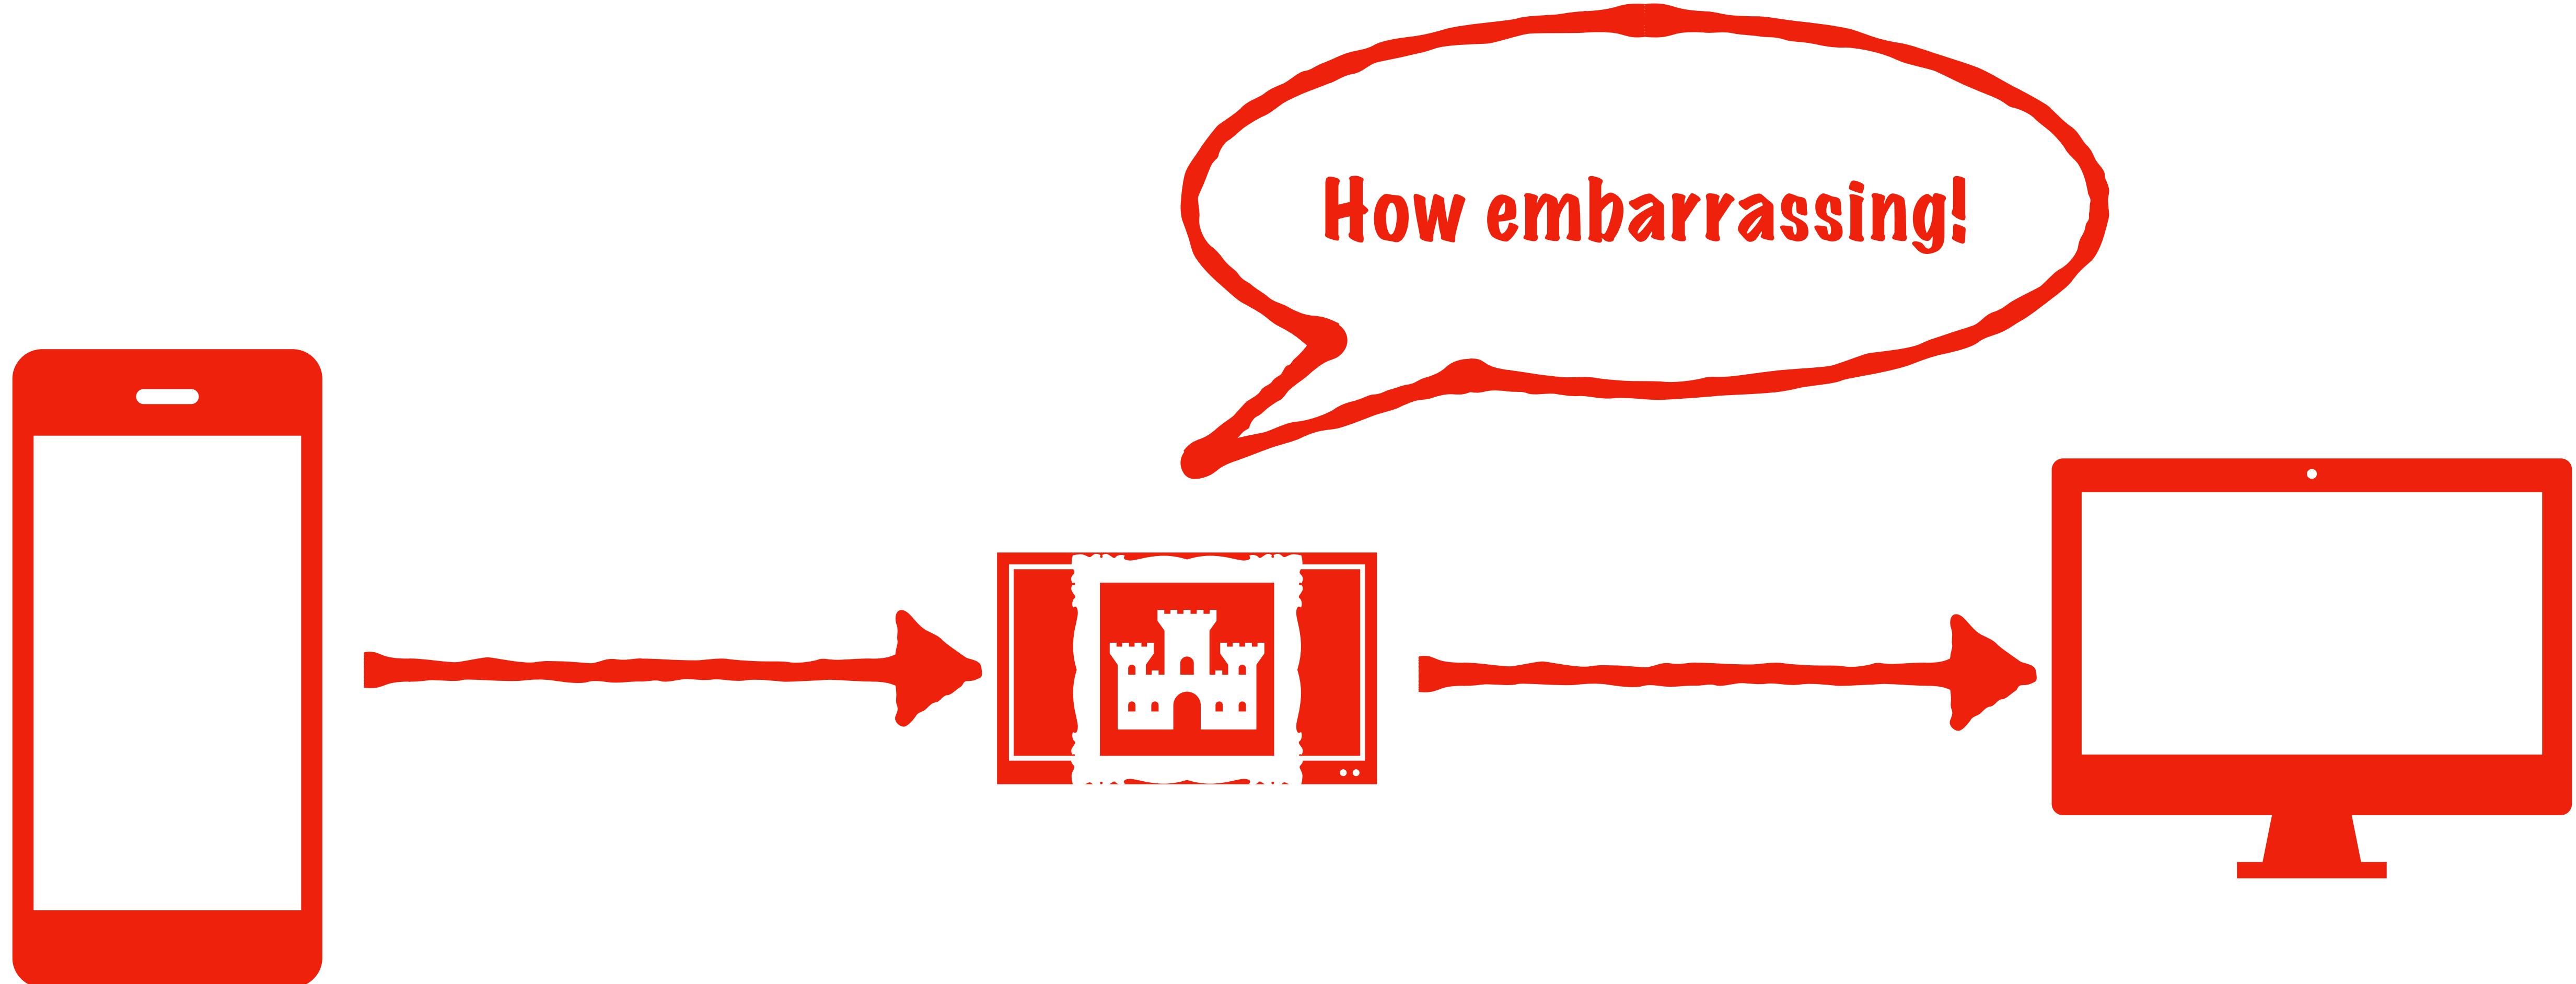


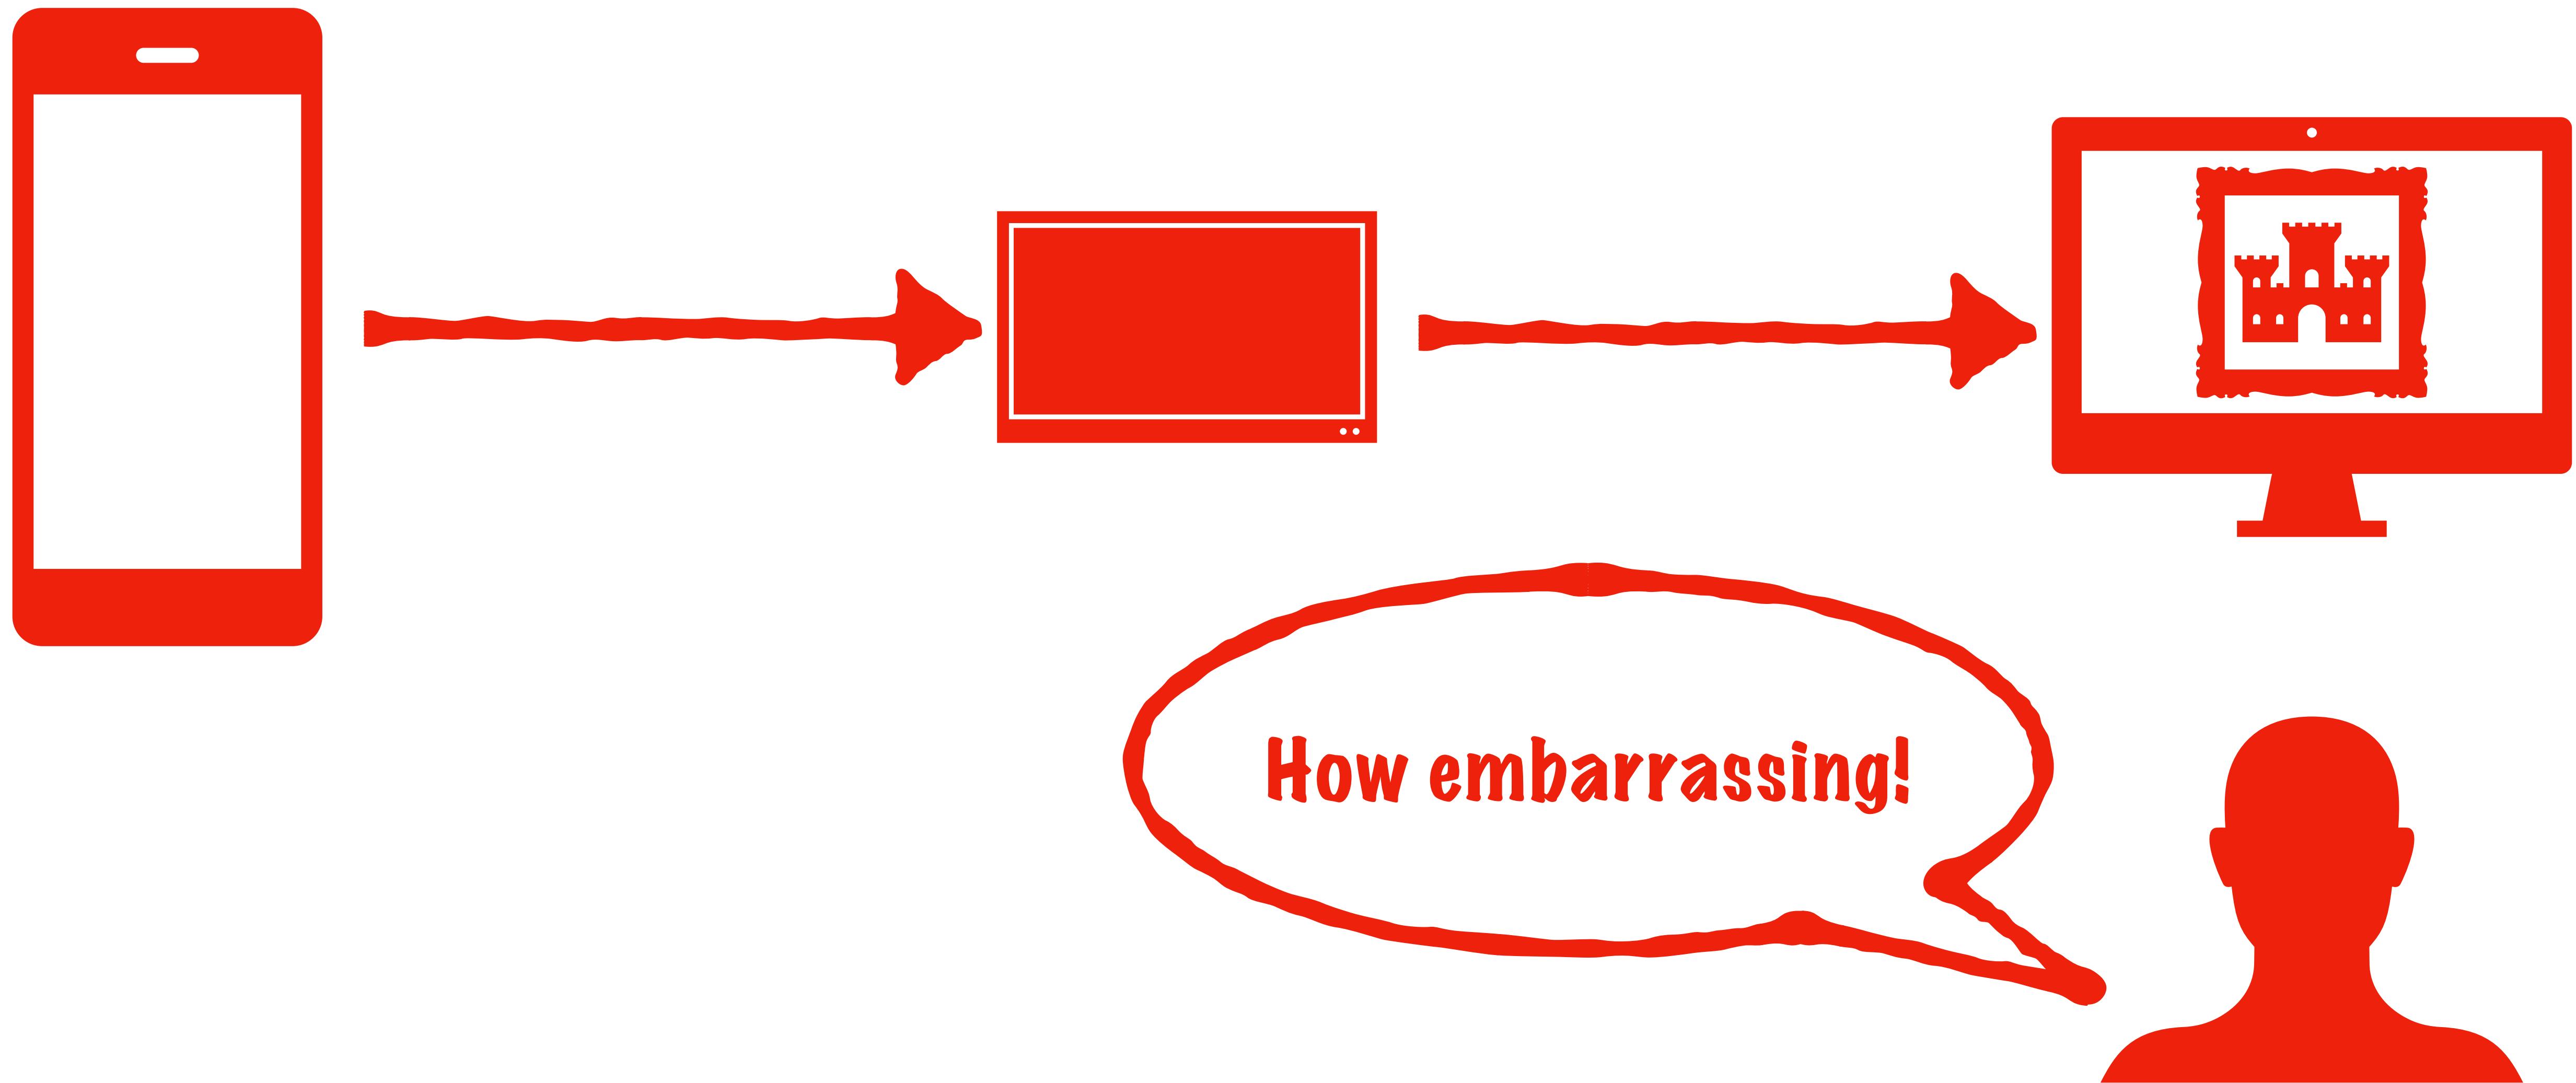
How embarrassing!

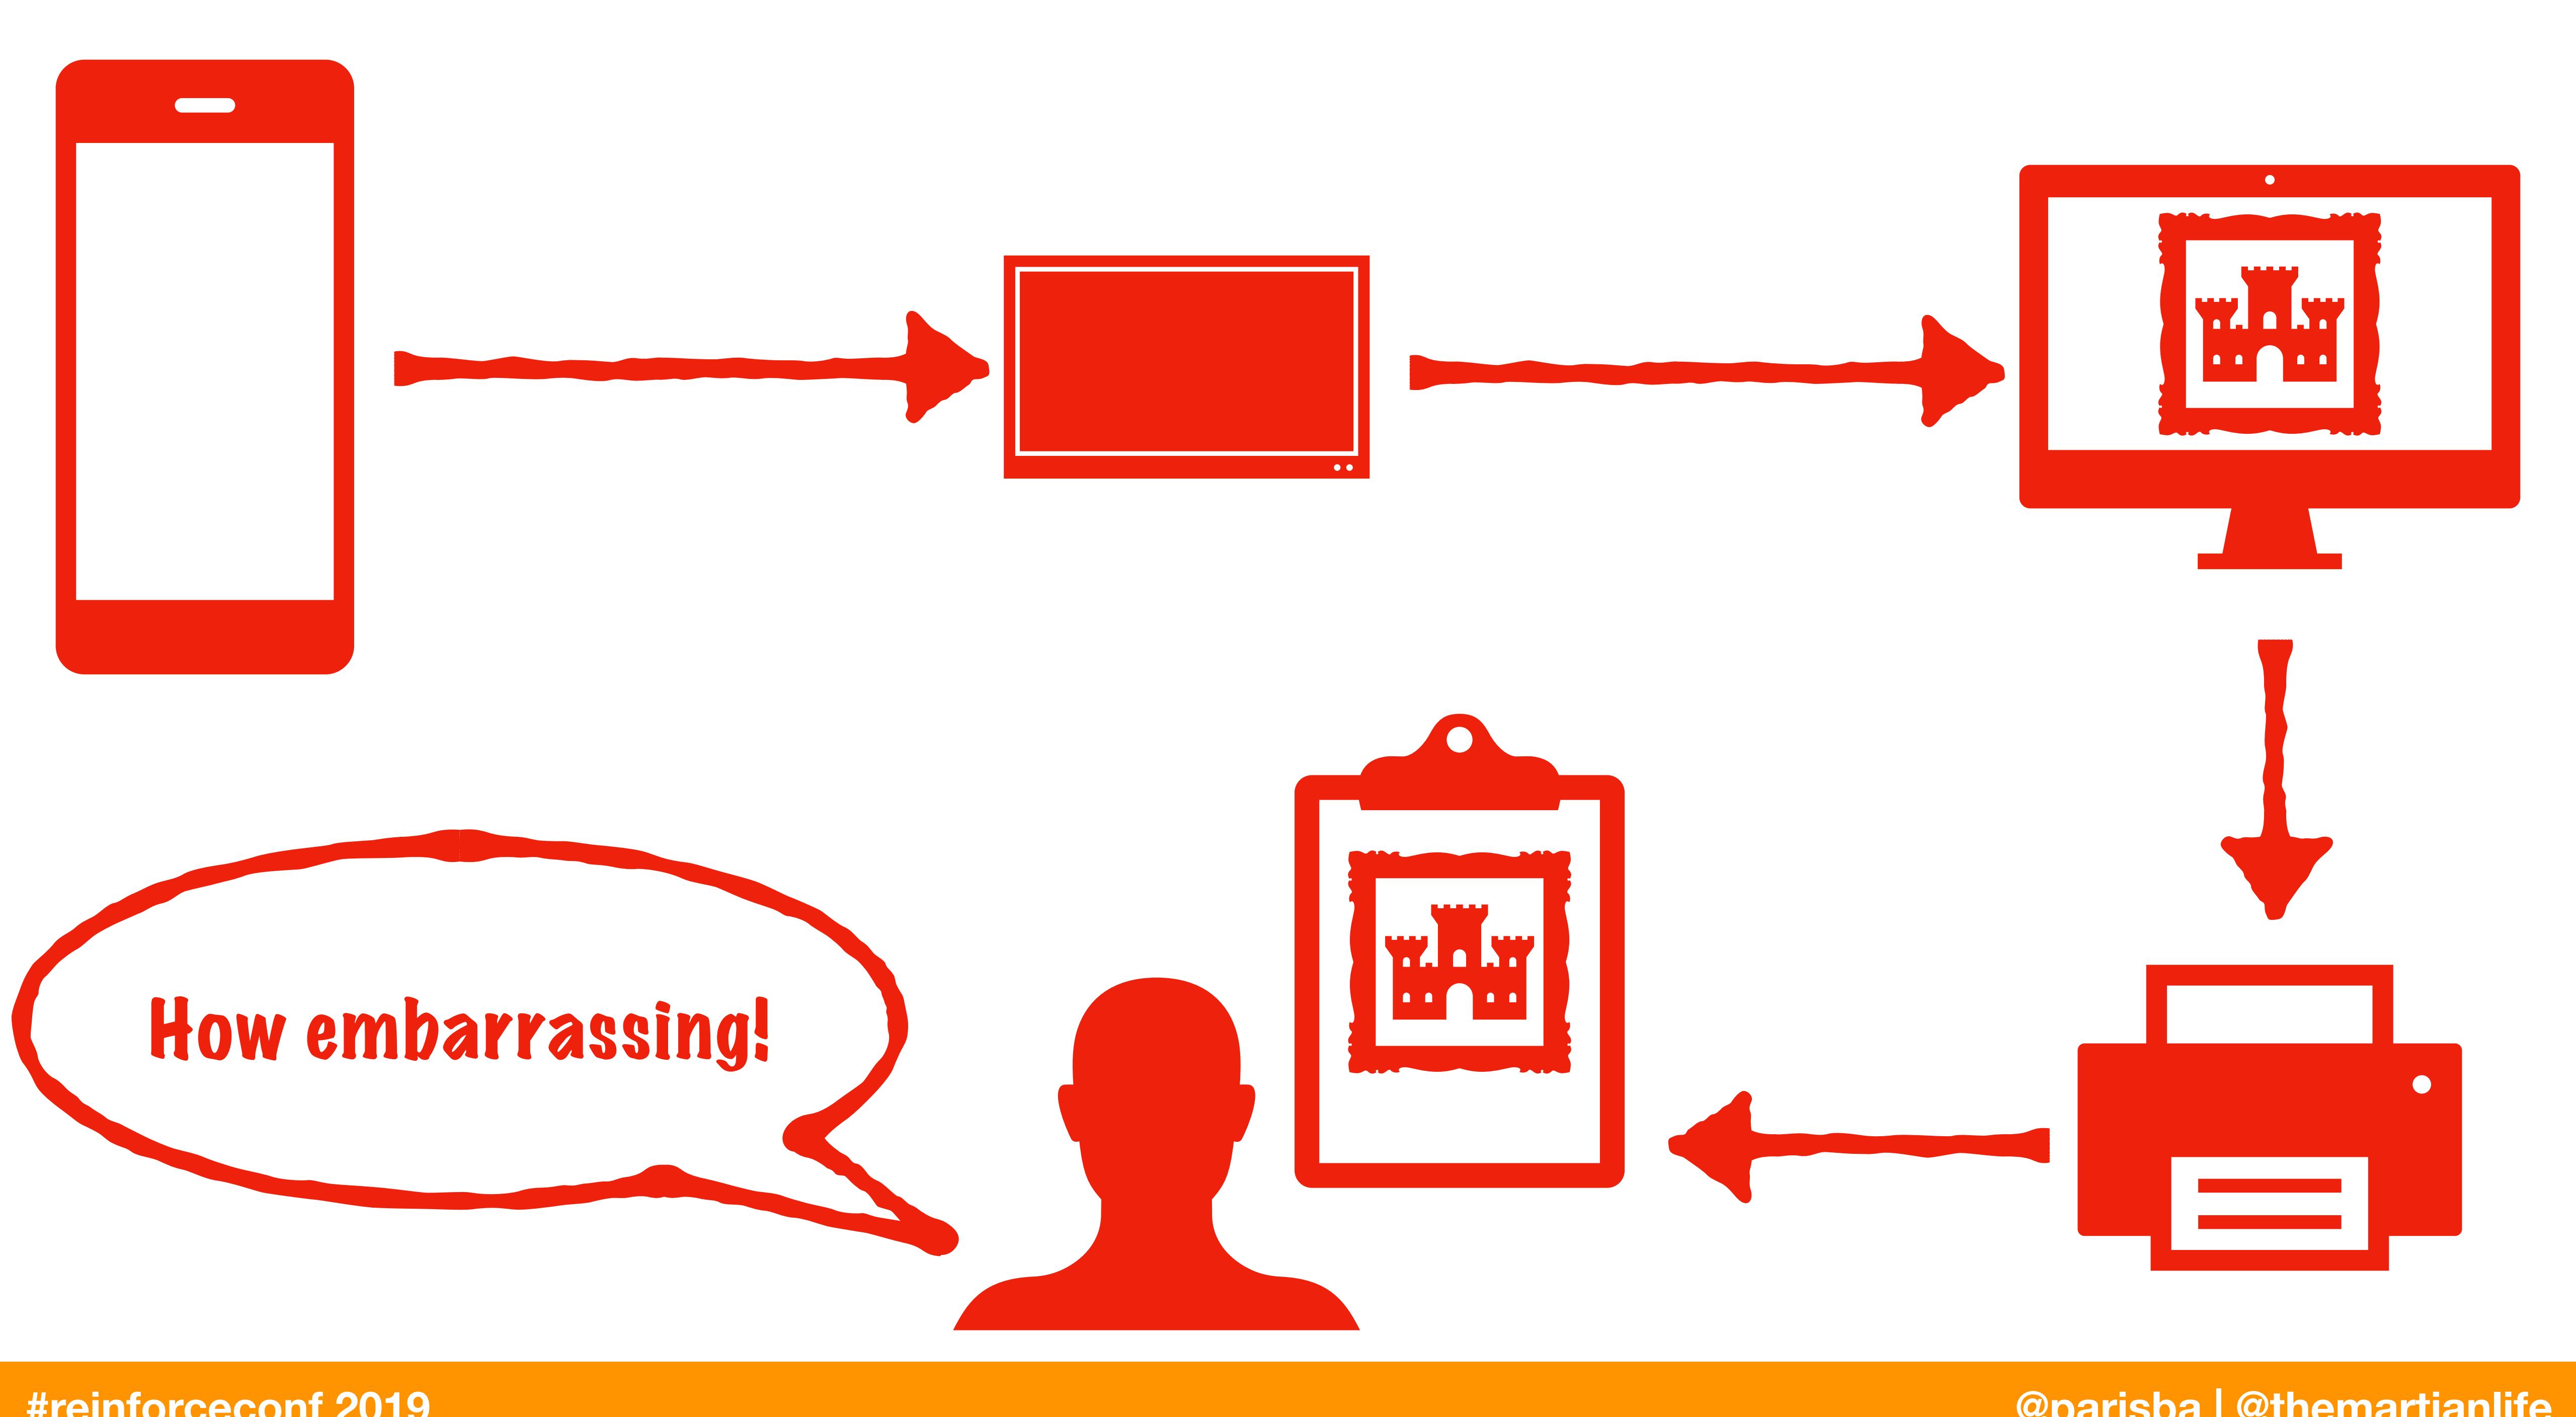


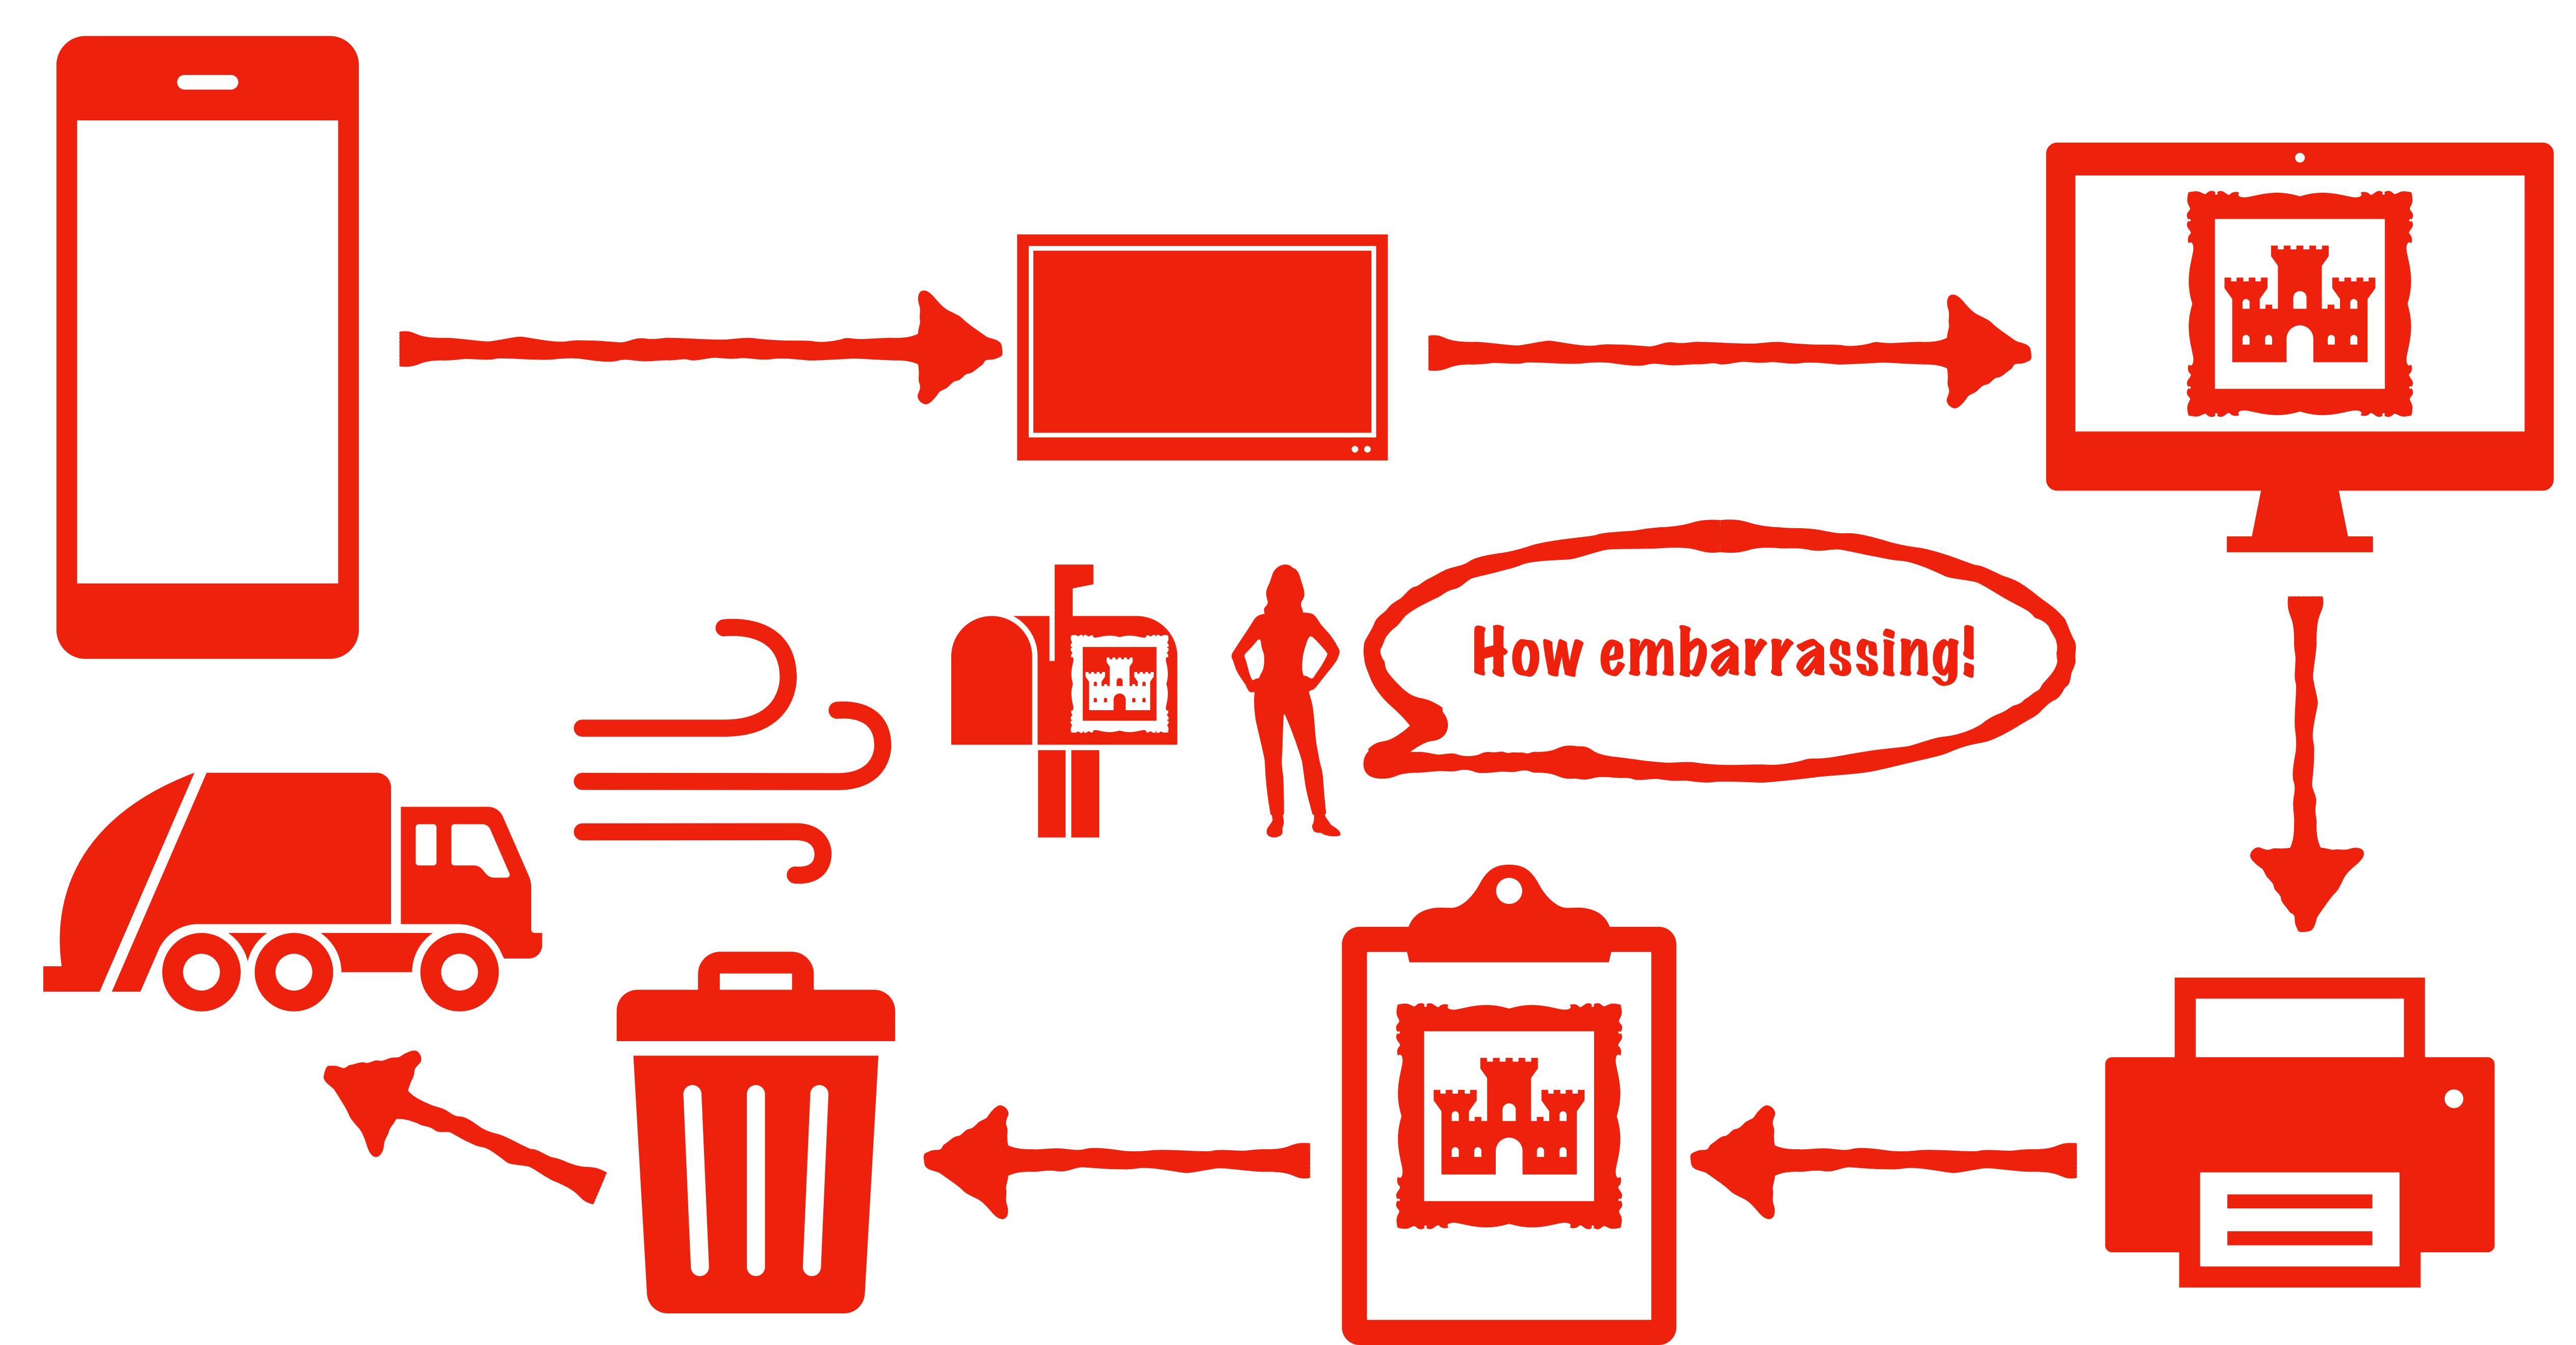
How embarrassing!





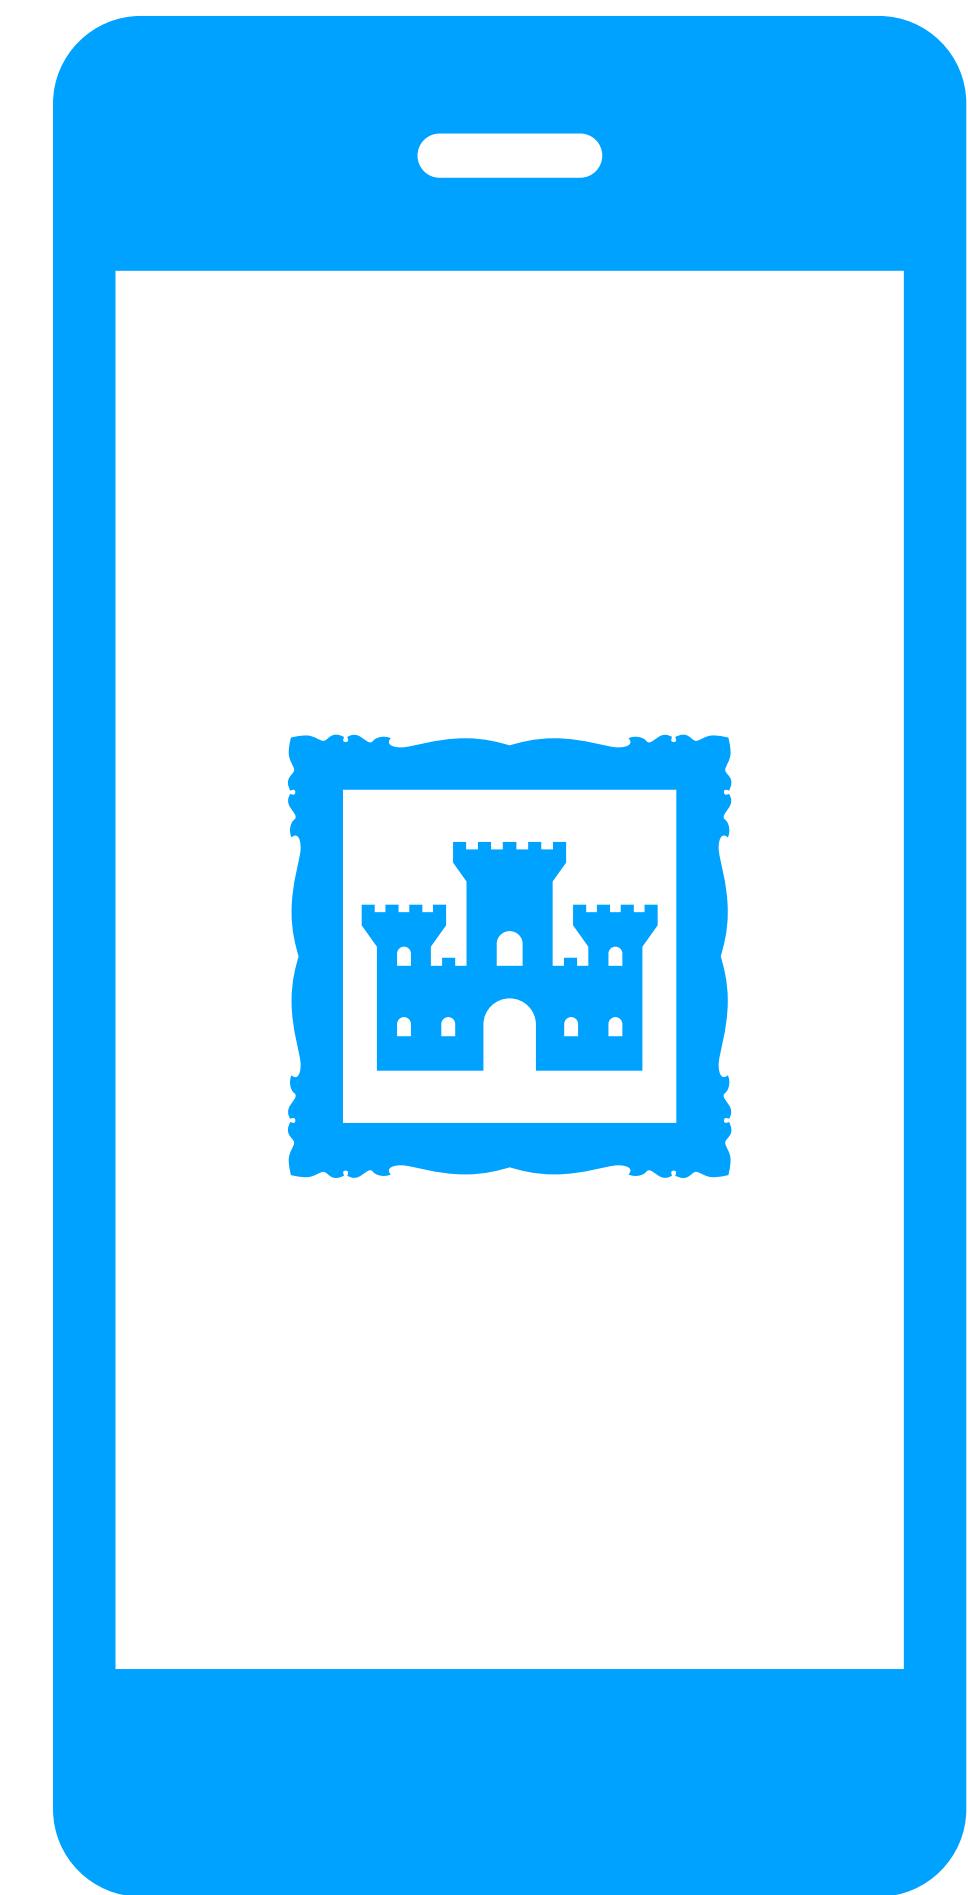






# No Control

(which might terrify you)

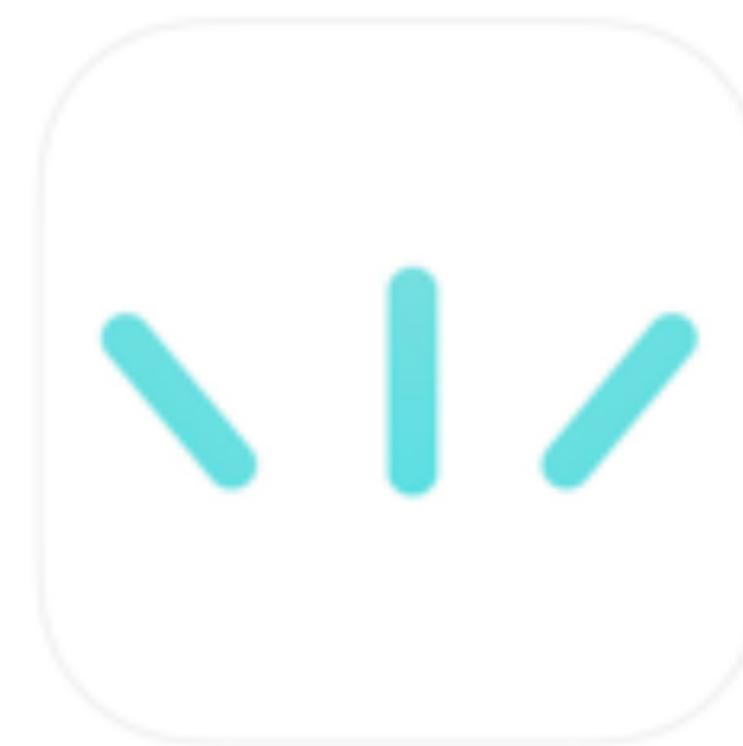


How embarrassing!

# Inability to share

==

# As good as secret



## MDacne - Custom Acne Treatment

Acne Treatment Crafted For You

MDalgorithms Inc.

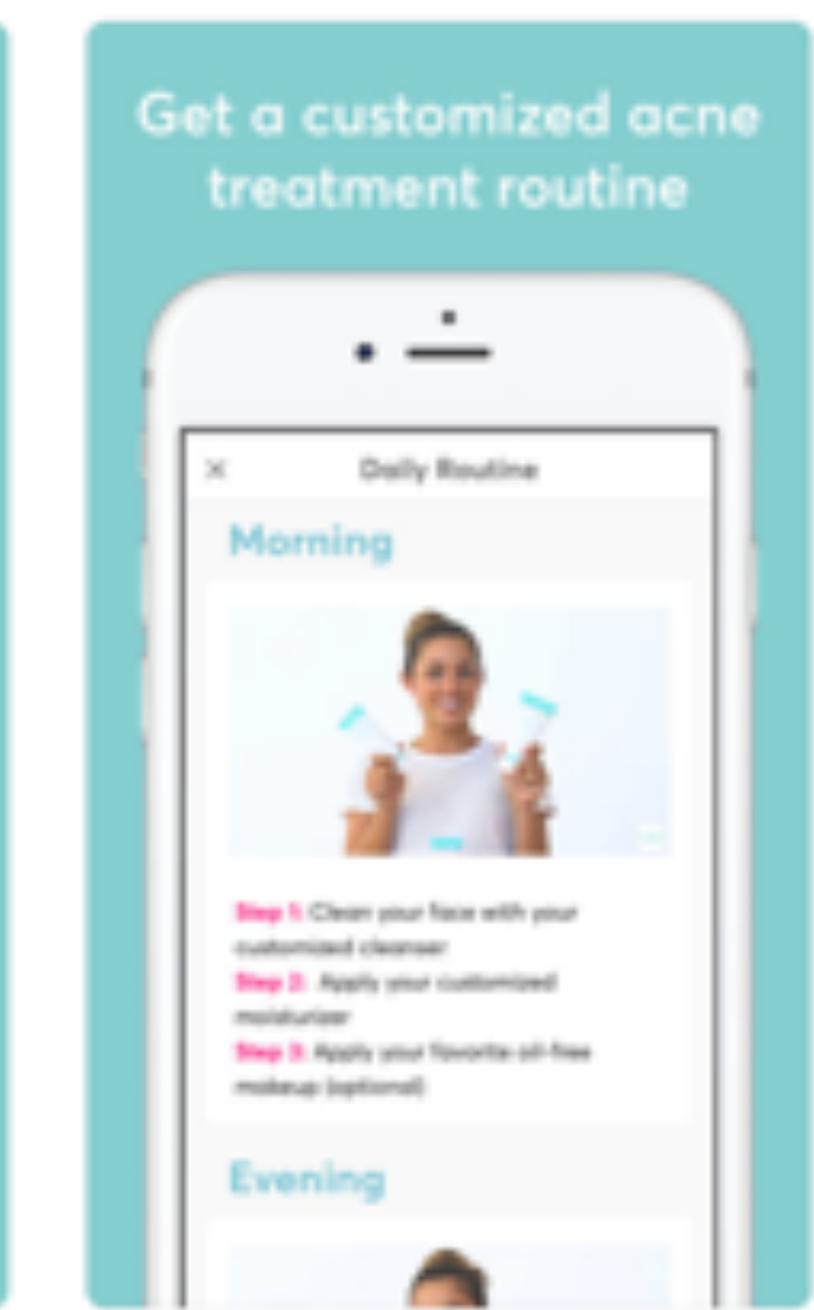
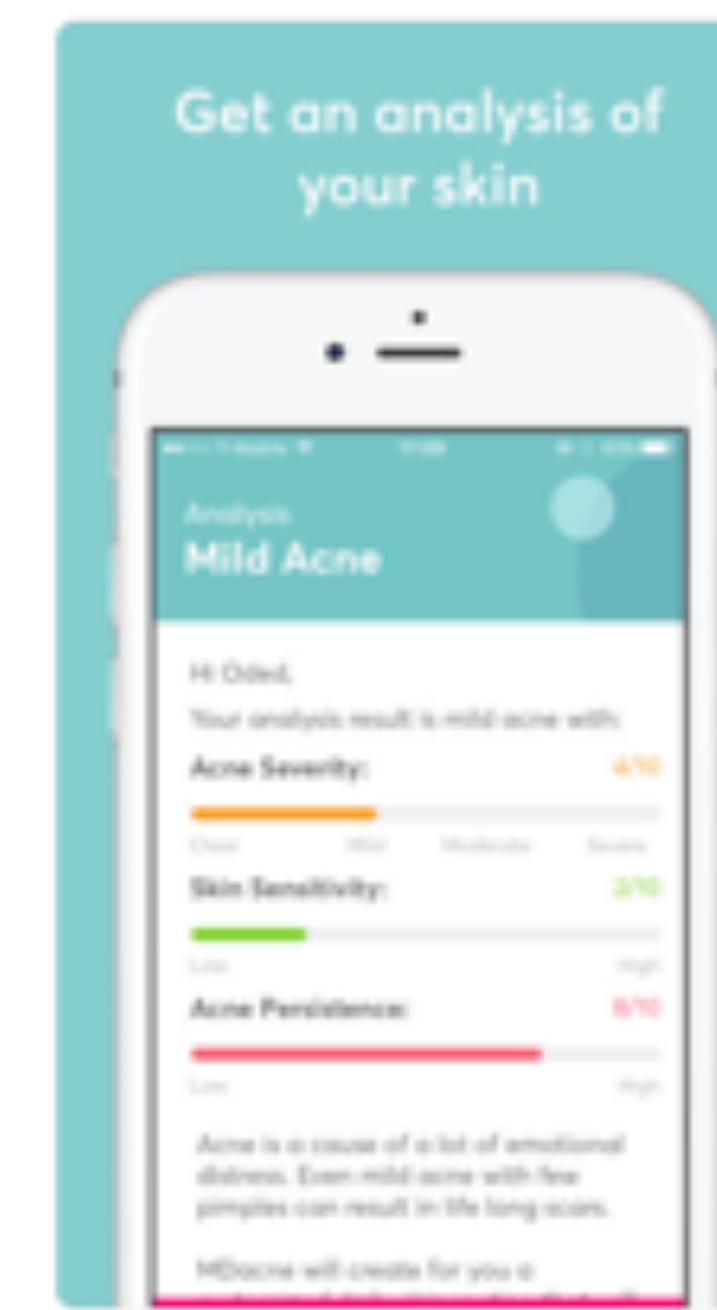
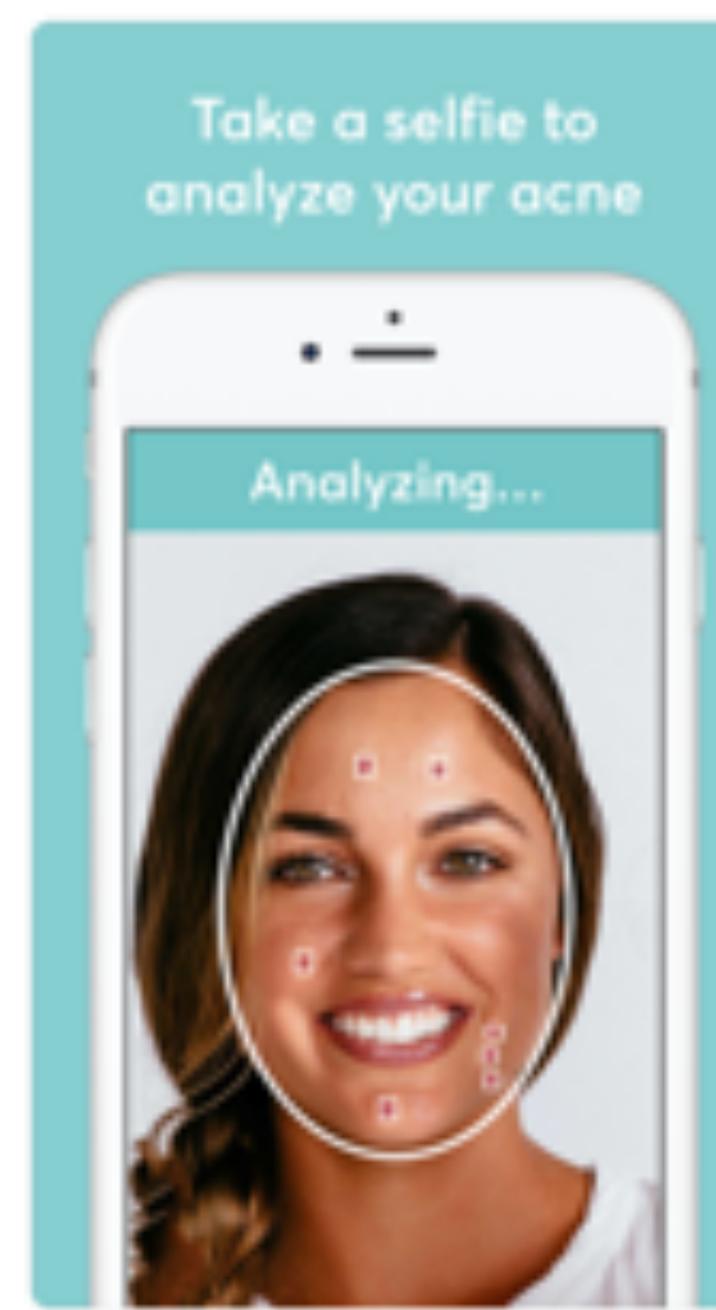
#6 in Medical

★★★★★ 4.3, 1.6K Ratings

Free

### Screenshots

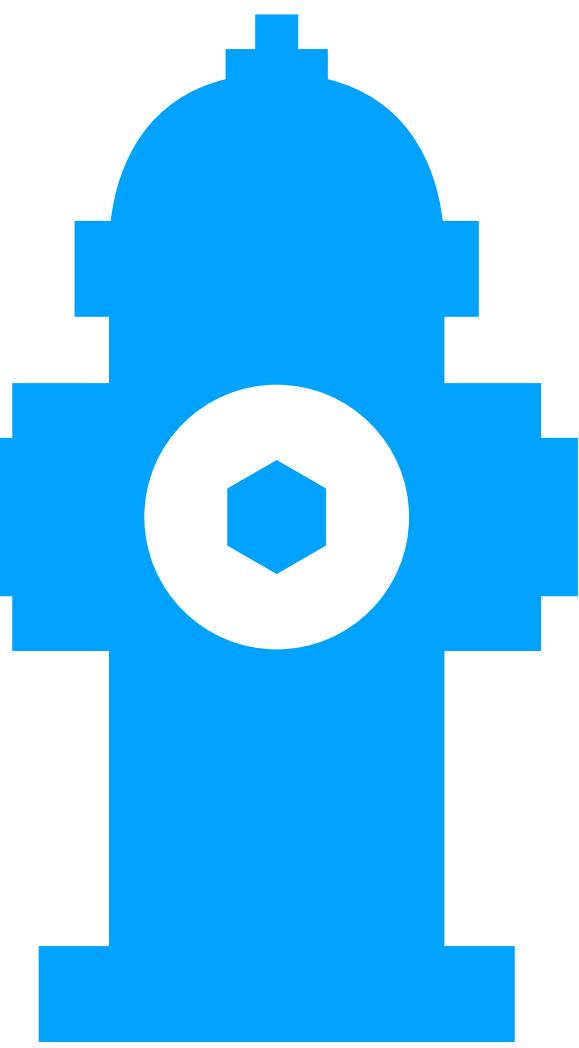
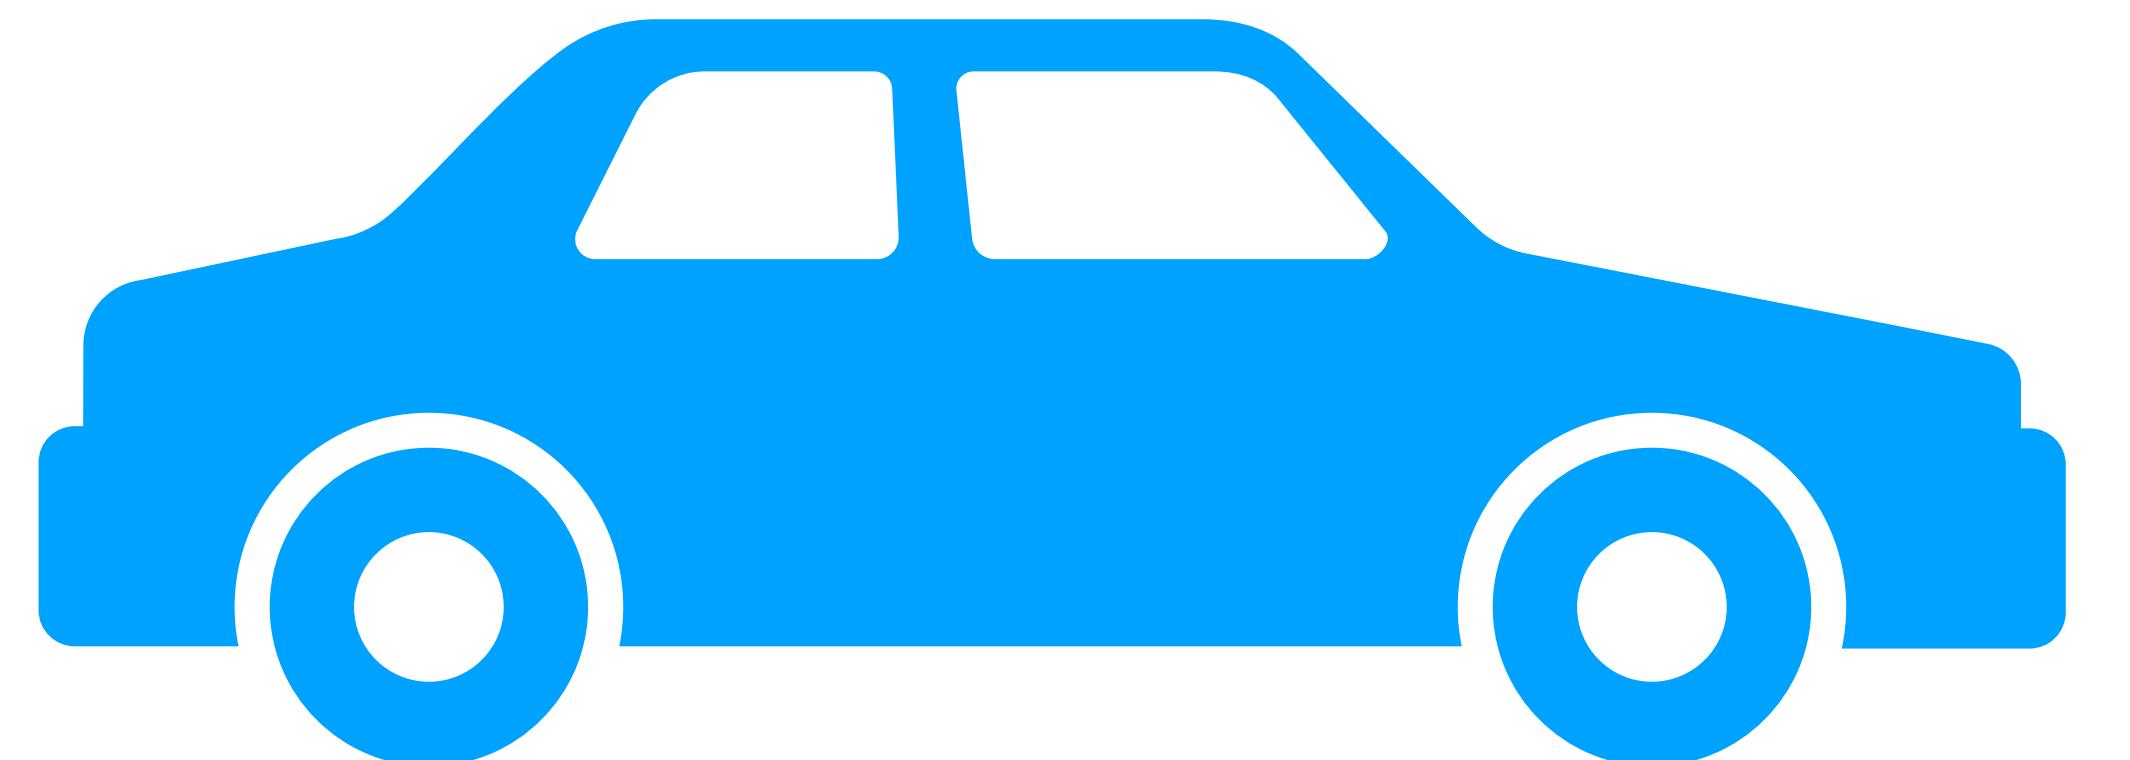
[iPhone](#) [iPad](#)



# Privacy

# **“Disconnected AI”**

Hey Model, should I  
brake?





PlantVillage

## Nuru

We have built an AI assistant called Nuru, which is Swahili for light. Nuru is an Artificially Intelligent that has been developed with the UN FAO, CGIAR, and other publicly funded institutions. As an assistant Nuru has learned to diagnose multiple diseases in Cassava, fall armyworm infections in African Maize, potato disease and wheat disease. She is also diagnosing spotted lanternfly pests in Pennsylvania.



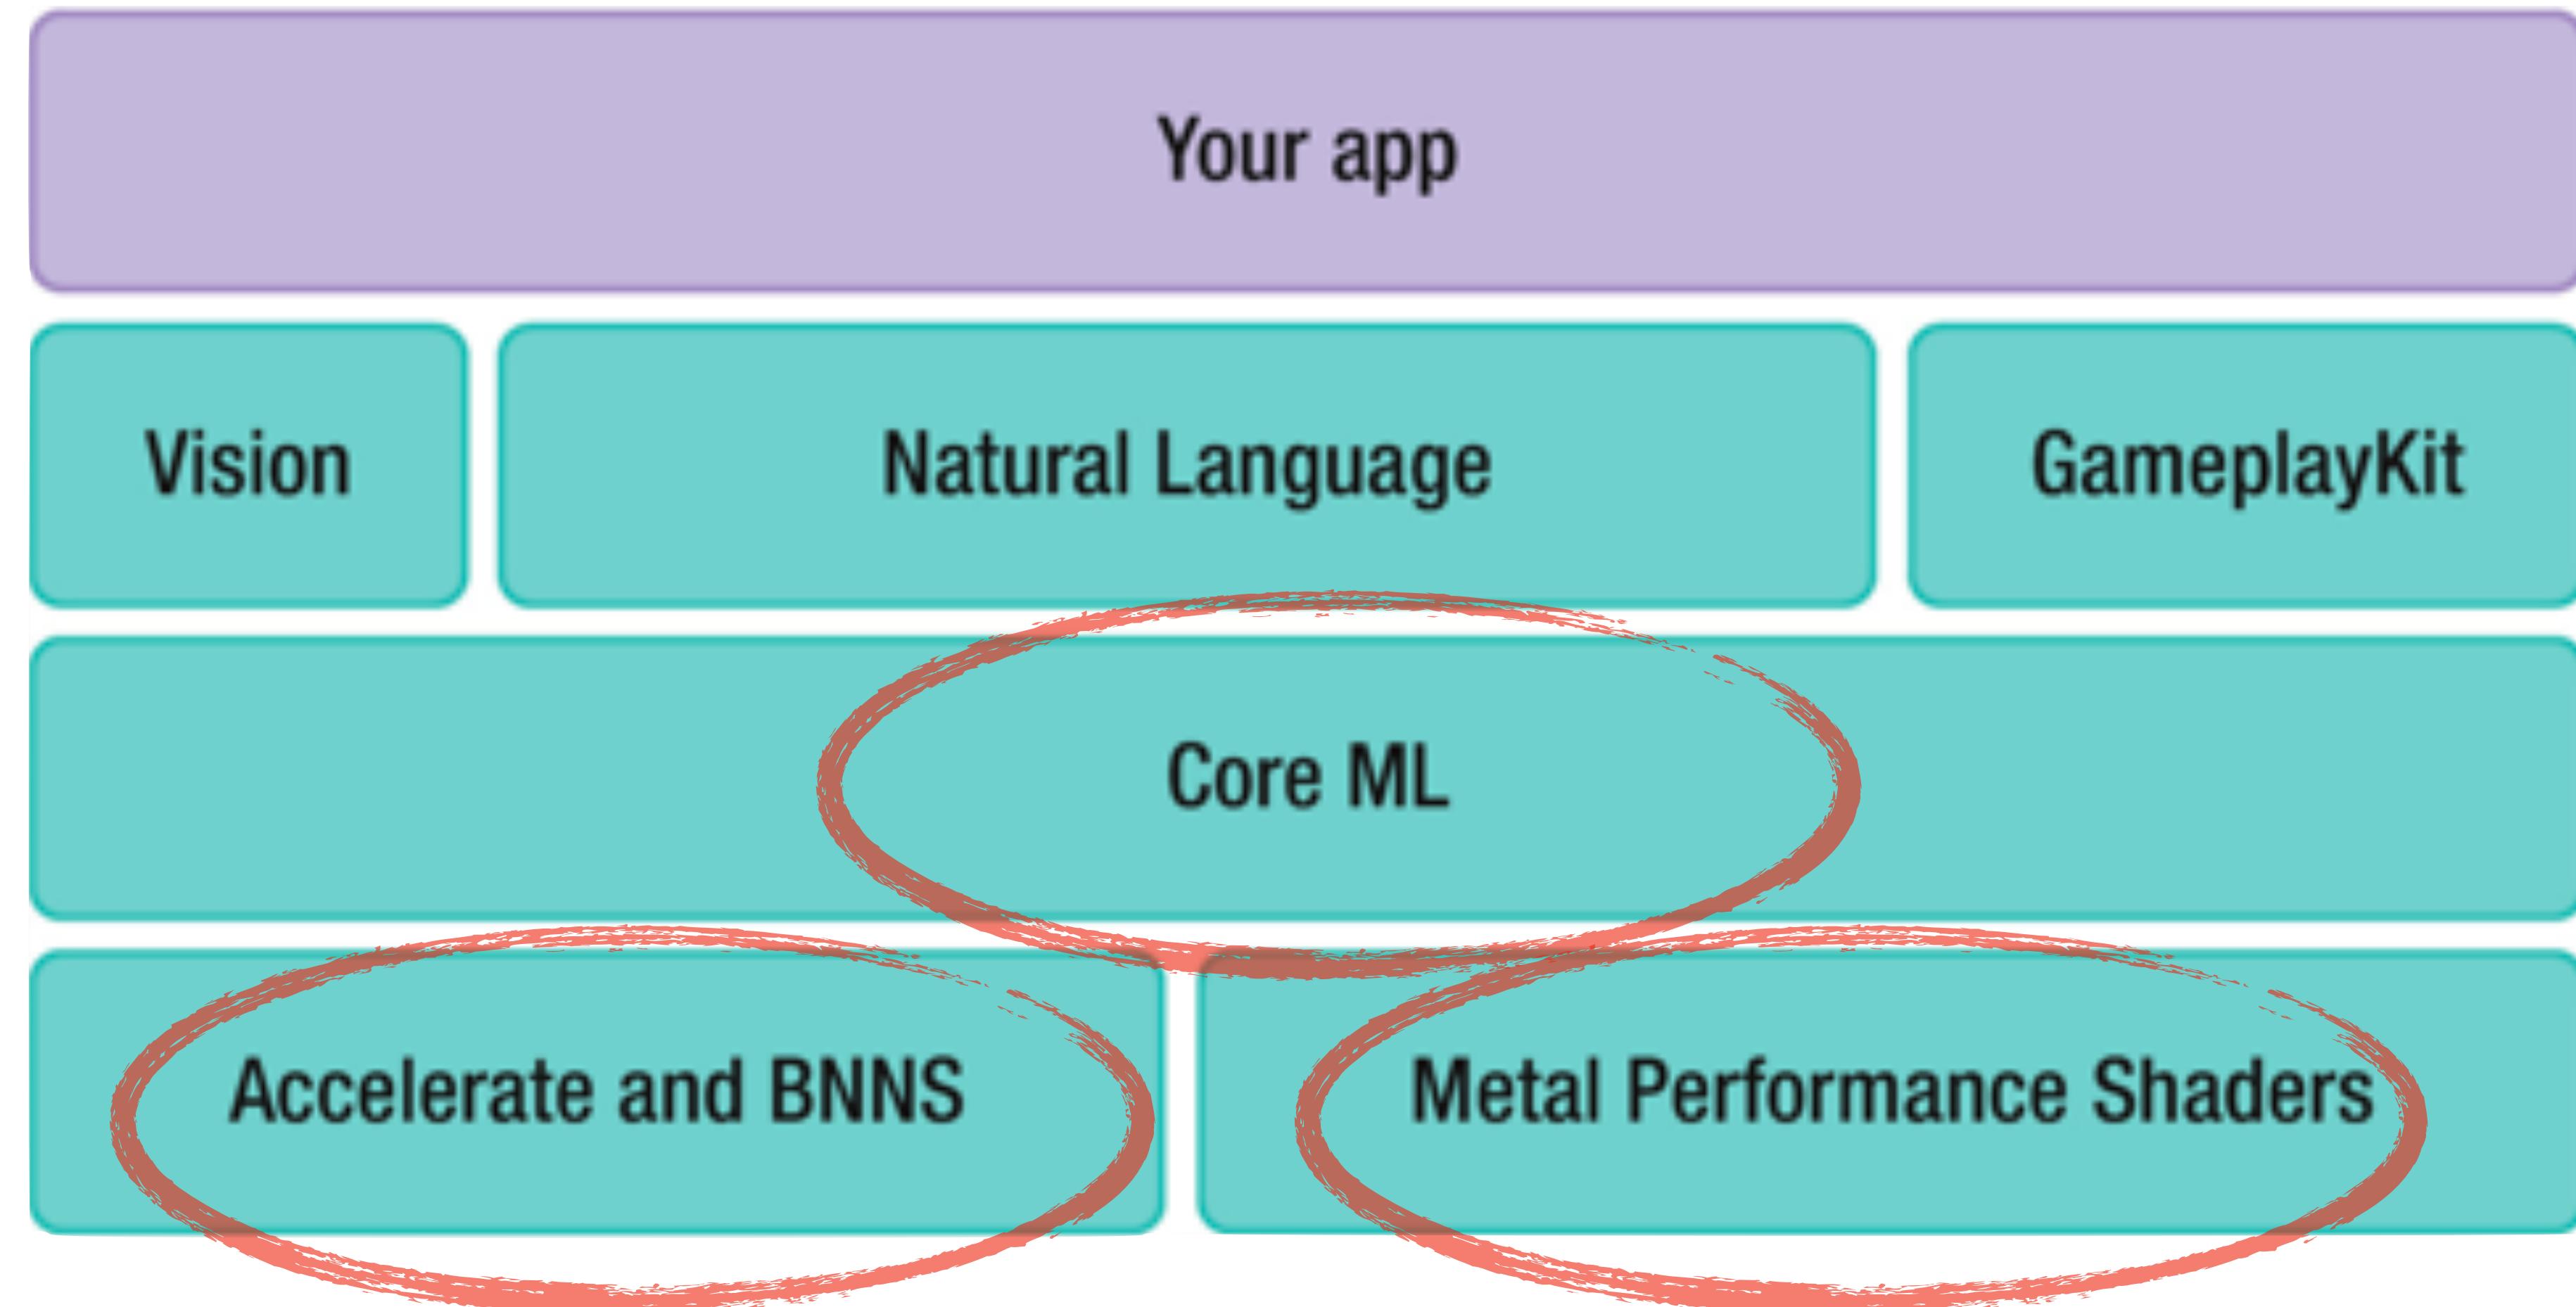
# Explain it with \$\$\$



# Privacy and Security

GDPR compliance is easier!

# Tools



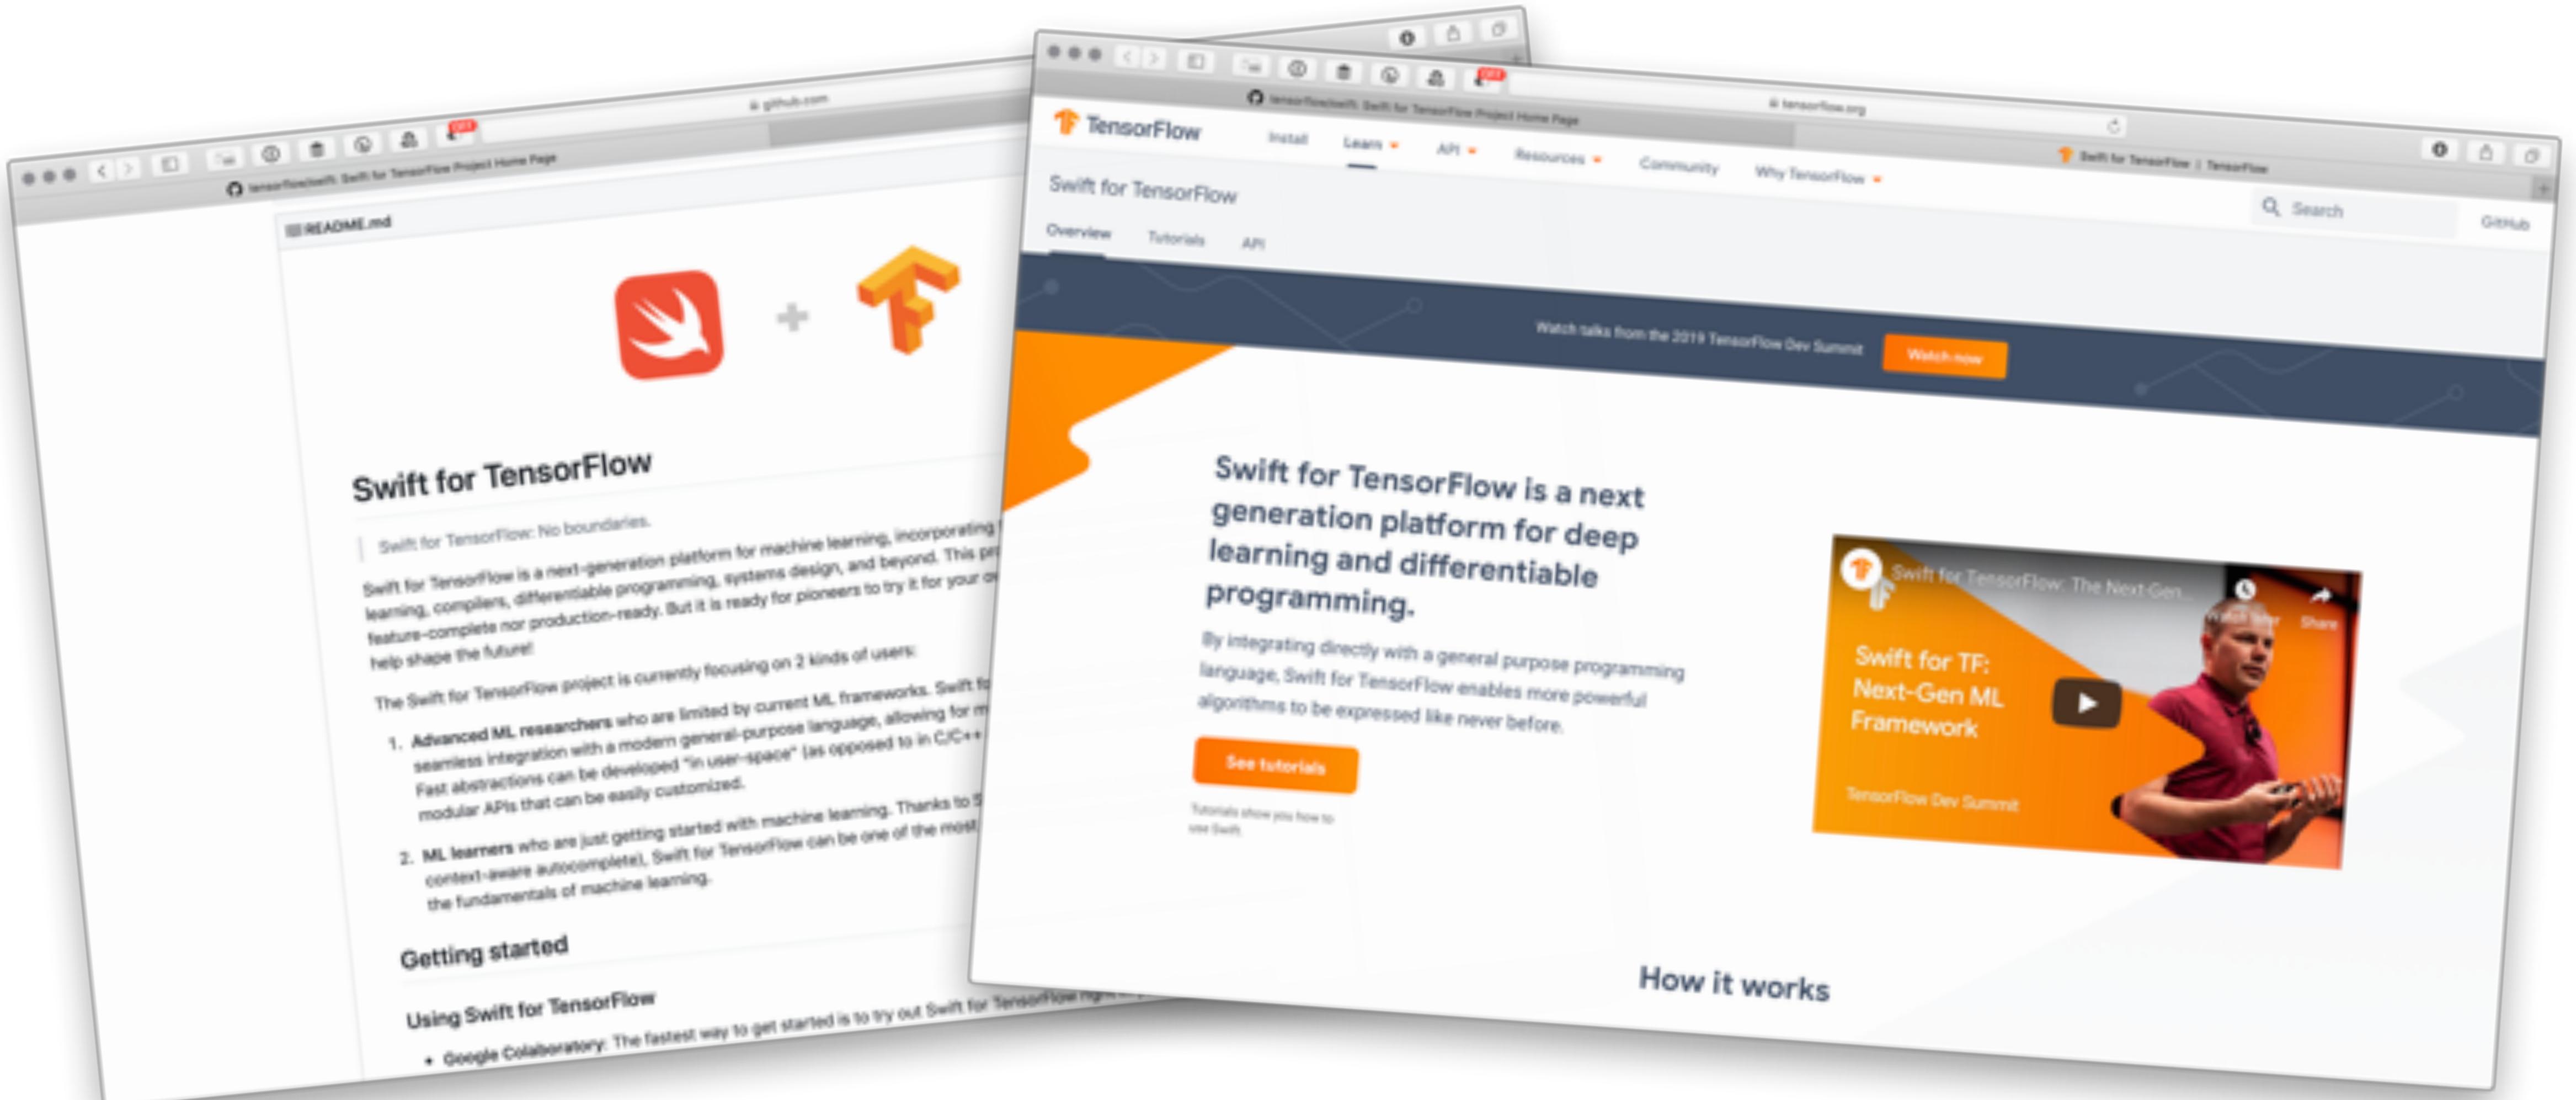


**XGBoost**



Caffe





# Hardware

Search or enter website name

NVIDIA, AUTONOMOUS MACHINES Develop Downloads Community Learn SDKs Buy Join Login

Home > Autonomous Machines > Buy > Jetson Nano Developer Kit

# Jetson Nano Developer Kit



The power of modern AI is now available for makers, learners, and embedded developers everywhere, for just \$99.

NVIDIA® Jetson Nano™ Developer Kit is a small, powerful computer that lets you run multiple neural networks in parallel for applications like image classification, object detection, segmentation, and speech processing. AI in an easy-to-use platform that runs in as little as 5-watts.

It's simpler than ever to get started! Just insert a microSD card with the system image, boot the developer kit, and enjoy using the same NVIDIA JetPack SDK used across the [NVIDIA Jetson™ family of products](#).

JetPack is compatible with NVIDIA's world-leading AI platform for training and deploying AI software, and reduces complexity and effort for developers by supporting many popular AI frameworks, like TensorFlow, PyTorch, Caffe, and MxNet. It also includes a full desktop Linux environment and out-of-the-box support for a variety of popular peripherals, add-ons, and ready-to-use projects.

Get started today with the Jetson Nano Developer Kit. We look forward to seeing what you create!

## Technical Specifications

The screenshot shows the Google Cloud website for the Edge TPU. At the top, there's a navigation bar with links for Google Cloud, Why Google, Solutions, Products, Pricing, Getting started, Docs, Support, and Sign in. Below the navigation is a search bar and two buttons: 'Contact sales' and 'Try free'. The main heading is 'Edge TPU' with the subtext 'EARLY ACCESS'. A subtext below it reads 'Google's purpose-built ASIC designed to run inference at the edge.' A blue 'REQUEST ACCESS' button is visible. The main content area features a large text block: 'With the explosive growth of connected devices, combined with a demand for privacy/confidentiality, low latency and bandwidth constraints, AI models trained in the cloud increasingly need to be run at the edge.' To the right of this text is a close-up image of a US penny coin, with two small black square components (representing the Edge TPU) placed on its surface. Below the main text is a smaller text box containing the following description: 'Google's purpose-built ASIC designed to run AI at the edge. It delivers high performance in a small physical and power footprint, enabling the deployment of high-accuracy AI at the edge.'

Google Cloud Why Google Solutions Products Pricing Getting started

cloud.google.com

Contact sales Try free

EARLY ACCESS

# Edge TPU

Google's purpose-built ASIC designed to run inference at the edge.

REQUEST ACCESS

With the explosive growth of connected devices, combined with a demand for privacy/confidentiality, low latency and bandwidth constraints, AI models trained in the cloud increasingly need to be run at the edge.

Google's purpose-built ASIC designed to run AI at the edge. It delivers high performance in a small physical and power footprint, enabling the deployment of high-accuracy AI at the edge.

# Tensor Processing Unit

[https://en.wikipedia.org/wiki/Tensor\\_processing\\_unit](https://en.wikipedia.org/wiki/Tensor_processing_unit)

# Task-based AI/ML

ML Task	Description
Recommender	Personalize choices for users
Image Classification	Label images
Object Detection	Recognize objects within images
Style Transfer	Stylize images
Activity Classification	Detect an activity using sensors
Image Similarity	Find similar images
Classifiers	Predict a label
Regression	Predict numeric values
Clustering	Group similar datapoints together
Text Classifier	Analyze sentiment of messages

# Task-based AI/ML

**New interest and motivations**

- + better software
- + better hardware
- + task focused frameworks

=> On-device ML 

*Part 1*

## On-device Machine Learning

*Part 2*

## Neural Style Transfer

**Thank You!**

*art*

Demonstration

*Part 4*

## The Future

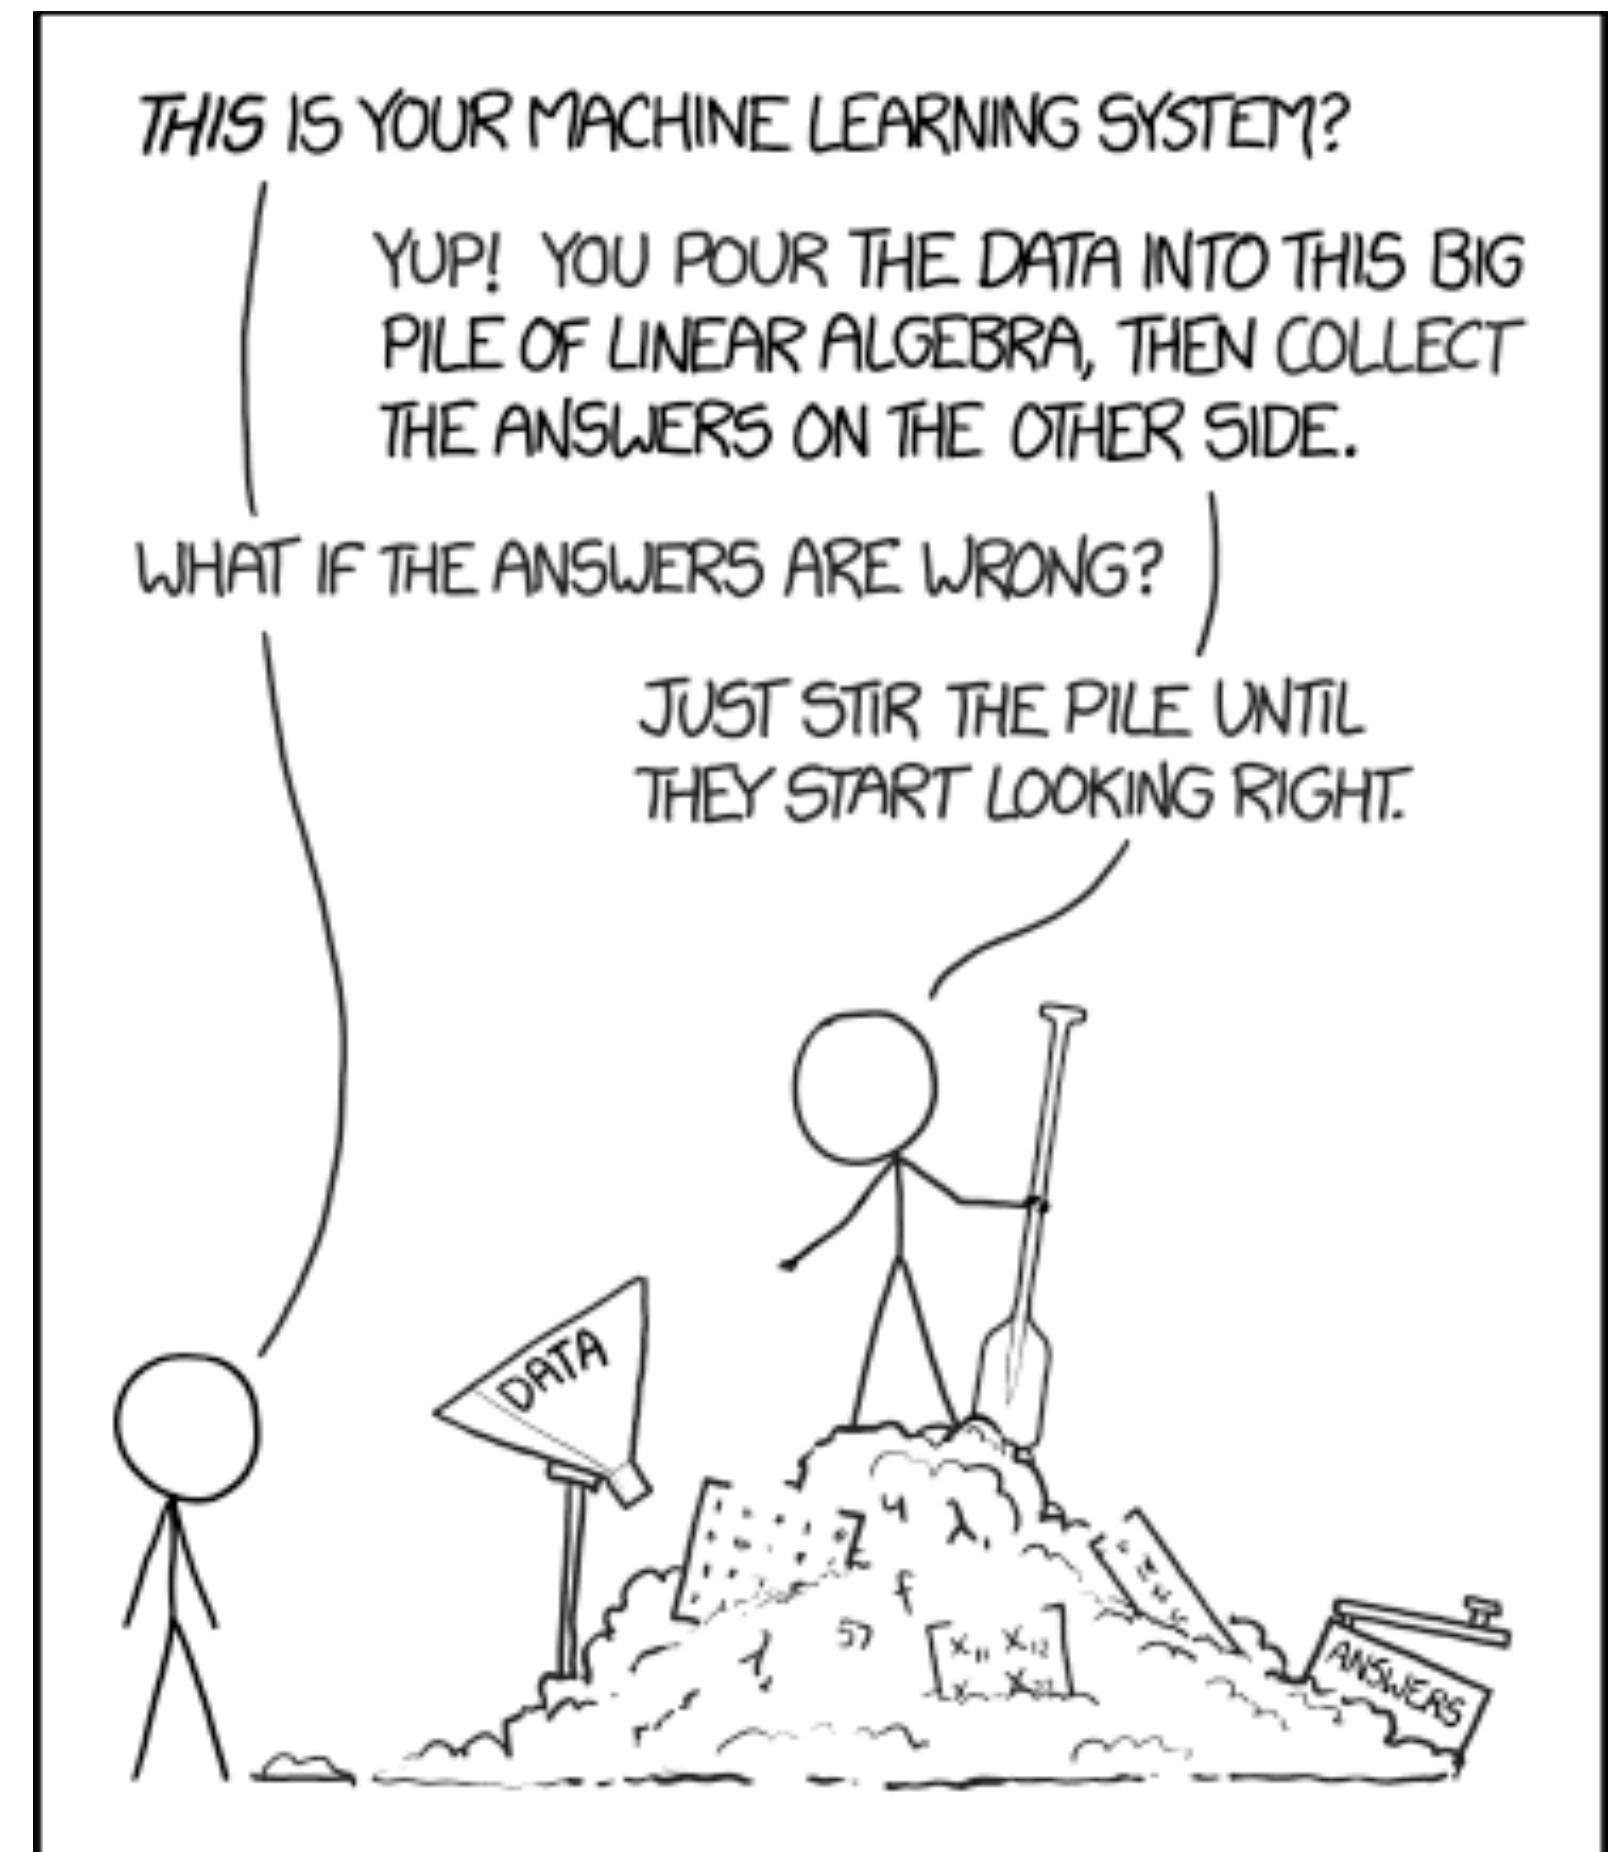
# Goals of this talk

- Discuss how useful, effective, and sensible the idea of on-device artificial intelligence is
- Have a brief diversion to discuss what Style Transfer is, and gloss over how it works
- Demonstrate that on-device AI works, by implementing Style Transfer for it on iOS
- Show off how useful task-focused approaches to AI can be, and how useful Turi Create is

# Great resources

(You should check them out)

- <https://developer.apple.com/videos/play/wwdc2018/712/>
- <https://github.com/apple/turicreate/>
- [https://github.com/alexsosn/iOS\\_ML](https://github.com/alexsosn/iOS_ML)
- <https://www.manning.com/books/grokking-deep-learning>
- <https://medium.com/artists-and-machine-intelligence/neural-artistic-style-transfer-a-comprehensive-look-f54d8649c199>



# Stay in touch

- Tweet at us 🐥
- Stalk us on the internet:
  - ▶ [www.paris.id.au](http://www.paris.id.au)
  - ▶ [www.themartianlife.com](http://www.themartianlife.com)
- Say g'day around the conf! 🙌
- Slides + code👉 [github.com/  
thesecretlab/ReinforceConf2019](https://github.com/thesecretlab/ReinforceConf2019)

